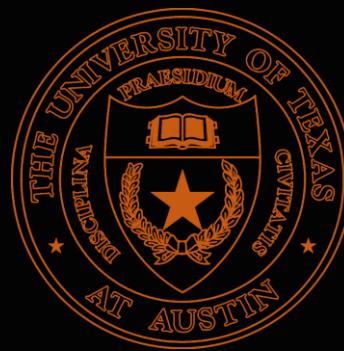
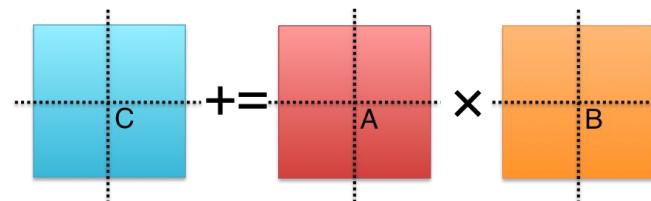


The Science of
High-Performance
Computing Group



STRASSEN

Implementing Strassen-like Fast Matrix Multiplication Algorithms with BLIS



Jianyu Huang, Leslie Rice

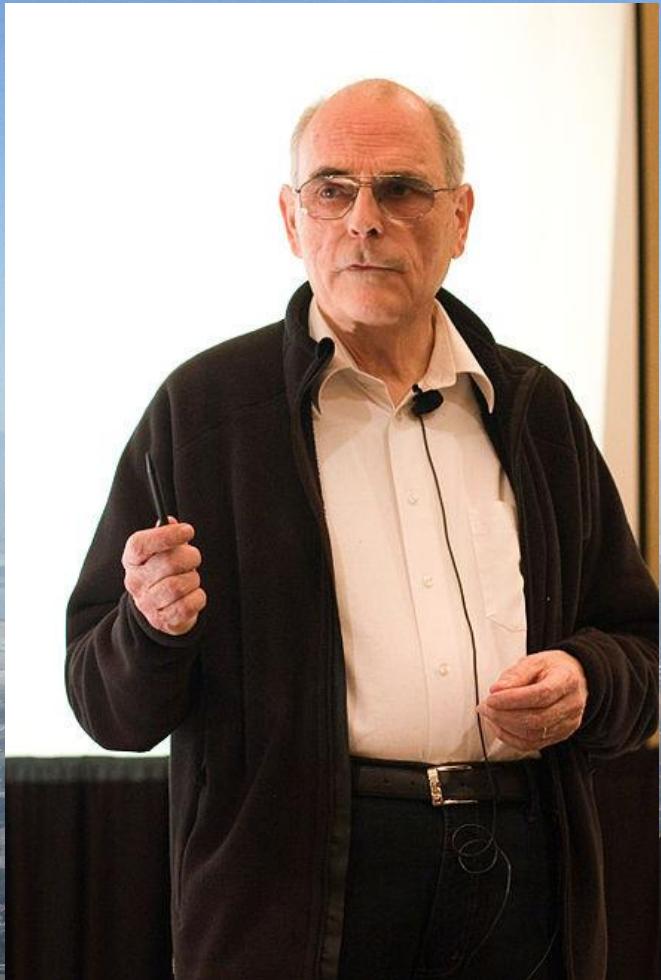
Joint work with Tyler M. Smith, Greg M. Henry, Robert A. van de Geijn

BLIS Retreat 2016

*Overlook of the Bay Area. Photo taken in Mission Peak Regional Preserve, Fremont, CA. Summer 2014.

STRASSEN, from 30,000 feet

Volker Strassen
(Born in 1936, aged 80)



Original Strassen Paper (1969)

Numer. Math. 13, 354–356 (1969)

Gaussian Elimination is not Optimal

VOLKER STRASSEN*

Received December 12, 1968

1. Below we will give an algorithm which computes the coefficients of the product of two square matrices A and B of order n from the coefficients of A and B with less than $4.7 \cdot n^{\log 7}$ arithmetical operations (all logarithms in this paper are for base 2, thus $\log 7 \approx 2.8$; the usual method requires approximately $2n^3$ arithmetical operations). The algorithm induces algorithms for inverting a matrix of order n , solving a system of n linear equations in n unknowns, computing a determinant of order n etc. all requiring less than $\text{const } n^{\log 7}$ arithmetical operations.

This fact should be compared with the result of KLYUYEV and KOKOVKIN-SCHERBAK [1] that Gaussian elimination for solving a system of linear equations is optimal if one restricts oneself to operations upon rows and columns as a whole. We also note that WINOGRAD [2] modifies the usual algorithms for matrix multiplication and inversion and for solving systems of linear equations, trading roughly half of the multiplications for additions and subtractions.

It is a pleasure to thank D. BRILLINGER for inspiring discussions about the present subject and St. COOK and B. PARLETT for encouraging me to write this paper.

2. We define algorithms $\alpha_{m,k}$ which multiply matrices of order $m2^k$, by induction on k : $\alpha_{m,0}$ is the usual algorithm for matrix multiplication (requiring m^3 multiplications and $m^2(m-1)$ additions). $\alpha_{m,k}$ already being known, define $\alpha_{m,k+1}$ as follows:

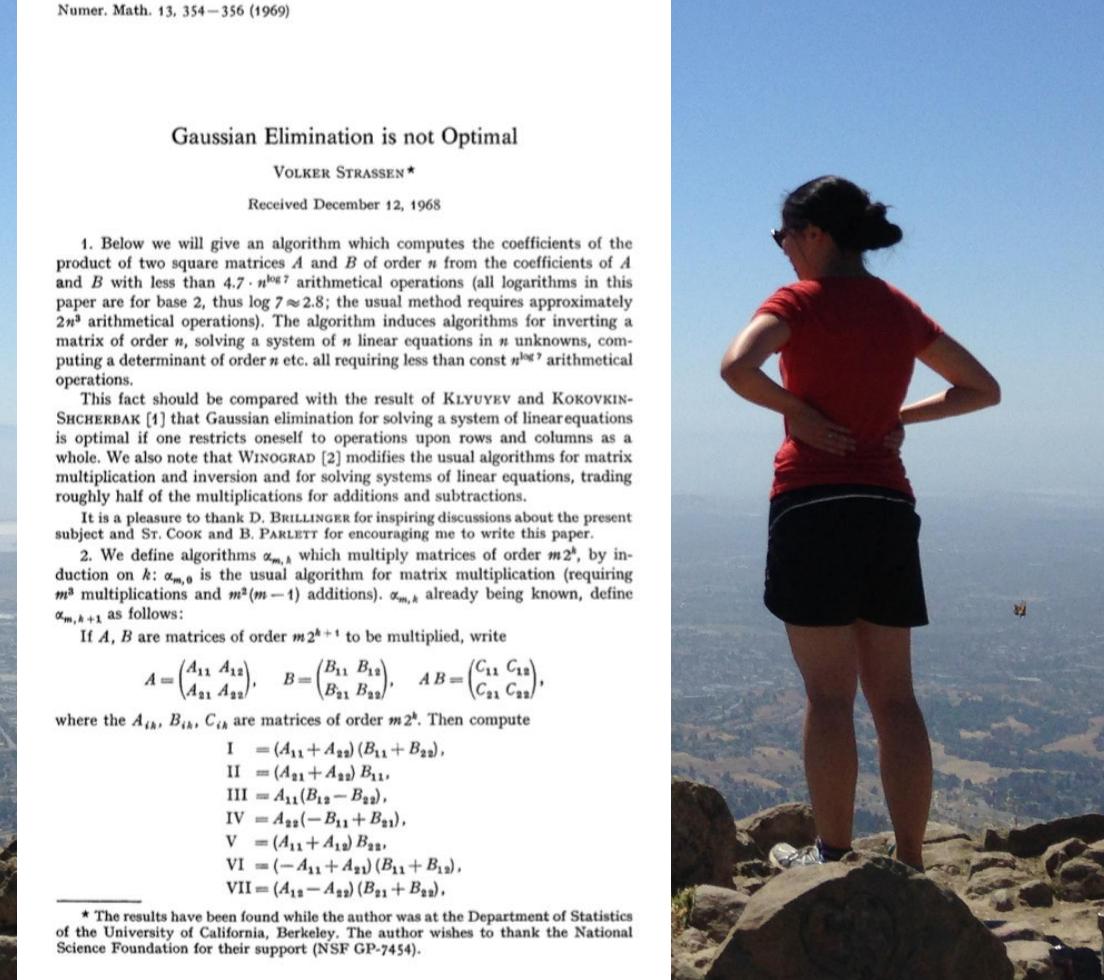
If A, B are matrices of order $m2^{k+1}$ to be multiplied, write

$$A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}, \quad B = \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}, \quad AB = \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix},$$

where the A_{ik}, B_{ik}, C_{ik} are matrices of order $m2^k$. Then compute

$$\begin{aligned} \text{I} &= (A_{11} + A_{21})(B_{11} + B_{21}), \\ \text{II} &= (A_{21} + A_{22})B_{11}, \\ \text{III} &= A_{11}(B_{12} - B_{22}), \\ \text{IV} &= A_{22}(-B_{11} + B_{21}), \\ \text{V} &= (A_{11} + A_{12})B_{22}, \\ \text{VI} &= (-A_{11} + A_{21})(B_{11} + B_{12}), \\ \text{VII} &= (A_{12} - A_{22})(B_{21} + B_{22}), \end{aligned}$$

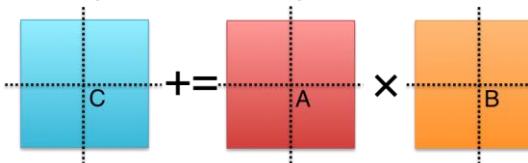
* The results have been found while the author was at the Department of Statistics of the University of California, Berkeley. The author wishes to thank the National Science Foundation for their support (NSF GP-7454).



One-level Strassen's Algorithm (In theory)

Assume m , n , and k are all even. A , B , and C are $m \times k$, $k \times n$, $m \times n$ matrices, respectively. Letting

$$C = \begin{pmatrix} C_{00} & C_{01} \\ C_{10} & C_{11} \end{pmatrix}, A = \begin{pmatrix} A_{00} & A_{01} \\ A_{10} & A_{11} \end{pmatrix}, B = \begin{pmatrix} B_{00} & B_{01} \\ B_{10} & B_{11} \end{pmatrix}$$



We can compute $C := C + AB$ by

Direct Computation

$$\begin{aligned} C_{00} &:= A_{00}B_{00} + A_{01}B_{10} + C_{00}; \\ C_{01} &:= A_{00}B_{01} + A_{01}B_{11} + C_{01}; \\ C_{10} &:= A_{10}B_{00} + A_{11}B_{10} + C_{10}; \\ C_{11} &:= A_{10}B_{01} + A_{11}B_{11} + C_{11}; \end{aligned}$$

8 multiplications, 8 additions

Strassen's Algorithm

$$\begin{aligned} M_0 &:= (A_{00} + A_{11})(B_{00} + B_{11}); \\ M_1 &:= (A_{10} + A_{11})B_{00}; \\ M_2 &:= A_{00}(B_{01} - B_{11}); \\ M_3 &:= A_{11}(B_{10} - B_{00}); \\ M_4 &:= (A_{00} + A_{01})B_{11}; \\ M_5 &:= (A_{10} - A_{00})(B_{00} + B_{01}); \\ M_6 &:= (A_{01} - A_{11})(B_{10} + B_{11}); \\ C_{00} &:= M_0 + M_3 - M_4 + M_7 + C_{00}; \\ C_{01} &:= M_2 + M_4 + C_{01}; \\ C_{10} &:= M_1 + M_3 + C_{10}; \\ C_{11} &:= M_0 - M_1 + M_2 + M_5 + C_{11}. \end{aligned}$$

7 multiplications, 22 additions

*Strassen, Volker. "Gaussian elimination is not optimal." *Numerische Mathematik* 13, no. 4 (1969): 354-356.

Multi-level Strassen's Algorithm (In theory)

$$M_0 := (A_{00} + A_{11})(B_{00} + B_{11});$$

$$M_1 := (A_{10} + A_{11})B_{00};$$

$$M_2 := A_{00}(B_{01} - B_{11});$$

$$M_3 := A_{11}(B_{10} - B_{00});$$

$$M_4 := (A_{00} + A_{01})B_{11};$$

$$M_5 := (A_{10} - A_{00})(B_{00} + B_{01});$$

$$M_6 := (A_{01} - A_{11})(B_{10} + B_{11});$$

$$C_{00} += M_0 + M_3 - M_4 + M_6$$

$$C_{01} += M_2 + M_4$$

$$C_{10} += M_1 + M_3$$

$$C_{11} += M_0 - M_1 + M_2 + M_5$$

- One-level Strassen (**1+14.3%** speedup)
 - 8 multiplications → 7 multiplications ;
- Two-level Strassen (**1+30.6%** speedup)
 - 64 multiplications → 49 multiplications;
- d -level Strassen (**$n^3/n^{2.803}$** speedup)
 - 8^d multiplications → 7^d multiplications;
If originally $m = n = k = 2^d$, where d is an integer,
then the cost becomes
 $(7/8)^{\log_2(n)} 2n^3 = n^{\log_2(7/8)} 2n^3 \approx 2n^{2.807}$ flops.

Multi-level Strassen's Algorithm (In theory)

$$M_0 := (A_{00} + A_{11})(B_{00} + B_{11});$$

$$M_1 := (A_{10} + A_{11})B_{00};$$

$$M_2 := A_{00}(B_{01} - B_{11});$$

$$M_3 := A_{11}(B_{10} - B_{00});$$

$$M_4 := (A_{00} + A_{01})B_{11};$$

$$M_5 := (A_{10} - A_{00})(B_{00} + B_{01});$$

$$M_6 := (A_{01} - A_{11})(B_{10} + B_{11});$$

$$C_{00} += M_0 + M_3 - M_4 + M_6$$

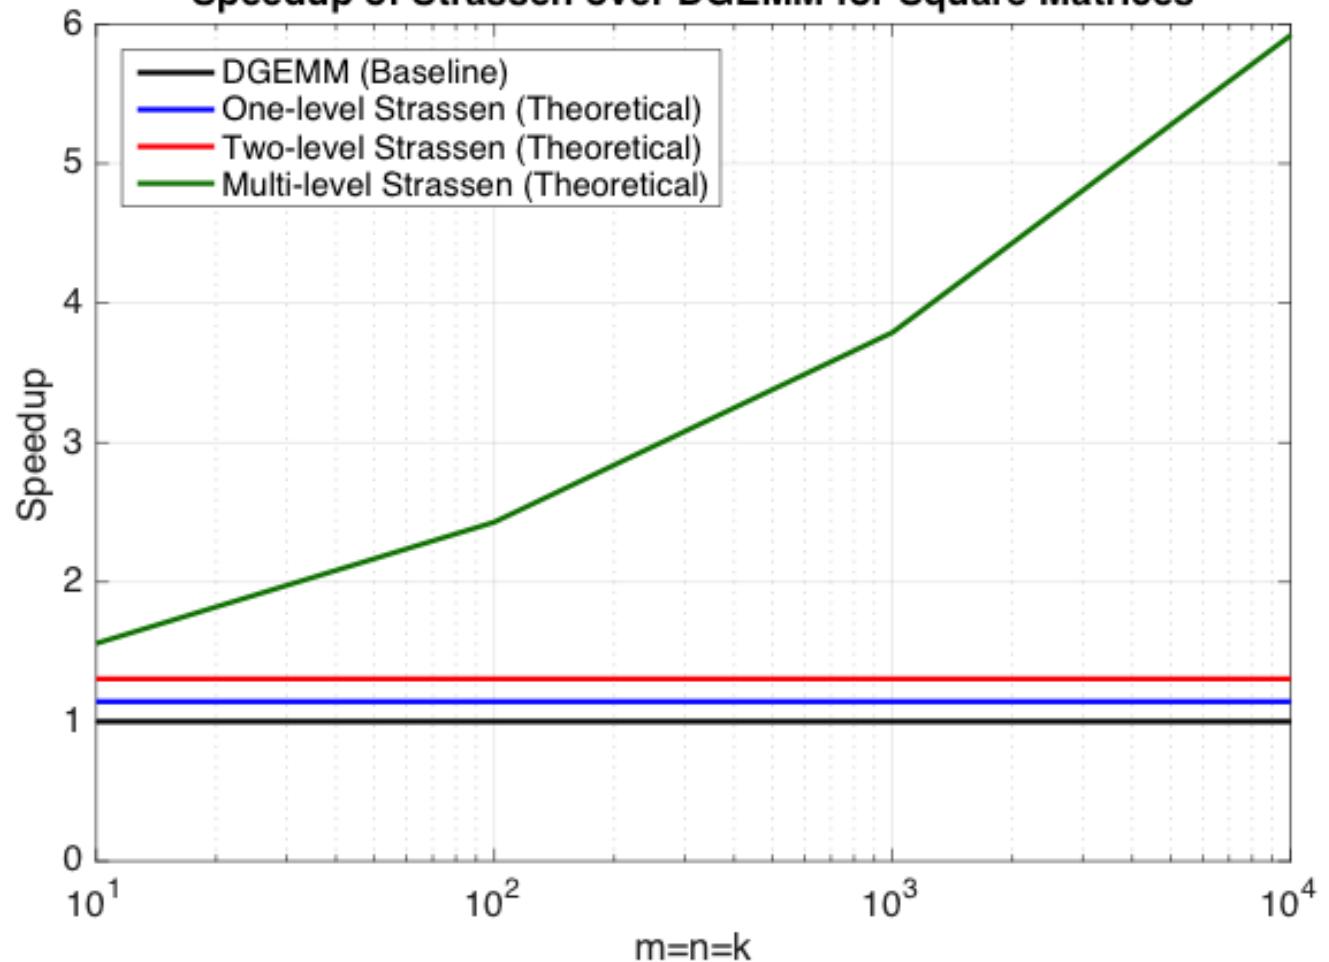
$$C_{01} += M_2 + M_4$$

$$C_{10} += M_1 + M_3$$

$$C_{11} += M_0 - M_1 + M_2 + M_5$$

- One-level Strassen ($1+14.3\%$ speedup)
 - 8 multiplications → 7 multiplications ;
- Two-level Strassen ($1+30.6\%$ speedup)
 - 64 multiplications → 49 multiplications;
- d -level Strassen ($n^3/n^{2.803}$ speedup)
 - 8^d multiplications → 7^d multiplications;

Speedup of Strassen over DGEMM for Square Matrices



Strassen's Algorithm (In practice)

$$M_0 := (A_{00} + A_{11})(B_{00} + B_{11});$$

$$M_1 := (A_{10} + A_{11})B_{00};$$

$$M_2 := A_{00}(B_{01} - B_{11});$$

$$M_3 := A_{11}(B_{10} - B_{00});$$

$$M_4 := (A_{00} + A_{01})B_{11};$$

$$M_5 := (A_{10} - A_{00})(B_{00} + B_{01});$$

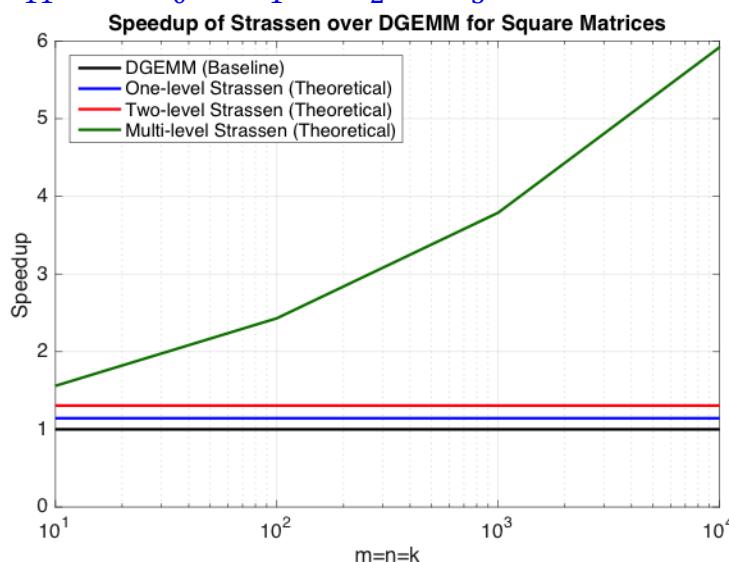
$$M_6 := (A_{01} - A_{11})(B_{10} + B_{11});$$

$$C_{00} += M_0 + M_3 - M_4 + M_6$$

$$C_{01} += M_2 + M_4$$

$$C_{10} += M_1 + M_3$$

$$C_{11} += M_0 - M_1 + M_2 + M_5$$



Strassen's Algorithm (In practice)

$$M_0 := \overbrace{(A_{00} + A_{11})(B_{00} + B_{11})}^{T_0};$$

$$M_1 := \overbrace{(A_{10} + A_{11})B_{00}}^{T_2};$$

$$M_2 := A_{00} \overbrace{(B_{01} - B_{11})}^{T_3};$$

$$M_3 := A_{11} \overbrace{(B_{10} - B_{00})}^{T_4};$$

$$M_4 := \overbrace{(A_{00} + A_{01})B_{11}}^{T_5};$$

$$M_5 := \overbrace{(A_{10} - A_{00})(B_{00} + B_{01})}^{T_6};$$

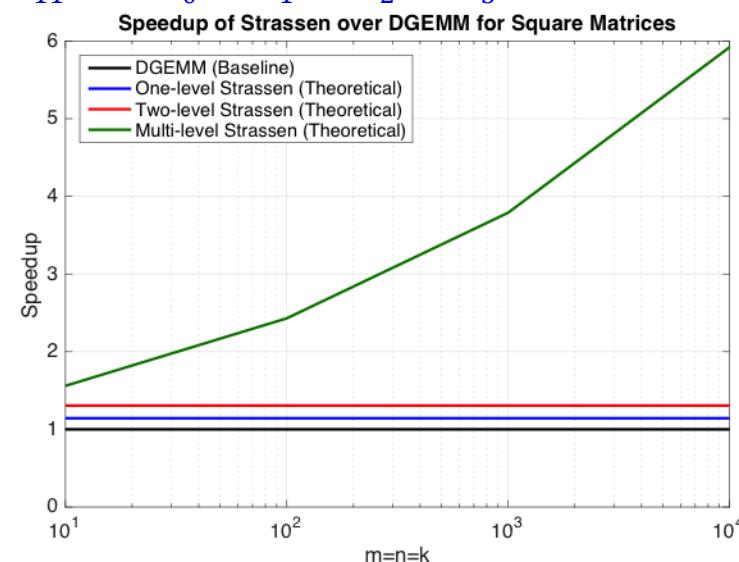
$$M_6 := \overbrace{(A_{01} - A_{11})(B_{10} + B_{11})}^{T_8};$$

$$C_{00} += M_0 + M_3 - M_4 + M_6$$

$$C_{01} += M_2 + M_4$$

$$C_{10} += M_1 + M_3$$

$$C_{11} += M_0 - M_1 + M_2 + M_5$$



- One-level Strassen (1+14.3% speedup)
 - 7 multiplications + 22 additions;
- Two-level Strassen (1+30.6% speedup)
 - 49 multiplications + 344 additions;

Strassen's Algorithm (In practice)

$$M_0 := \overbrace{(A_{00} + A_{11})(B_{00} + B_{11})}^{T_0 + T_1};$$

$$M_1 := \overbrace{(A_{10} + A_{11})B_{00}}^{T_2};$$

$$M_2 := A_{00}\overbrace{(B_{01} - B_{11})}^{T_3};$$

$$M_3 := A_{11}\overbrace{(B_{10} - B_{00})}^{T_4};$$

$$M_4 := \overbrace{(A_{00} + A_{01})B_{11}}^{T_5};$$

$$M_5 := \overbrace{(A_{10} - A_{00})(B_{00} + B_{01})}^{T_6 + T_7};$$

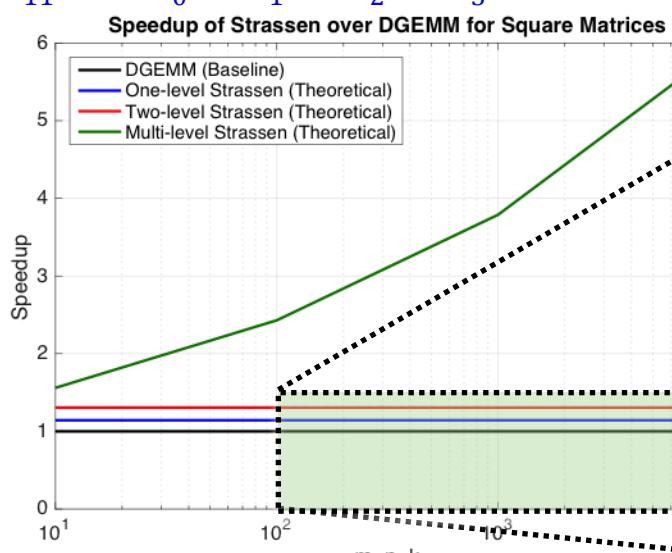
$$M_6 := \overbrace{(A_{01} - A_{11})(B_{10} + B_{11})}^{T_8 + T_9};$$

$$C_{00} += M_0 + M_3 - M_4 + M_6$$

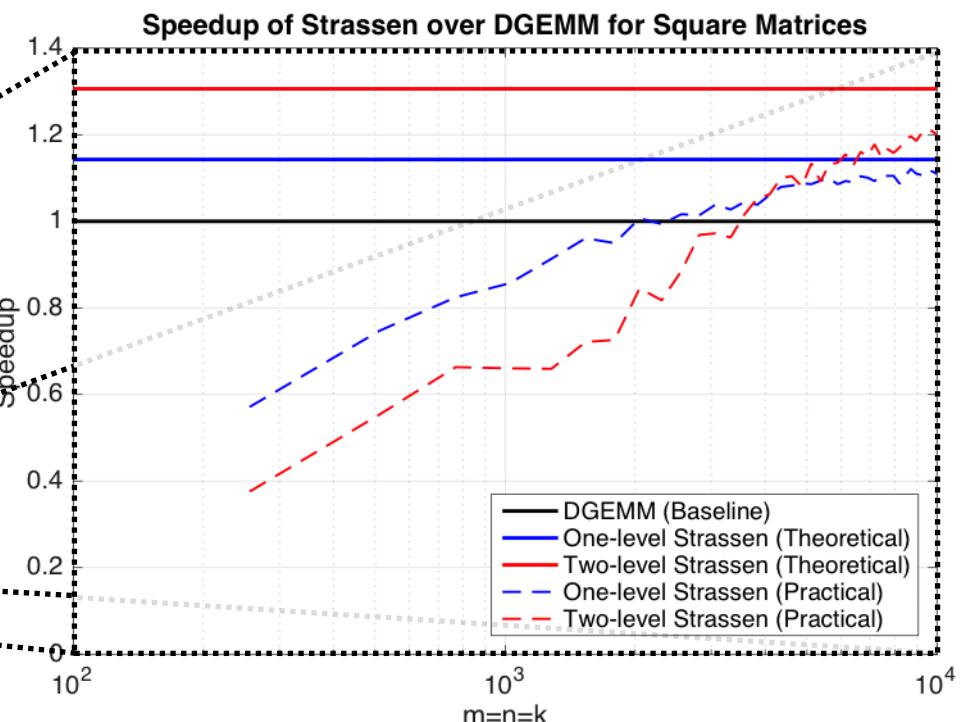
$$C_{01} += M_2 + M_4$$

$$C_{10} += M_1 + M_3$$

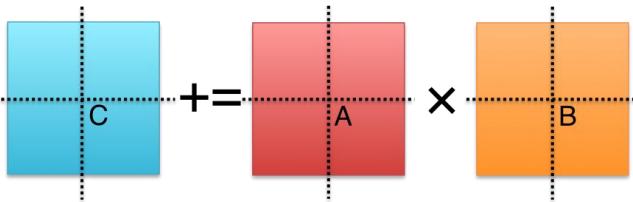
$$C_{11} += M_0 - M_1 + M_2 + M_5$$



- One-level Strassen (1+14.3% speedup)
 - 7 multiplications + 22 additions;
- Two-level Strassen (1+30.6% speedup)
 - 49 multiplications + 344 additions;
- d -level Strassen ($n^3/n^{2.803}$ speedup)
 - Numerical unstable; Not achievable



To achieve practical high performance of Strassen's algorithm.....



**Conventional
Implementations**

**Our
Implementations**

Matrix Size

Must be large



Matrix Shape

Must be square

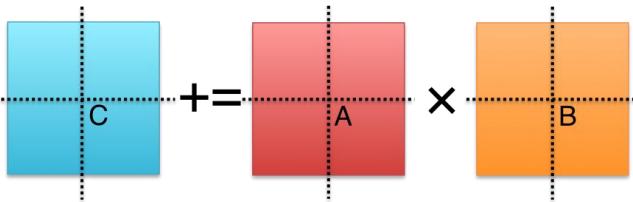


**No Additional
Workspace**



Parallelism

To achieve practical high performance of Strassen's algorithm.....



Conventional
Implementations

Our
Implementations

Matrix Size

Must be large



Matrix Shape

Must be square



No Additional
Workspace

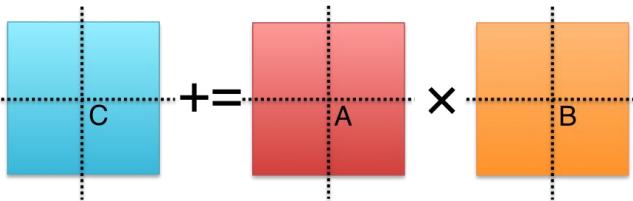


Parallelism

Usually task parallelism



To achieve practical high performance of Strassen's algorithm.....



Conventional
Implementations

Our
Implementations

Matrix Size

Must be large



$$c \quad += \quad A \times B$$



Matrix Shape

Must be square



$$C \quad += \quad A \times B$$



No Additional
Workspace



Parallelism

Usually task parallelism



Can be data parallelism



Outline

- Review of State-of-the-art GEMM in BLIS
- Strassen's Algorithm Reloaded
- Theoretical Model and Practical Performance
- Extension to Other BLAS-3 Operation
- Extension to Other Fast Matrix Multiplication
- Conclusion

Level-3 BLAS Matrix-Matrix Multiplication (GEMM)

- (General) matrix-matrix multiplication (GEMM) is supported in the level-3 BLAS* interface as

```
dgemm( transa, transb, m, n, k,  
       alpha, A, lda, B, ldb,  
       beta, C, ldc )
```

- Ignoring `transa` and `transb`, GEMM computes

$$C := \alpha AB + \beta C;$$

- We consider the simplified version of GEMM

$$C := \alpha AB + C$$

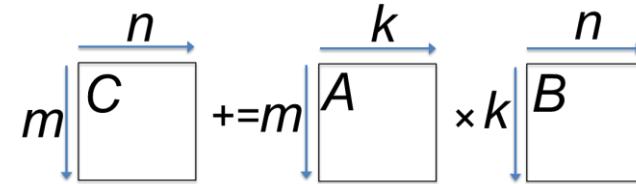
State-of-the-art GEMM in BLIS

- BLAS-like Library Instantiation Software (**BLIS**) is a portable framework for instantiating BLAS-like dense linear algebra libraries.
 - Field Van Zee, and Robert van de Geijn. "BLIS: A Framework for Rapidly Instantiating BLAS Functionality." *ACM TOMS* 41.3 (2015): 14.
- BLIS provides a refactoring of **GotoBLAS** algorithm (best-known approach) to implement **GEMM**.
 - Kazushige Goto, and Robert van de Geijn. "High-performance implementation of the level-3 BLAS." *ACM TOMS* 35.1 (2008): 4.
 - Kazushige Goto, and Robert van de Geijn. "Anatomy of high-performance matrix multiplication." *ACM TOMS* 34.3 (2008): 12.
- GEMM implementation in BLIS has 6-layers of loops. The outer 5 loops are written in **C**. The inner-most loop (micro-kernel) is written in **assembly** for high performance.
 - Partition matrices into smaller blocks to fit into the different memory hierarchy.
 - The order of these loops is designed to utilize the cache reuse rate.

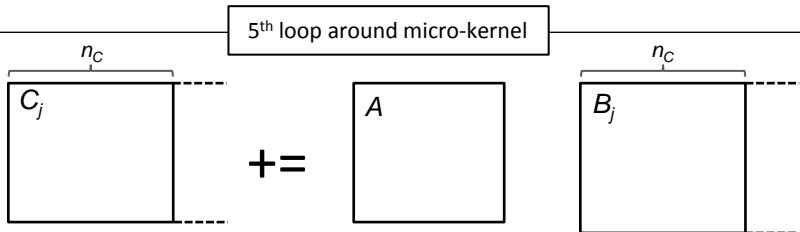
State-of-the-art GEMM in BLIS

- BLAS-like Library Instantiation Software (**BLIS**) is a portable framework for instantiating BLAS-like dense linear algebra libraries.
 - Field Van Zee, and Robert van de Geijn. "BLIS: A Framework for Rapidly Instantiating BLAS Functionality." *ACM TOMS* 41.3 (2015): 14.
- BLIS provides a refactoring of **GotoBLAS** algorithm (best-known approach) to implement **GEMM**.
 - Kazushige Goto, and Robert van de Geijn. "High-performance implementation of the level-3 BLAS." *ACM TOMS* 35.1 (2008): 4.
 - Kazushige Goto, and Robert van de Geijn. "Anatomy of high-performance matrix multiplication." *ACM TOMS* 34.3 (2008): 12.
- GEMM implementation in BLIS has 6-layers of loops. The outer 5 loops are written in **C**. The inner-most loop (micro-kernel) is written in **assembly** for high performance.
 - Partition matrices into smaller blocks to fit into the different memory hierarchy.
 - The order of these loops is designed to utilize the cache reuse rate.
- BLIS opens the black box of GEMM, leading to many applications built on BLIS.
 - Chenhan D. Yu, Jianyu Huang, Woody Austin, Bo Xiao, and George Biros. "Performance Optimization for the **k-Nearest Neighbors** Kernel on x86 Architectures." In *SC'15*.
 - Jianyu Huang, Tyler Smith, Greg Henry, and Robert van de Geijn. "**Strassen**'s Algorithm Reloaded." To appear in *SC'16*.
 - Devin Matthews. "High-Performance **Tensor Contraction** without BLAS.", arXiv:1607.00291
 - Paul Springer, Paolo Bientinesi. "Design of a High-performance GEMM-like **Tensor-Tensor Multiplication**", arXiv:1607.00145

GotoBLAS algorithm for GEMM in **BLIS**

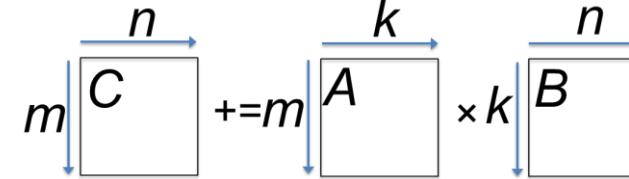


*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



□ main memory

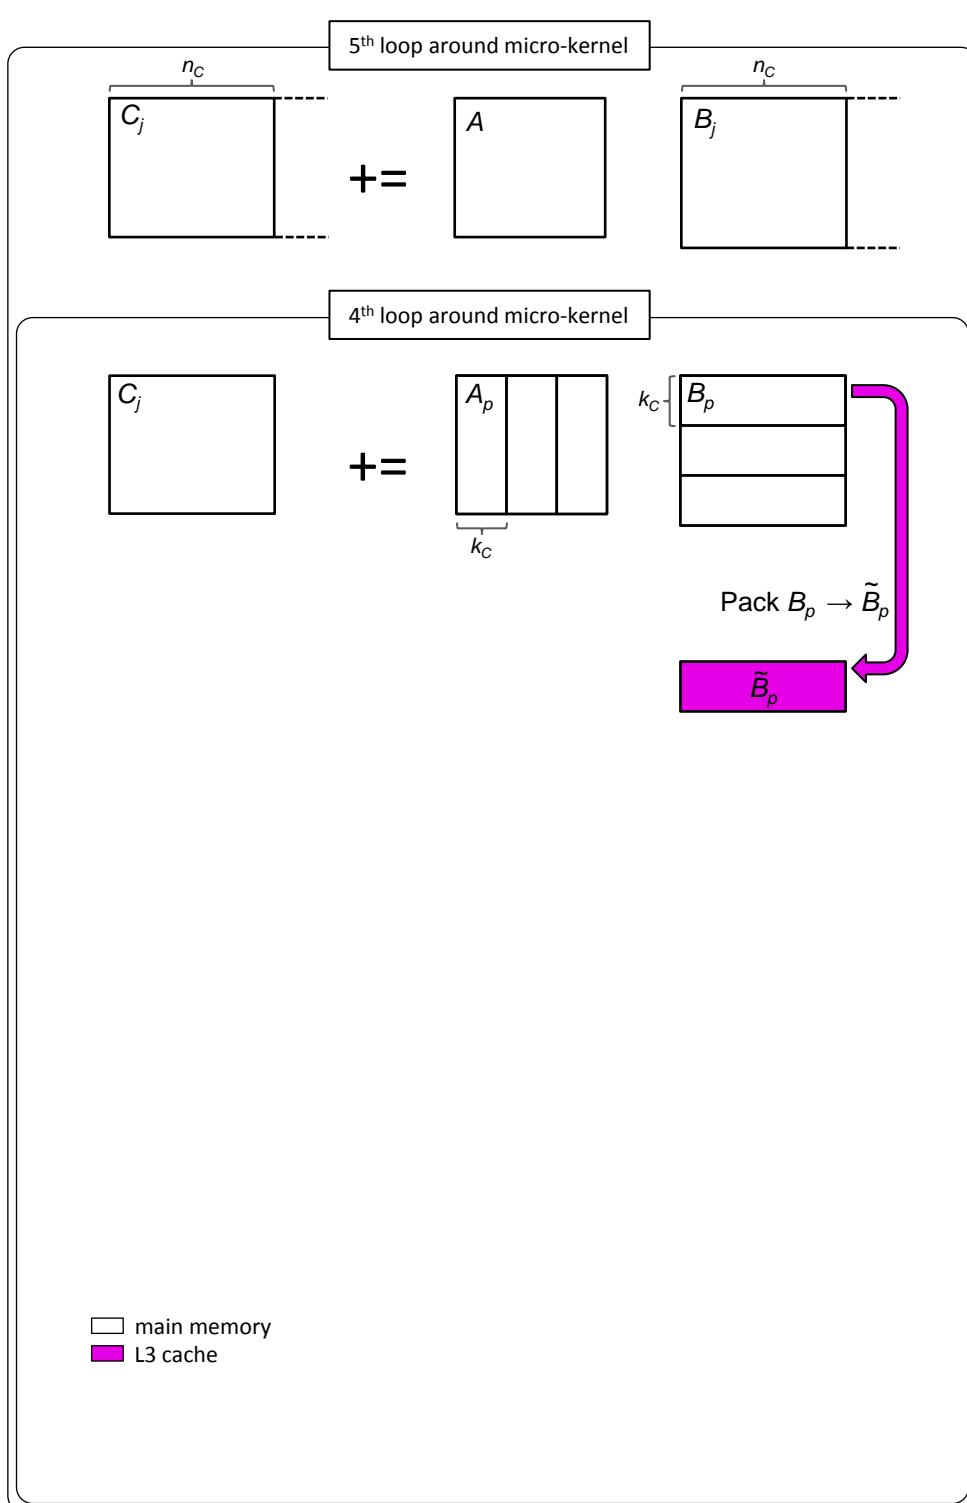
GotoBLAS algorithm for GEMM in **BLIS**



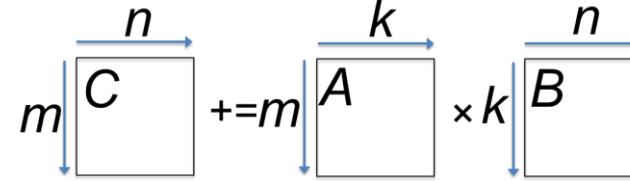
Loop 5 **for** $j_c = 0 : n-1$ **steps of** n_c
 $\mathcal{J}_c = j_c : j_c + n_c - 1$

endfor

*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



GotoBLAS algorithm for GEMM in **BLIS**



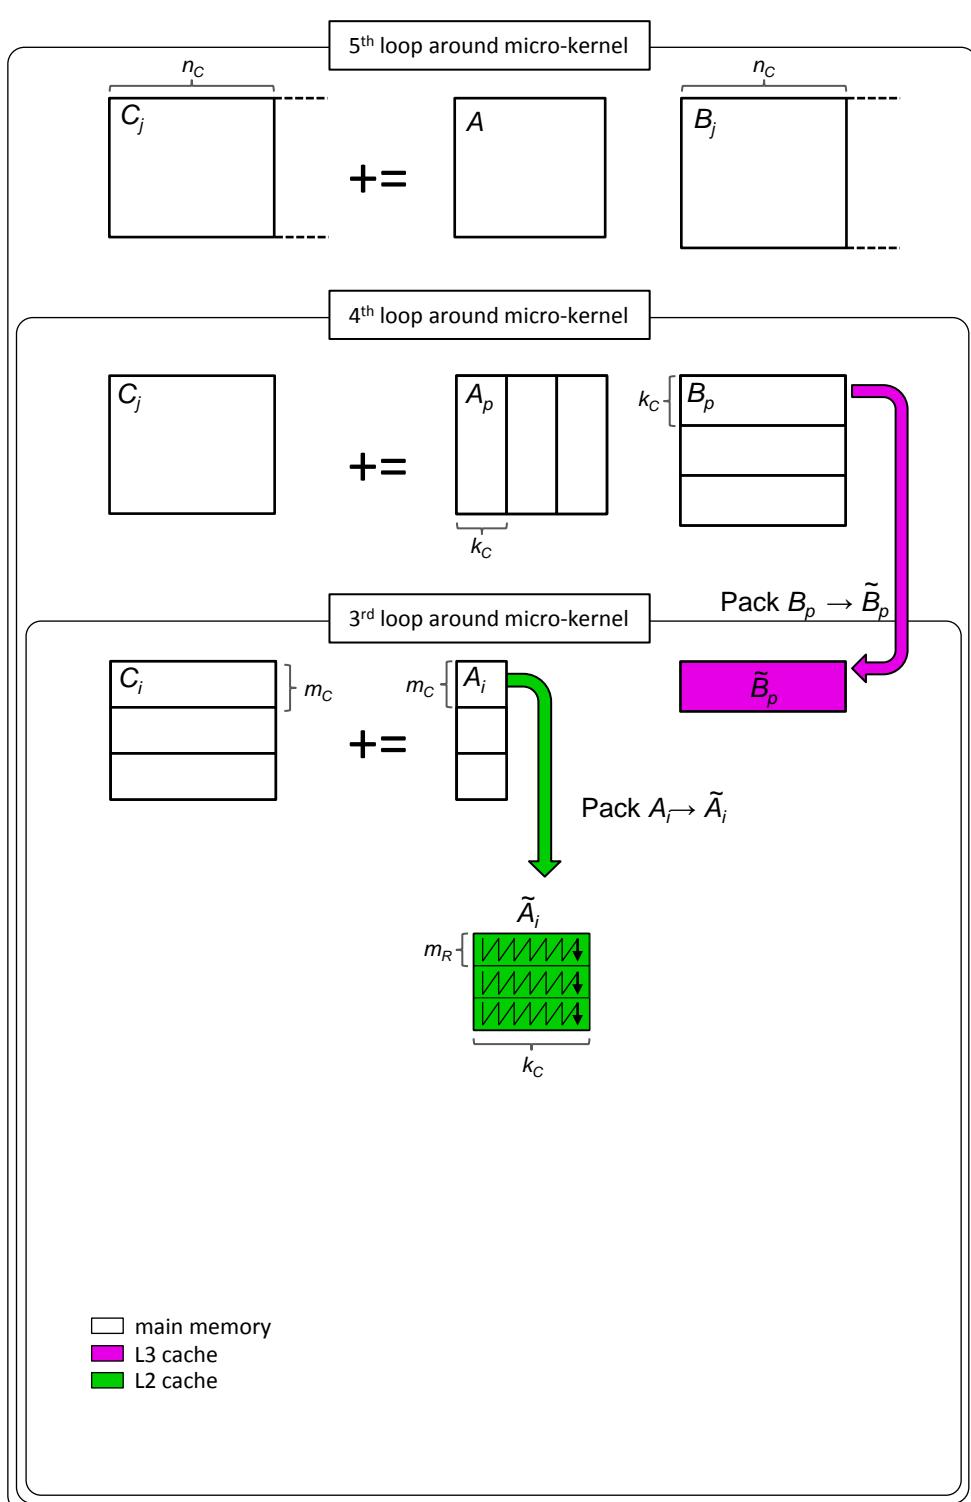
```

Loop 5 for  $j_c = 0 : n-1$  steps of  $n_c$ 
   $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
  for  $p_c = 0 : k-1$  steps of  $k_c$ 
     $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
     $B(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
  
```

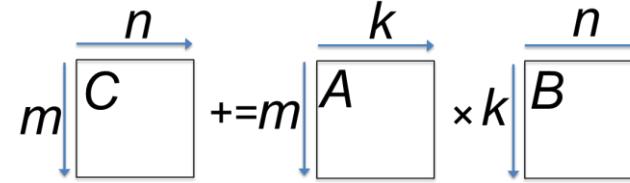
```

endfor
endfor
  
```

*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



GotoBLAS algorithm for GEMM in **BLIS**



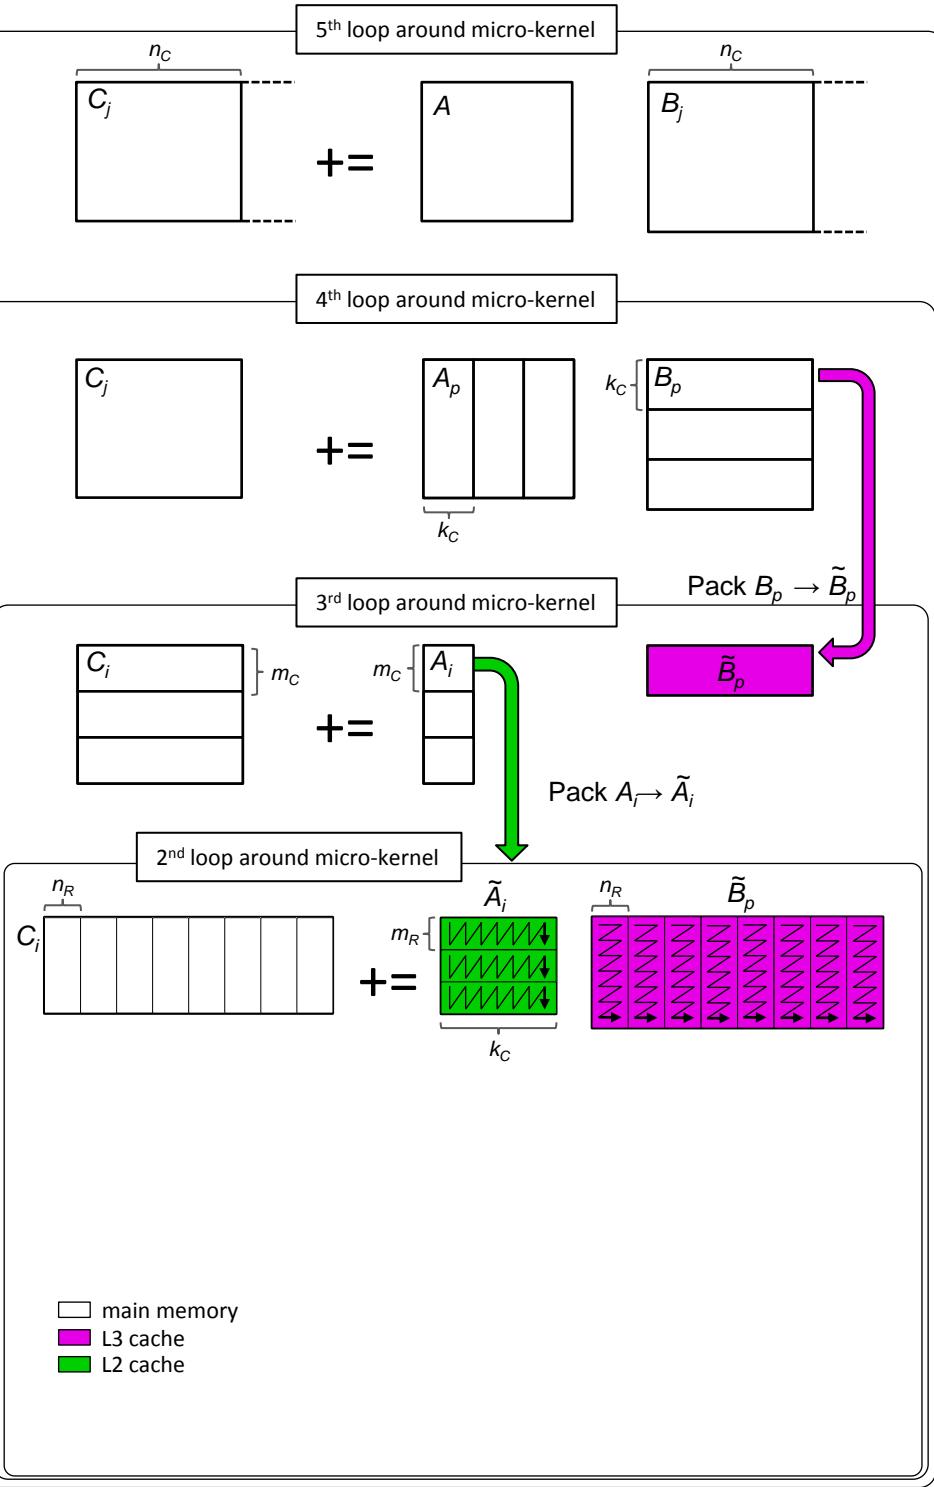
```

Loop 5 for  $j_c = 0 : n-1$  steps of  $n_c$ 
   $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
  for  $p_c = 0 : k-1$  steps of  $k_c$ 
     $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
     $B(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
    for  $i_c = 0 : m-1$  steps of  $m_c$ 
       $\mathcal{I}_c = i_c : i_c + m_c - 1$ 
       $A(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$ 

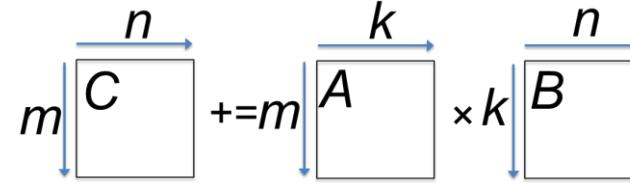
endfor
endfor
endfor

```

*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



GotoBLAS algorithm for GEMM in **BLIS**



```

Loop 5 for  $j_c = 0 : n - 1$  steps of  $n_c$ 
     $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
    for  $p_c = 0 : k - 1$  steps of  $k_c$ 
         $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
         $B(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
        for  $i_c = 0 : m - 1$  steps of  $m_c$ 
             $\mathcal{I}_c = i_c : i_c + m_c - 1$ 
             $A(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$ 
            // macro-kernel
            for  $j_r = 0 : n_c - 1$  steps of  $n_r$ 
                 $\mathcal{J}_r = j_r : j_r + n_r - 1$ 

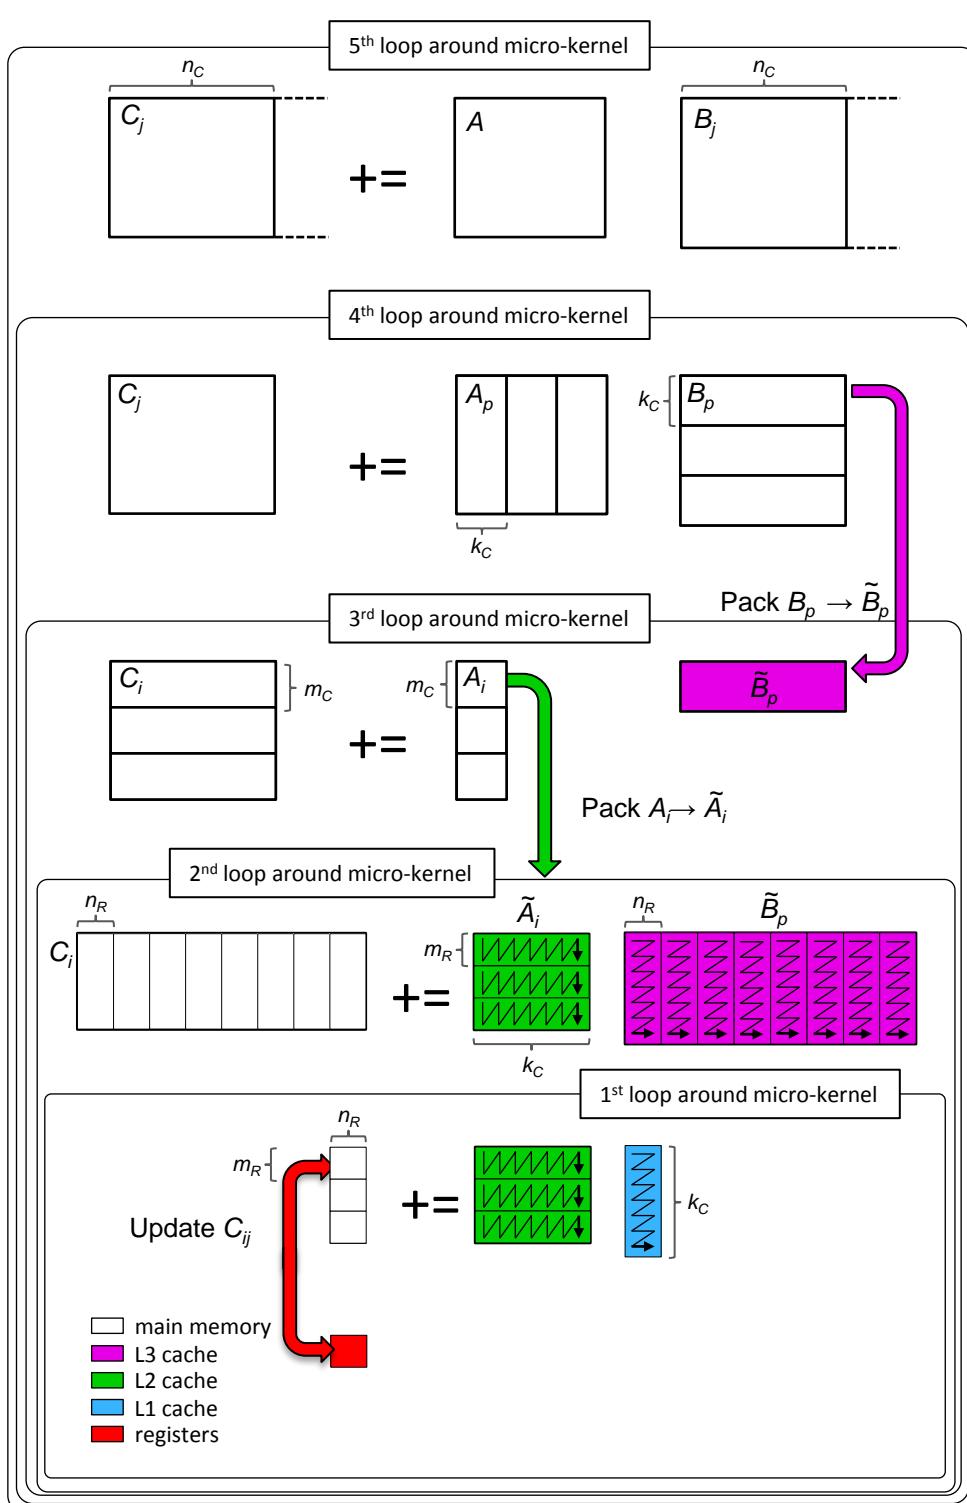
```

```

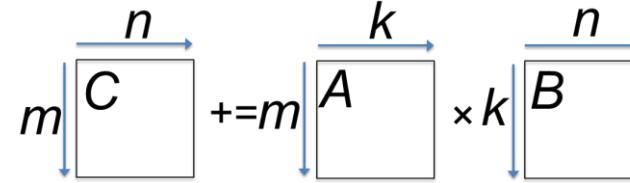
endfor
endfor
endfor
endfor

```

*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



GotoBLAS algorithm for GEMM in **BLIS**

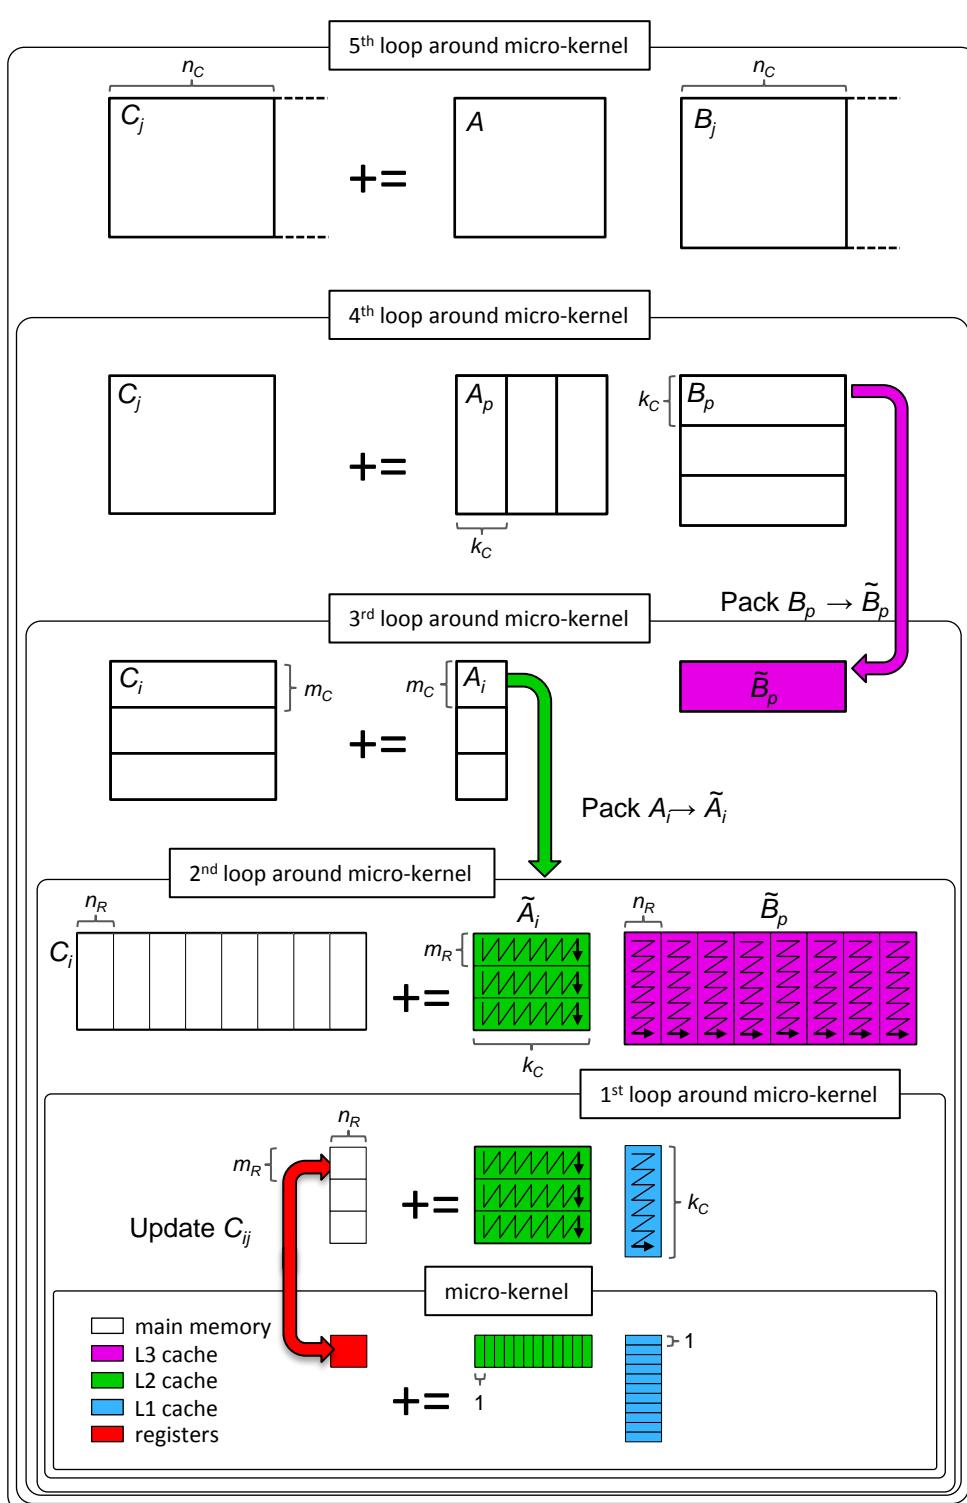


```

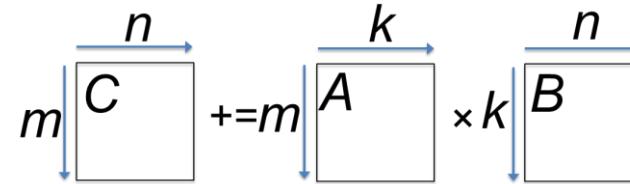
Loop 5 for  $j_c = 0 : n-1$  steps of  $n_c$ 
     $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
    for  $p_c = 0 : k-1$  steps of  $k_c$ 
         $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
         $B(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
        for  $i_c = 0 : m-1$  steps of  $m_c$ 
             $\mathcal{I}_c = i_c : i_c + m_c - 1$ 
             $A(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$ 
            // macro-kernel
            for  $j_r = 0 : n_c-1$  steps of  $n_r$ 
                 $\mathcal{J}_r = j_r : j_r + n_r - 1$ 
                for  $i_r = 0 : m_c-1$  steps of  $m_r$ 
                     $\mathcal{I}_r = i_r : i_r + m_r - 1$ 
                    endfor
                    endfor
                    endfor
                    endfor
                    endfor

```

*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



GotoBLAS algorithm for GEMM in **BLIS**

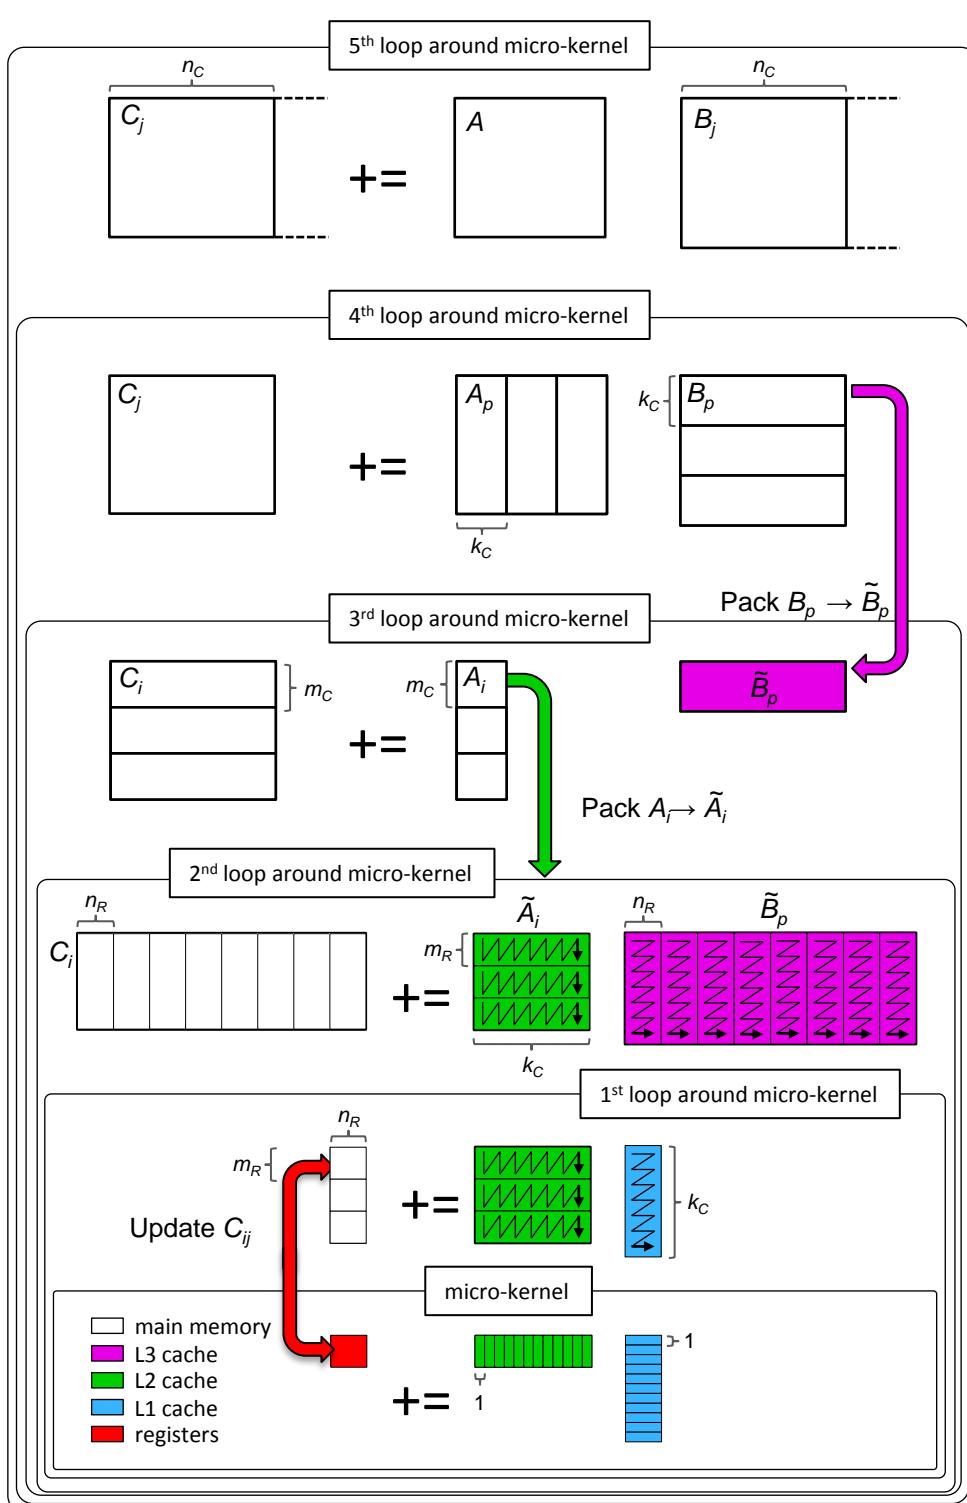


```

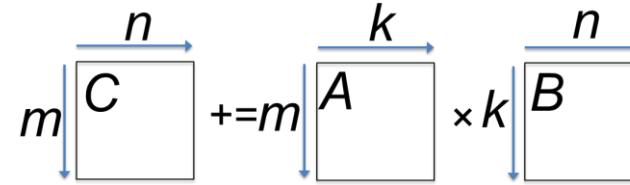
n      k      n
m | C   += m | A   x k | B
      ↓      ↓      ↓
Loop 5  for  $j_c = 0 : n - 1$  steps of  $n_c$ 
         $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
        for  $p_c = 0 : k - 1$  steps of  $k_c$ 
             $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
             $B(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
            for  $i_c = 0 : m - 1$  steps of  $m_c$ 
                 $\mathcal{I}_c = i_c : i_c + m_c - 1$ 
                 $A(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$ 
                // macro-kernel
                for  $j_r = 0 : n_c - 1$  steps of  $n_r$ 
                     $\mathcal{J}_r = j_r : j_r + n_r - 1$ 
                    for  $i_r = 0 : m_c - 1$  steps of  $m_r$ 
                         $\mathcal{I}_r = i_r : i_r + m_r - 1$ 
                        // micro-kernel
                        for  $p_r = 0 : p_c - 1$  steps of 1
                             $C_c(\mathcal{I}_r, \mathcal{J}_r) += \alpha \tilde{A}_i(\mathcal{I}_r, p_r) \tilde{B}_p(p_r, \mathcal{J}_r)$ 
                        endfor
                    endfor
                endfor
            endfor
        endfor
    endfor
endfor

```

*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.



GotoBLAS algorithm for GEMM in **BLIS**



*Field G. Van Zee, and Tyler M. Smith. “Implementing high-performance complex matrix multiplication.” In *ACM Transactions on Mathematical Software (TOMS)*, accepted pending modifications.

Outline

- Review of State-of-the-art GEMM in BLIS
- **Strassen's Algorithm Reloaded**
- Theoretical Model and Practical Performance
- Extension to Other BLAS-3 Operation
- Extension to Other Fast Matrix Multiplication
- Conclusion

One-level Strassen's Algorithm Reloaded

$M_0 := \alpha(A_{00} + A_{11})(B_{00} + B_{11});$
 $M_1 := \alpha(A_{10} + A_{11})B_{00};$
 $M_2 := \alpha A_{00}(B_{01} - B_{11});$
 $M_3 := \alpha A_{11}(B_{10} - B_{00});$
 $M_4 := \alpha(A_{00} + A_{01})B_{11};$
 $M_5 := \alpha(A_{10} - A_{00})(B_{00} + B_{01});$
 $M_6 := \alpha(A_{01} - A_{11})(B_{10} + B_{11});$
 $C_{00} += M_0 + M_3 - M_4 + M_6$
 $C_{01} += M_2 + M_4$
 $C_{10} += M_1 + M_3$
 $C_{11} += M_0 - M_1 + M_2 + M_5$



$M_0 := \alpha(A_{00} + A_{11})(B_{00} + B_{11});$	$C_{00} += M_0; C_{11} += M_0;$
$M_1 := \alpha(A_{10} + A_{11})B_{00};$	$C_{10} += M_1; C_{11} -= M_1;$
$M_2 := \alpha A_{00}(B_{01} - B_{11});$	$C_{01} += M_2; C_{11} += M_2;$
$M_3 := \alpha A_{11}(B_{10} - B_{00});$	$C_{00} += M_3; C_{10} += M_3;$
$M_4 := \alpha(A_{00} + A_{01})B_{11};$	$C_{01} += M_4; C_{00} -= M_4;$
$M_5 := \alpha(A_{10} - A_{00})(B_{00} + B_{01});$	$C_{11} += M_5;$
$M_6 := \alpha(A_{01} - A_{11})(B_{10} + B_{11});$	$C_{00} += M_6;$

$M := \alpha(X + Y)(V + W);$	$C += M;$	$D += M;$
------------------------------	-----------	-----------

General operation for one-level Strassen:

$M := \alpha(X + \delta Y)(V + \varepsilon W);$	$C += \gamma_0 M; D += \gamma_1 M;$
$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}.$	

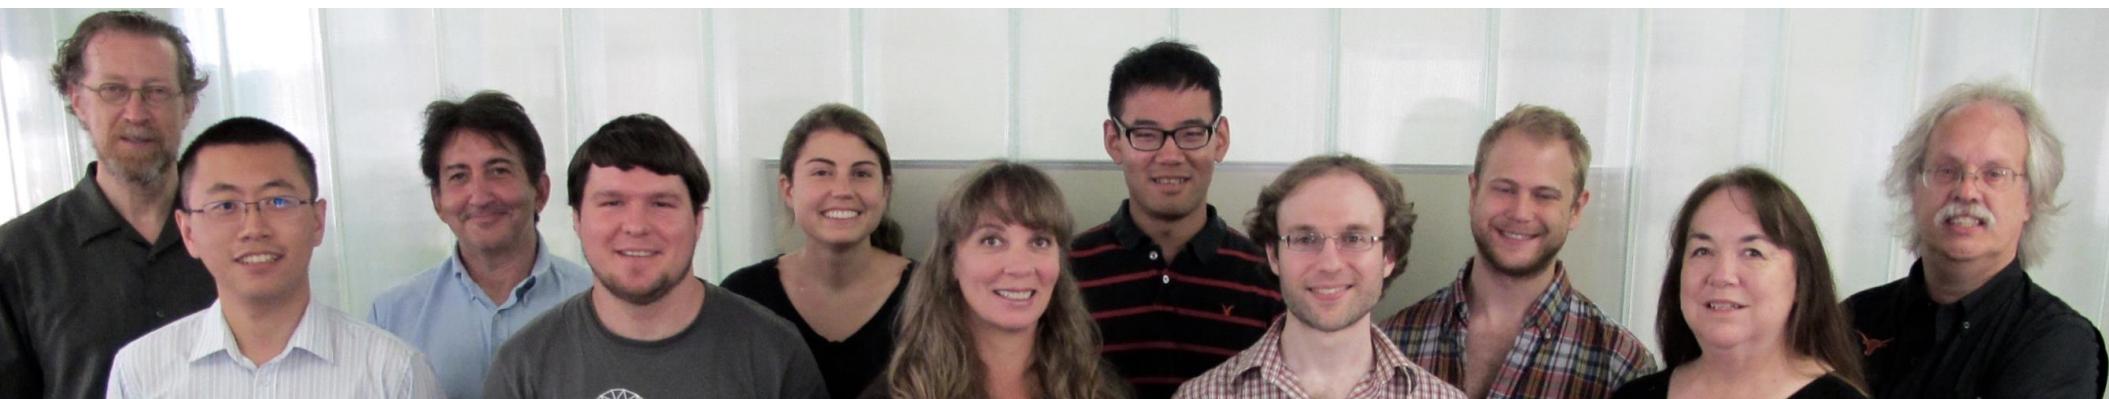
The Science of
High-Performance
Computing Group



High-performance implementation of the general operation?

$$M := \alpha(X + \delta Y)(V + \varepsilon W); \\ \gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}.$$

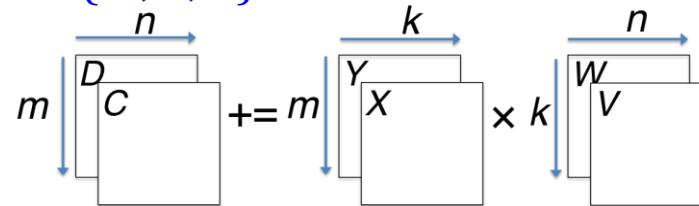
$$C += \gamma_0 M; D += \gamma_1 M;$$



*<http://shpc.ices.utexas.edu/index.html>

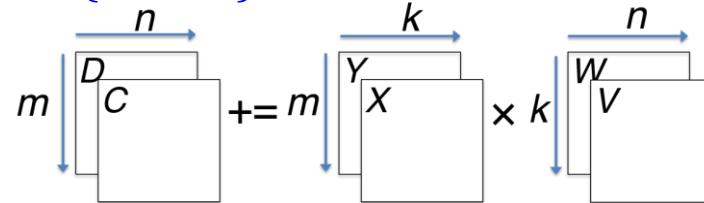
$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M;$$

$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}$.



$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M;$$

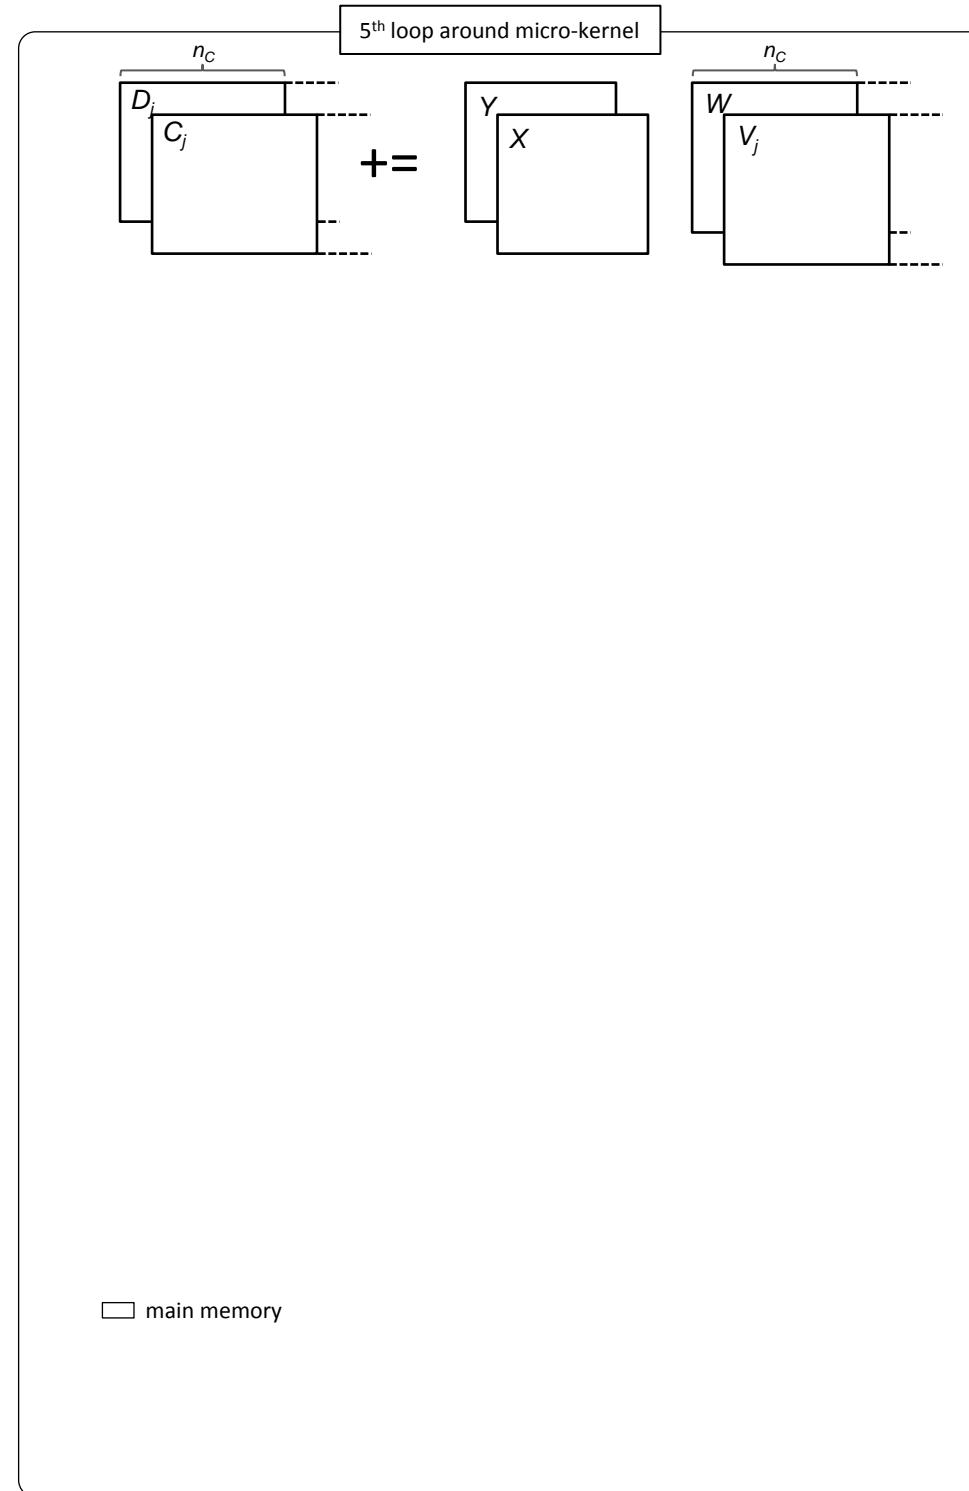
$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}$.



Loop 5 for $j_c = 0 : n - 1$ steps of n_c
 $\mathcal{J}_c = j_c : j_c + n_c - 1$

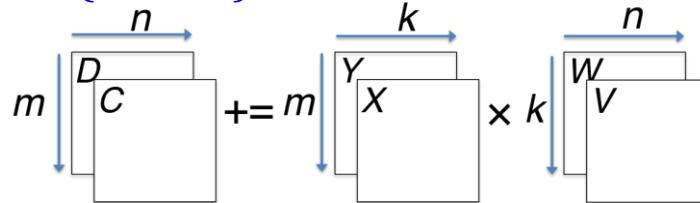
endfor

*Jianyu Huang, Tyler Smith, Greg Henry, and Robert van de Geijn.
“Strassen’s Algorithm Reloaded.” In SC’16.



$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M;$$

$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}$.



Loop 5 **for** $j_c = 0 : n - 1$ **steps of** n_c

$$\mathcal{J}_c = j_c : j_c + n_c - 1$$

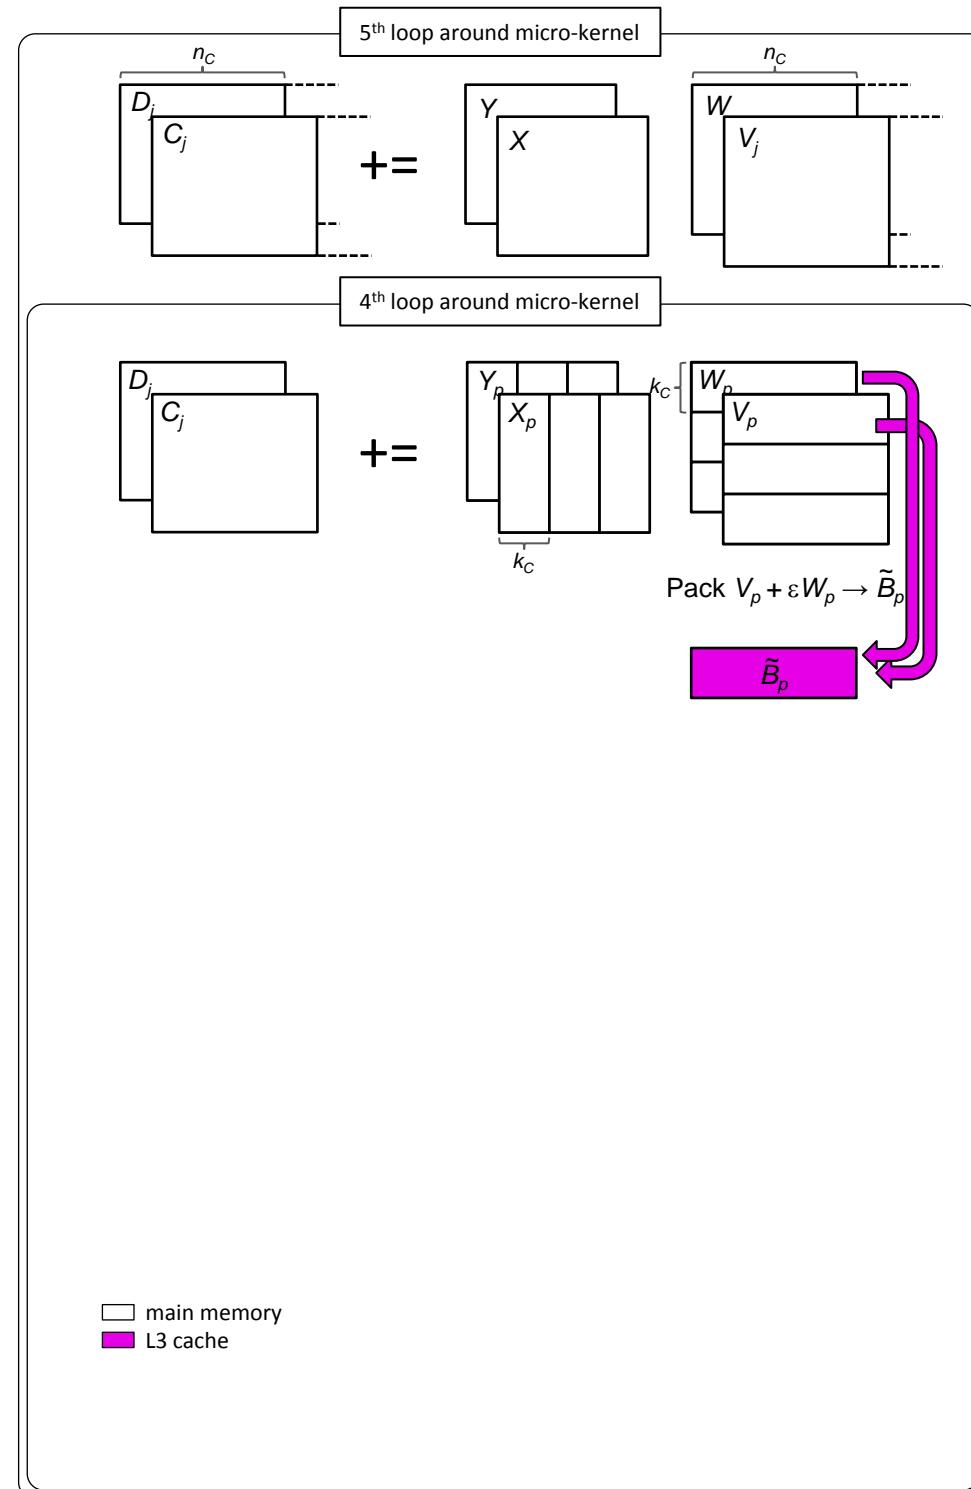
Loop 4 **for** $p_c = 0 : k - 1$ **steps of** k_c

$$\mathcal{P}_c = p_c : p_c + k_c - 1$$

$$V(\mathcal{P}_c, \mathcal{J}_c) + \varepsilon W(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$$

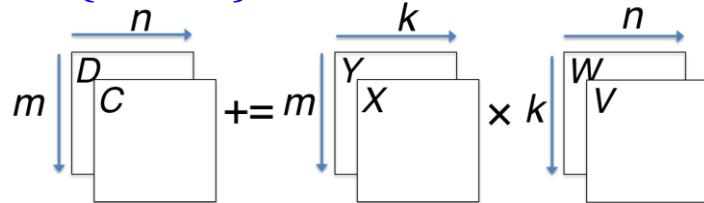
endfor
endfor

*Jianyu Huang, Tyler Smith, Greg Henry, and Robert van de Geijn.
“Strassen’s Algorithm Reloaded.” In SC’16.



$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M;$$

$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}$.



```

Loop 5  for  $j_c = 0 : n - 1$  steps of  $n_c$ 
         $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
Loop 4  for  $p_c = 0 : k - 1$  steps of  $k_c$ 
         $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
         $V(\mathcal{P}_c, \mathcal{J}_c) + \varepsilon W(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
Loop 3  for  $i_c = 0 : m - 1$  steps of  $m_c$ 
         $\mathcal{I}_c = i_c : i_c + m_c - 1$ 
         $X(\mathcal{I}_c, \mathcal{P}_c) + \delta Y(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$ 

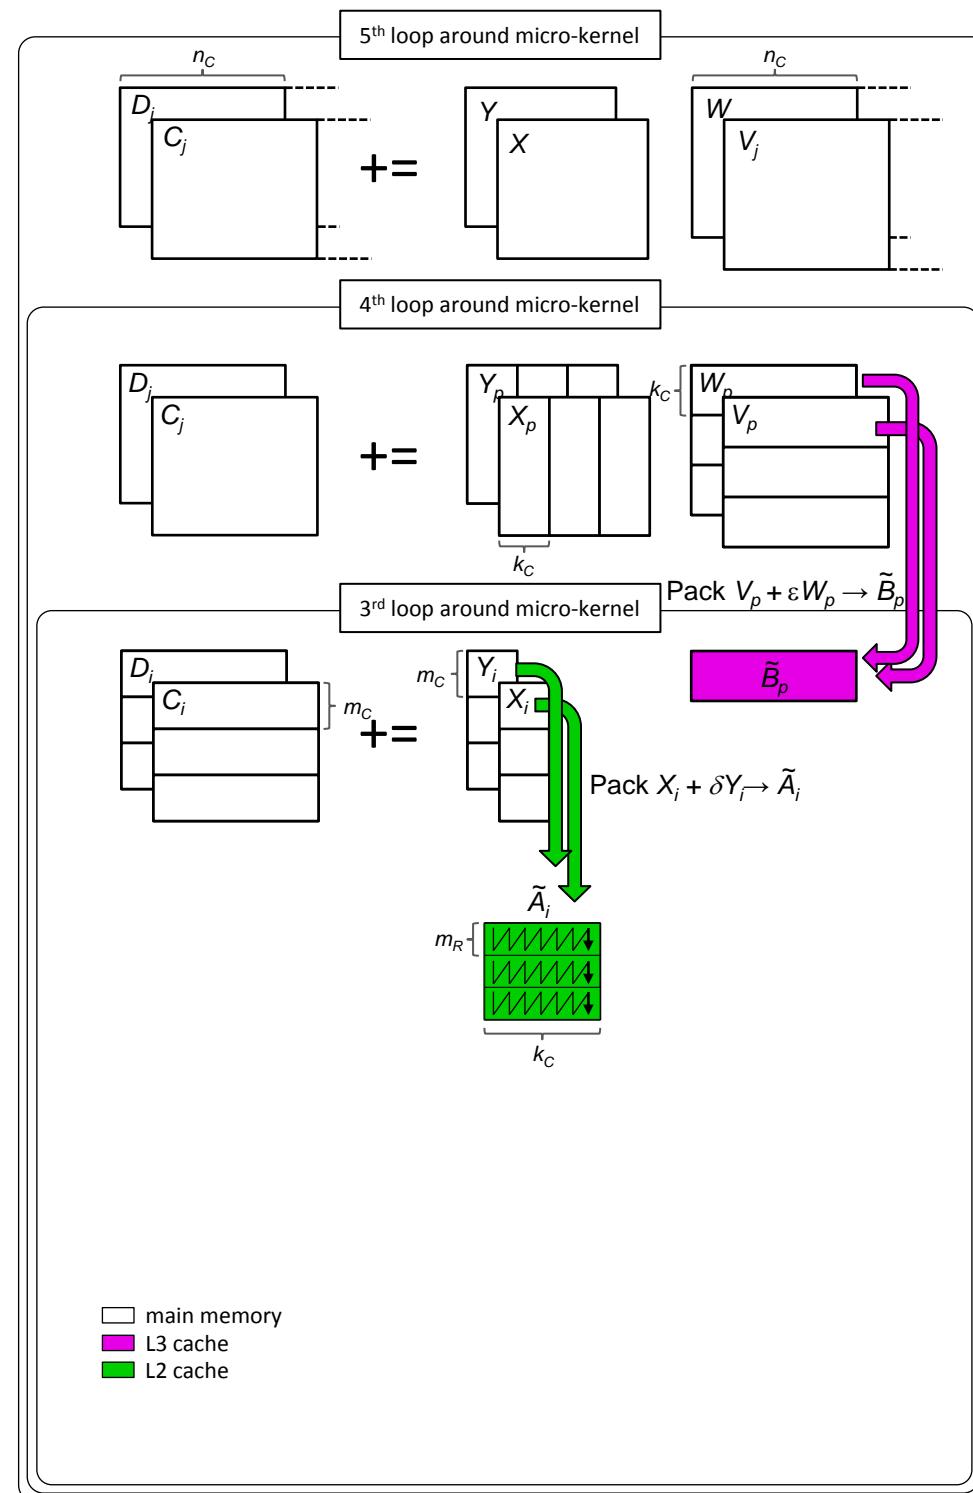
```

```

    endfor
    endfor
    endfor

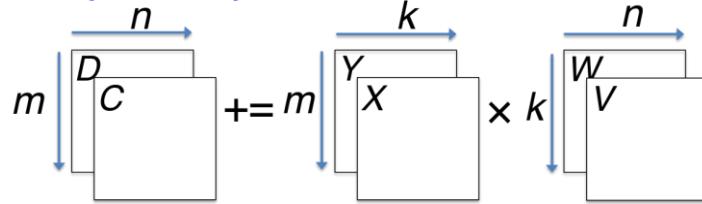
```

*Jianyu Huang, Tyler Smith, Greg Henry, and Robert van de Geijn.
“Strassen’s Algorithm Reloaded.” In SC’16.



$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M;$$

$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}$.



Loop 5 **for** $j_c = 0 : n - 1$ **steps of** n_c

$$\mathcal{J}_c = j_c : j_c + n_c - 1$$

Loop 4 **for** $p_c = 0 : k - 1$ **steps of** k_c

$$\mathcal{P}_c = p_c : p_c + k_c - 1$$

$$V(\mathcal{P}_c, \mathcal{J}_c) + \varepsilon W(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$$

Loop 3 **for** $i_c = 0 : m - 1$ **steps of** m_c

$$\mathcal{I}_c = i_c : i_c + m_c - 1$$

$$X(\mathcal{I}_c, \mathcal{P}_c) + \delta Y(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$$

// macro-kernel

Loop 2 **for** $j_r = 0 : n_c - 1$ **steps of** n_r

$$\mathcal{J}_r = j_r : j_r + n_r - 1$$

Loop 1 **for** $i_r = 0 : m_c - 1$ **steps of** m_r

$$\mathcal{I}_r = i_r : i_r + m_r - 1$$

// micro-kernel

Loop 0 **for** $p_r = 0 : p_c - 1$ **steps of** 1

$$M_r(\mathcal{I}_r, \mathcal{J}_r) += \tilde{A}_i(\mathcal{I}_r, p_r) \tilde{B}_p(p_r, \mathcal{J}_r)$$

endfor

$$C(\mathcal{I}_r + i_c, \mathcal{J}_r + j_c) += \alpha \gamma_0 M_r(\mathcal{I}_r, \mathcal{J}_r)$$

$$D(\mathcal{I}_r + i_c, \mathcal{J}_r + j_c) += \alpha \gamma_1 M_r(\mathcal{I}_r, \mathcal{J}_r)$$

endfor

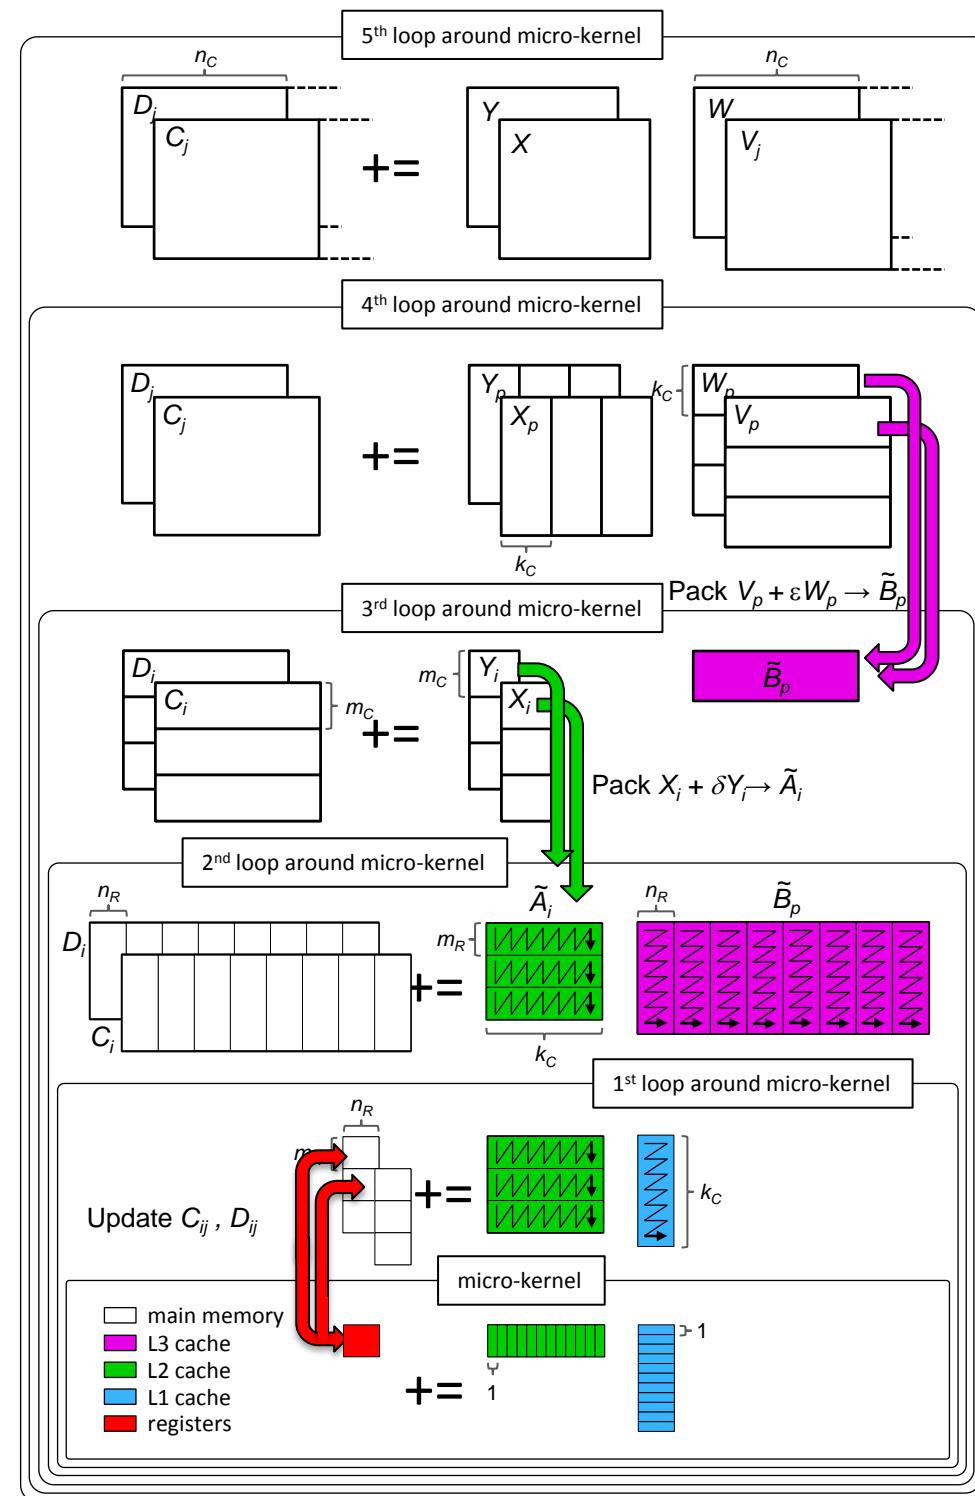
endfor

endfor

endfor

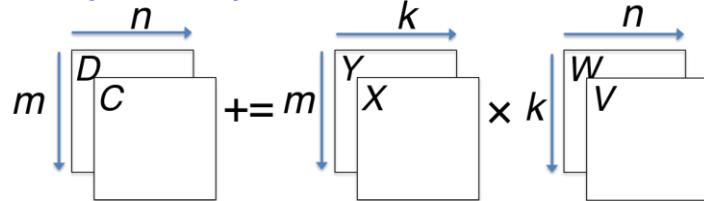
*Jianyu Huang, Tyler Smith, Greg Henry, and Robert van de Geijn.

“Strassen’s Algorithm Reloaded.” In SC’16.



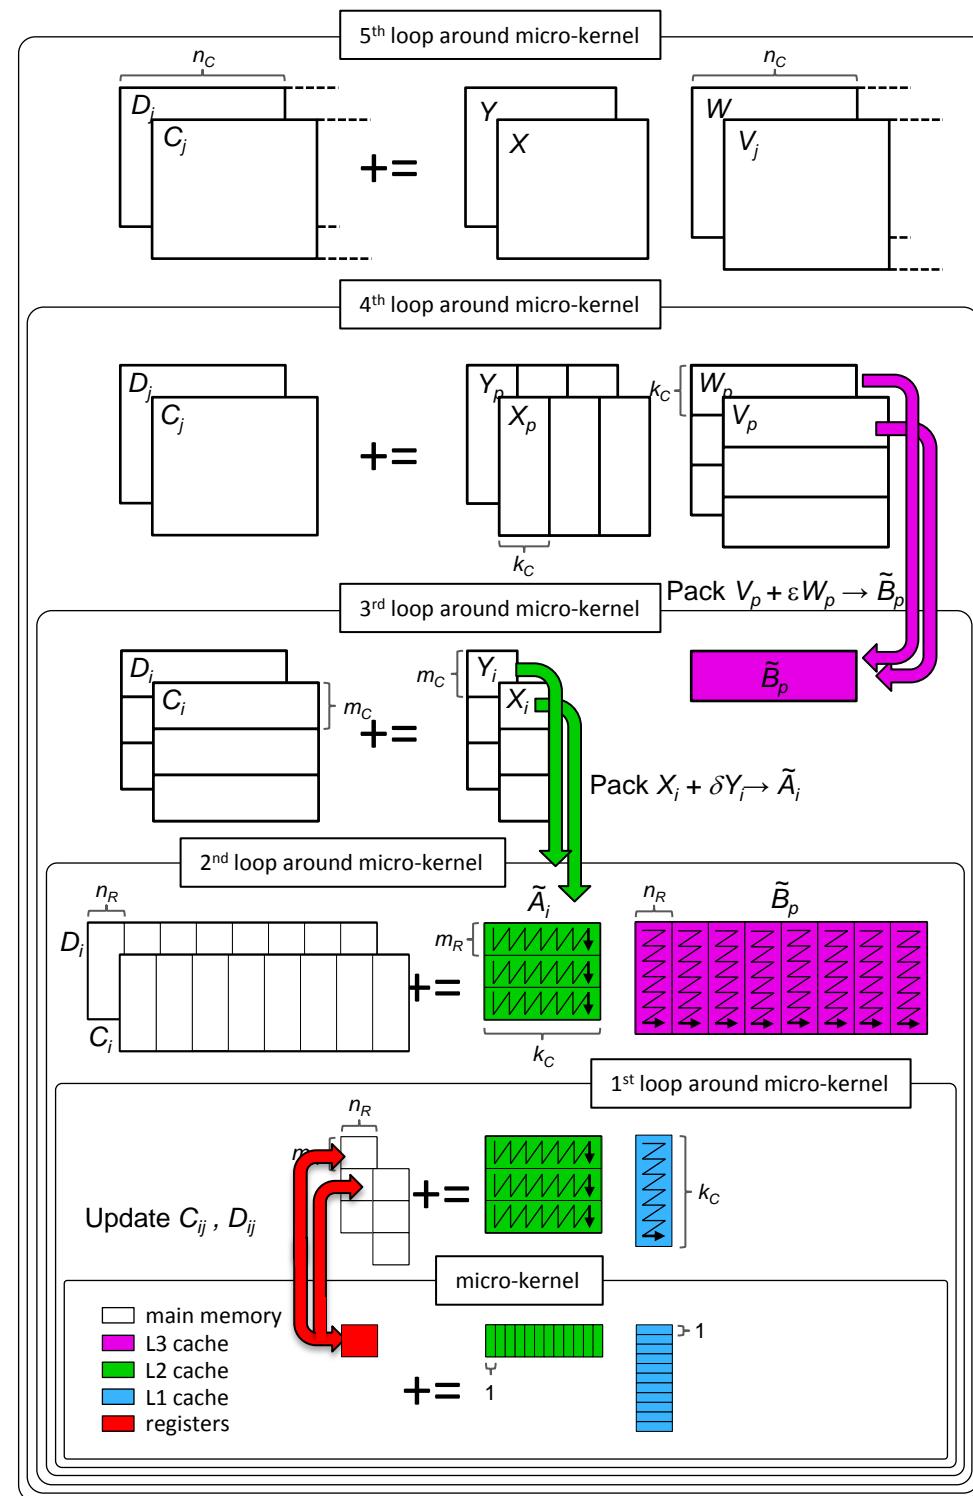
$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M;$$

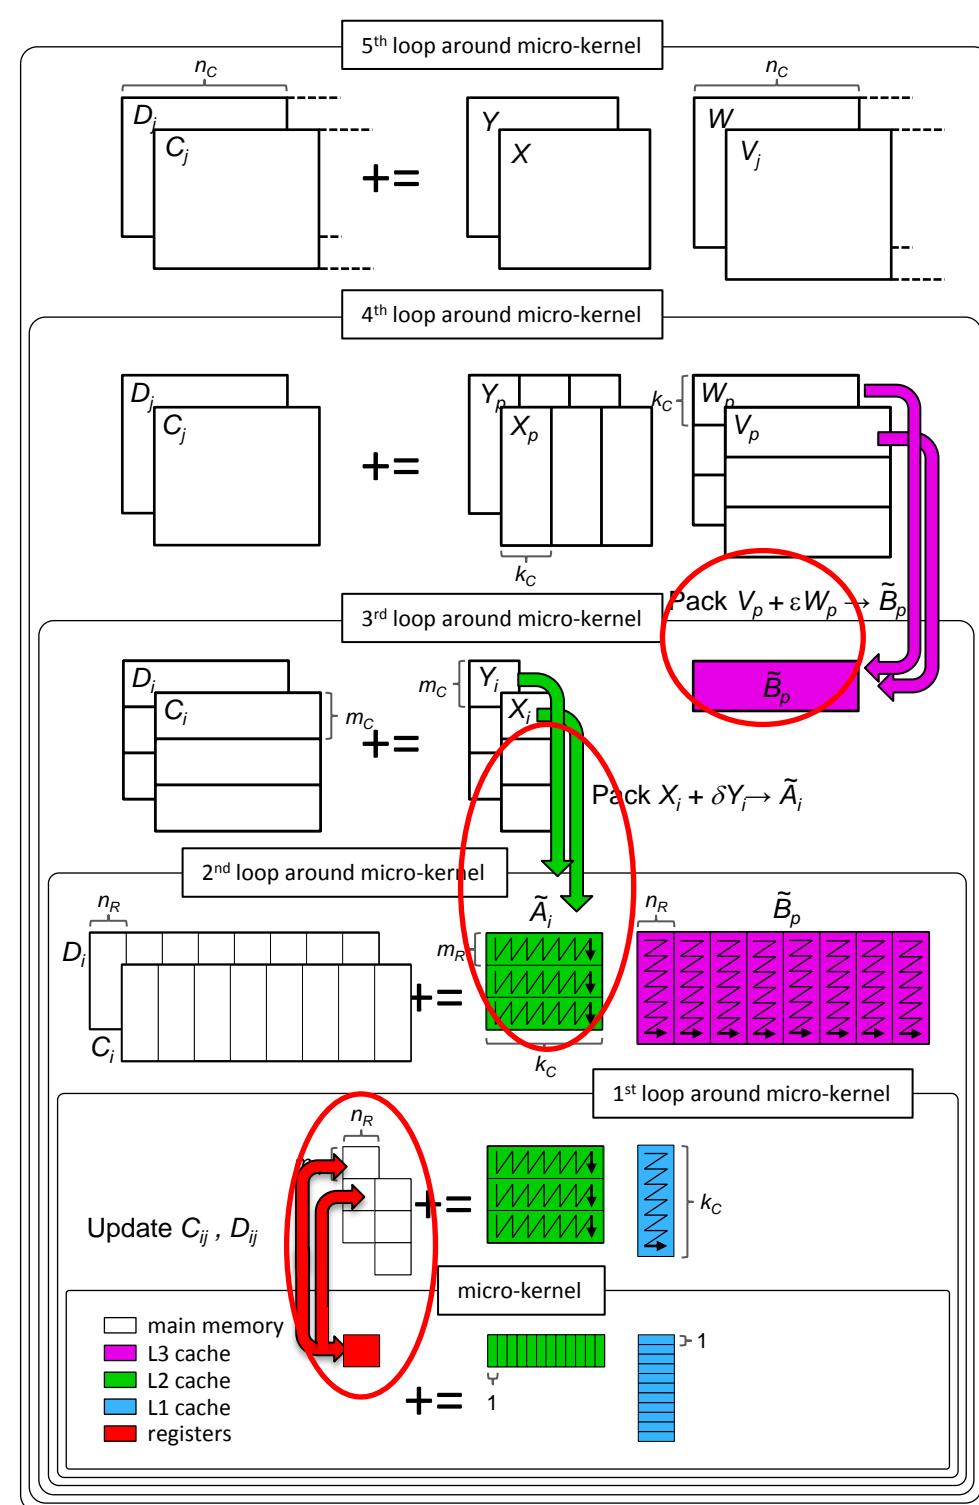
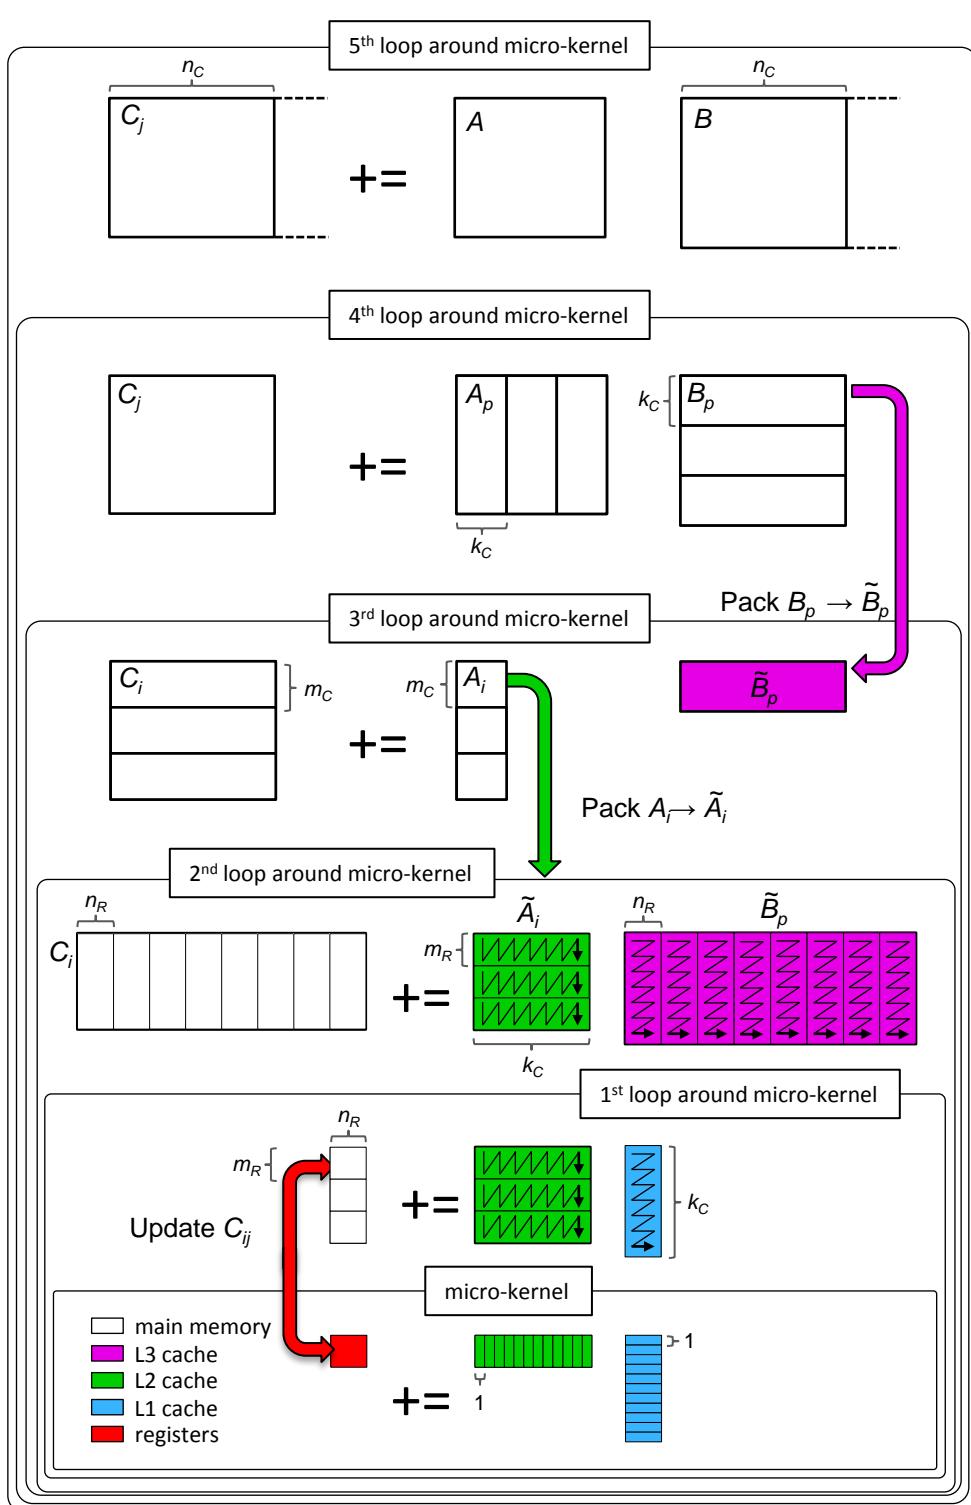
$\gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}$.



```

Loop 5  for  $j_c = 0 : n - 1$  steps of  $n_c$ 
         $\mathcal{J}_c = j_c : j_c + n_c - 1$ 
Loop 4  for  $p_c = 0 : k - 1$  steps of  $k_c$ 
         $\mathcal{P}_c = p_c : p_c + k_c - 1$ 
         $V(\mathcal{P}_c, \mathcal{J}_c) + \epsilon W(\mathcal{P}_c, \mathcal{J}_c) \rightarrow \tilde{B}_p$ 
Loop 3  for  $i_c = 0 : m - 1$  steps of  $m_c$ 
         $\mathcal{I}_c = i_c : i_c + m_c - 1$ 
         $X(\mathcal{I}_c, \mathcal{P}_c) + \delta Y(\mathcal{I}_c, \mathcal{P}_c) \rightarrow \tilde{A}_i$ 
        // macro-kernel
Loop 2  for  $j_r = 0 : n_c - 1$  steps of  $n_r$ 
         $\mathcal{J}_r = j_r : j_r + n_r - 1$ 
Loop 1  for  $i_r = 0 : m_c - 1$  steps of  $m_r$ 
         $\mathcal{I}_r = i_r : i_r + m_r - 1$ 
        //micro-kernel
Loop 0  for  $p_r = 0 : p_c - 1$  steps of 1
         $M_r(\mathcal{I}_r, \mathcal{J}_r) += \tilde{A}_i(\mathcal{I}_r, p_r) \tilde{B}_p(p_r, \mathcal{J}_r)$ 
        endfor
         $C(\mathcal{I}_r + i_c, \mathcal{J}_r + j_c) += \alpha \gamma_0 M_r(\mathcal{I}_r, \mathcal{J}_r)$ 
         $D(\mathcal{I}_r + i_c, \mathcal{J}_r + j_c) += \alpha \gamma_1 M_r(\mathcal{I}_r, \mathcal{J}_r)$ 
        endfor
        endfor
        endfor
        endfor
    
```





Two-level Strassen's Algorithm Reloaded

Assume m , n , and k are all multiples of 4. Letting

$$C = \left(\begin{array}{cc|cc} C_{0,0} & C_{0,1} & C_{0,2} & C_{0,3} \\ C_{1,0} & C_{1,1} & C_{1,2} & C_{1,3} \\ \hline C_{2,0} & C_{2,1} & C_{2,2} & C_{2,3} \\ C_{3,0} & C_{3,1} & C_{3,2} & C_{3,3} \end{array} \right), A = \left(\begin{array}{cc|cc} A_{0,0} & A_{0,1} & A_{0,2} & A_{0,3} \\ A_{1,0} & A_{1,1} & A_{1,2} & A_{1,3} \\ \hline A_{2,0} & A_{2,1} & A_{2,2} & A_{2,3} \\ A_{3,0} & A_{3,1} & A_{3,2} & A_{3,3} \end{array} \right), B = \left(\begin{array}{cc|cc} B_{0,0} & B_{0,1} & B_{0,2} & B_{0,3} \\ B_{1,0} & B_{1,1} & B_{1,2} & B_{1,3} \\ \hline B_{2,0} & B_{2,1} & B_{2,2} & B_{2,3} \\ B_{3,0} & B_{3,1} & B_{3,2} & B_{3,3} \end{array} \right),$$

where $C_{i,j}$ is $\frac{m}{4} \times \frac{n}{4}$, $A_{i,p}$ is $\frac{m}{4} \times \frac{k}{4}$, and $B_{p,j}$ is $\frac{k}{4} \times \frac{n}{4}$.

Two-level Strassen's Algorithm Reloaded (Continue)

$M_0 := \alpha(A_{0,0} + A_{2,2} + A_{1,1} + A_{3,3})(B_{0,0} + B_{2,2} + B_{1,1} + B_{3,3});$	$C_{0,0} += M_0;$	$C_{1,1} += M_0;$	$C_{2,2} += M_0;$	$C_{3,3} += M_0;$
$M_1 := \alpha(A_{1,0} + A_{3,2} + A_{1,1} + A_{3,3})(B_{0,0} + B_{2,2});$	$C_{1,0} += M_1;$	$C_{1,1} -= M_1;$	$C_{3,2} += M_1;$	$C_{3,3} -= M_1;$
$M_2 := \alpha(A_{0,0} + A_{2,2})(B_{0,1} + B_{2,3} + B_{1,1} + B_{3,3});$	$C_{0,1} += M_2;$	$C_{1,1} += M_2;$	$C_{2,3} += M_2;$	$C_{3,3} += M_2;$
$M_3 := \alpha(A_{1,1} + A_{3,3})(B_{1,0} + B_{3,2} + B_{0,0} + B_{2,2});$	$C_{0,0} += M_3;$	$C_{1,0} += M_3;$	$C_{2,2} += M_3;$	$C_{3,2} += M_3;$
$M_4 := \alpha(A_{0,0} + A_{2,2} + A_{0,1} + A_{2,3})(B_{1,1} + B_{3,3});$	$C_{0,0} -= M_4;$	$C_{0,1} += M_4;$	$C_{2,2} -= M_4;$	$C_{2,3} += M_4;$
$M_5 := \alpha(A_{1,0} + A_{3,2} + A_{0,0} + A_{2,2})(B_{0,0} + B_{2,2} + B_{0,1} + B_{2,3});$	$C_{1,1} += M_5;$	$C_{3,3} += M_5;$		
$M_6 := \alpha(A_{0,1} + A_{2,3} + A_{1,1} + A_{3,3})(B_{1,0} + B_{3,2} + B_{1,1} + B_{3,3});$	$C_{0,0} += M_6;$	$C_{2,2} += M_6;$		
$M_7 := \alpha(A_{2,0} + A_{2,2} + A_{3,1} + A_{3,3})(B_{0,0} + B_{1,1});$	$C_{2,0} += M_7;$	$C_{3,1} += M_7;$	$C_{2,2} -= M_7;$	$C_{3,3} -= M_7;$
$M_8 := \alpha(A_{3,0} + A_{3,2} + A_{3,1} + A_{3,3})(B_{0,0});$	$C_{3,0} += M_8;$	$C_{3,1} -= M_8;$	$C_{3,2} -= M_8;$	$C_{3,3} += M_8;$
$M_9 := \alpha(A_{2,0} + A_{2,2})(B_{0,1} + B_{1,1});$	$C_{2,1} += M_9;$	$C_{3,1} += M_9;$	$C_{2,3} -= M_9;$	$C_{3,3} -= M_9;$
$M_{10} := \alpha(A_{3,1} + A_{3,3})(B_{1,0} + B_{0,0});$	$C_{2,0} += M_{10};$	$C_{3,0} += M_{10};$	$C_{2,2} -= M_{10};$	$C_{3,2} -= M_{10};$
.
$M_{40} := \alpha(A_{3,0} + A_{1,0} + A_{2,0} + A_{0,0})(B_{0,0} + B_{0,2} + B_{0,1} + B_{0,3});$	$C_{3,3} += M_{40};$			
$M_{41} := \alpha(A_{2,1} + A_{0,1} + A_{3,1} + A_{1,1})(B_{1,0} + B_{1,2} + B_{1,1} + B_{1,3});$	$C_{2,2} += M_{41};$			
$M_{42} := \alpha(A_{0,2} + A_{2,2} + A_{1,3} + A_{3,3})(B_{2,0} + B_{2,2} + B_{3,1} + B_{3,3});$	$C_{0,0} += M_{42};$	$C_{1,1} += M_{42};$		
$M_{43} := \alpha(A_{1,2} + A_{3,2} + A_{1,3} + A_{3,3})(B_{2,0} + B_{2,2});$	$C_{1,0} += M_{43};$	$C_{1,1} -= M_{43};$		
$M_{44} := \alpha(A_{0,2} + A_{2,2})(B_{2,1} + B_{2,3} + B_{3,1} + B_{3,3});$	$C_{0,1} += M_{44};$	$C_{1,1} += M_{44};$		
$M_{45} := \alpha(A_{1,3} + A_{3,3})(B_{3,0} + B_{3,2} + B_{2,0} + B_{2,2});$	$C_{0,0} += M_{45};$	$C_{1,0} += M_{45};$		
$M_{46} := \alpha(A_{0,2} + A_{2,2} + A_{0,3} + A_{2,3})(B_{3,1} + B_{3,3});$	$C_{0,0} -= M_{46};$	$C_{0,1} += M_{46};$		
$M_{47} := \alpha(A_{1,2} + A_{3,2} + A_{0,2} + A_{2,2})(B_{2,0} + B_{2,2} + B_{2,1} + B_{2,3});$	$C_{1,1} += M_{47};$			
$M_{48} := \alpha(A_{0,3} + A_{2,3} + A_{1,3} + A_{3,3})(B_{3,0} + B_{3,2} + B_{3,1} + B_{3,3});$	$C_{0,0} += M_{48};$			
$M := \alpha(X_0 + X_1 + X_2 + X_3)(V + V_1 + V_2 + V_3);$	$C_0 += M;$	$C_1 += M;$	$C_2 += M;$	$C_3 += M;$

General operation for two-level Strassen:

$M := \alpha(X_0 + \delta_1 X_1 + \delta_2 X_2 + \delta_3 X_3)(V + \varepsilon_1 V_1 + \varepsilon_2 V_2 + \varepsilon_3 V_3);$	$C_0 += \gamma_0 M;$	$C_1 += \gamma_1 M;$	$C_2 += \gamma_2 M;$	$C_3 += \gamma_3 M;$
$\gamma_p \delta_p \varepsilon_i \in \{-1, 0, 1\}.$				

Additional Levels of Strassen Reloaded

- The general operation of one-level Strassen:

$$M := \alpha(X + \delta Y)(V + \varepsilon W); \quad C += \gamma_0 M; D += \gamma_1 M; \\ \gamma_0, \gamma_1, \delta, \varepsilon \in \{-1, 0, 1\}.$$

- The general operation of two-level Strassen:

$$M := \alpha(X_0 + \delta_1 X_1 + \delta_2 X_2 + \delta_3 X_3)(V + \varepsilon_1 V_1 + \varepsilon_2 V_2 + \varepsilon_3 V_3); \\ C_0 += \gamma_0 M; C_1 += \gamma_1 M; C_2 += \gamma_2 M; C_3 += \gamma_3 M; \\ \gamma_i, \delta_i, \varepsilon_i \in \{-1, 0, 1\}.$$

- The general operation needed to integrate k levels of Strassen is given by

$$M := \alpha \left(\sum_{s=0}^{l_X-1} \delta_s X_s \right) \left(\sum_{t=0}^{l_V-1} \varepsilon_t V_t \right); \\ C_r += \gamma_r M \text{ for } r = 0, \dots, l_C - 1; \\ \delta_i, \varepsilon_i, \gamma_i \in \{-1, 0, 1\}.$$

Building blocks

$$M := \alpha \left(\sum_{s=0}^{l_X-1} \delta_s X_s \right) \left(\sum_{t=0}^{l_V-1} \epsilon_t V_t \right);$$
$$C_r += \gamma_r M \text{ for } r = 0, \dots, l_C - 1;$$
$$\delta_i, \epsilon_i, \gamma_i \in \{-1, 0, 1\}.$$

BLIS framework

- A routine for packing B_p into \tilde{B}_p ➔
 - C/Intel intrinsics
- A routine for packing A_i into \tilde{A}_i ➔
 - C/Intel intrinsics
- A micro-kernel for updating an $m_R \times n_R$ submatrix of C . ➔
 - SIMD assembly (AVX, AVX512, etc)

Adapted to general operation

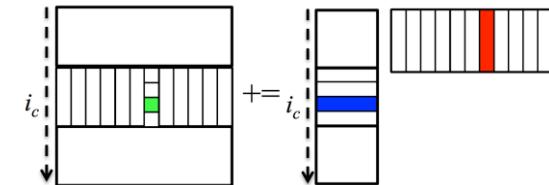
- Integrate the addition of multiple matrices V_t into \tilde{B}_p
- Integrate the addition of multiple matrices X_s into \tilde{A}_i
- Integrate the update of multiple submatrices of C .

Variations on a theme

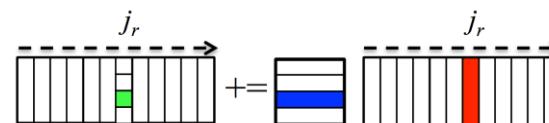
- Naïve Strassen
 - A traditional implementation with temporary buffers.
- AB Strassen
 - Integrate the addition of matrices into \tilde{A}_i and \tilde{B}_p .
- ABC Strassen
 - Integrate the addition of matrices into \tilde{A}_i and \tilde{B}_p .
 - Integrate the update of multiple submatrices of C in the micro-kernel.

Parallelization

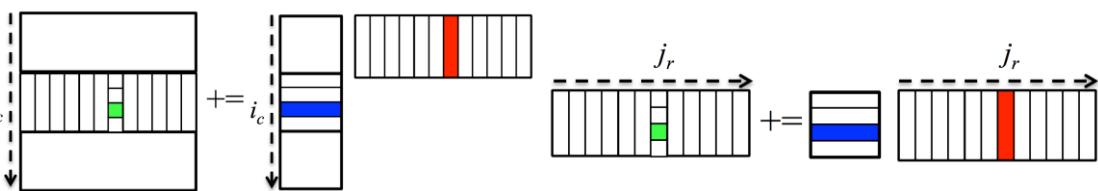
- 3rd loop (along m_C direction)



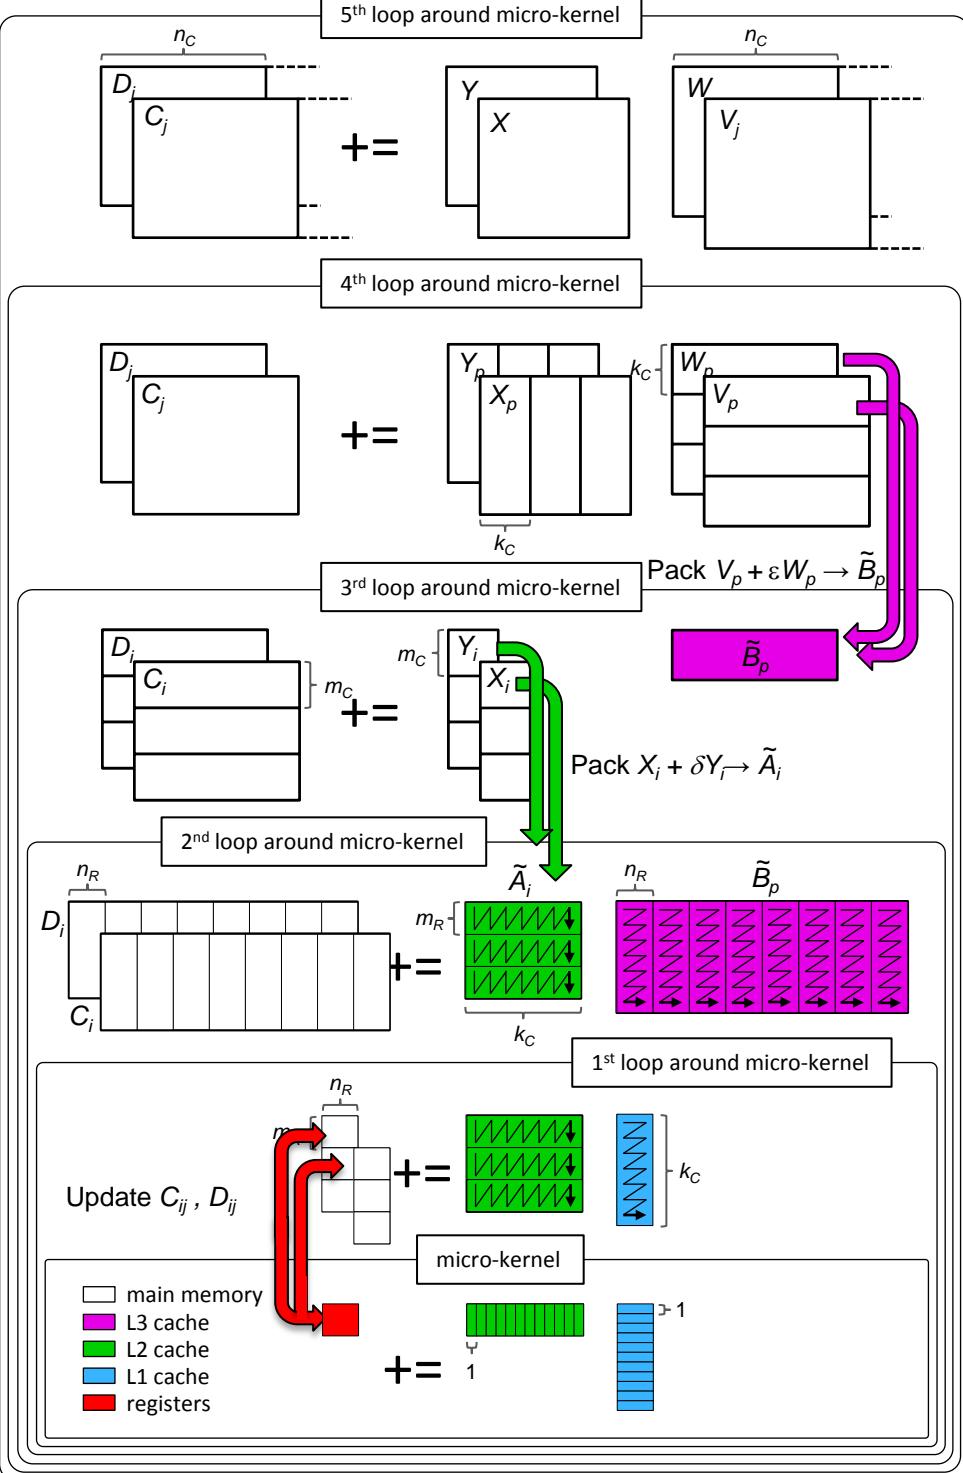
- 2nd loop (along n_R direction)



- both 3rd and 2nd loop



*Tyler M. Smith, Robert Van De Geijn, Mikhail Smelyanskiy, Jeff R. Hammond, and Field G. Van Zee. "Anatomy of high-performance many-threaded matrix multiplication." In *Parallel and Distributed Processing Symposium, 2014 IEEE 28th International*, pp. 1049-1059. IEEE, 2014.



Outline

- Review of State-of-the-art GEMM in BLIS
- Strassen's Algorithm Reloaded
- **Theoretical Model and Practical Performance**
- Extension to Other BLAS-3 Operation
- Extension to Other Fast Matrix Multiplication
- Conclusion

Performance Model

- Performance Metric

$$\text{Effective GFLOPS} = \frac{2 \cdot m \cdot n \cdot k}{\text{time (in seconds)}} \cdot 10^{-9}$$

- Total Time Breakdown

$$T = T_a + T_m$$

The equation $T = T_a + T_m$ is displayed above two blue downward-pointing arrows. The first arrow points to the term T_a and is labeled "Arithmetic Operations" below it. The second arrow points to the term T_m and is labeled "Memory Operations" below it.

Arithmetic Operations

$$T_a = T_a^x + T_a^{A+} + T_a^{B+} + T_a^{C+}$$

Submatrix
multiplication

Extra additions with
submatrices of A, B, C,
respectively

- DGEMM

- No extra additions

$$T_a = 2mnk \cdot \tau_a$$

- One-level Strassen (ABC, AB, Naïve)

- 7 submatrix multiplications
- 5 extra additions of submatrices of A and B
- 12 extra additions of submatrices of C

$$T_a = \left(7 \times 2 \frac{m}{2} \frac{n}{2} \frac{k}{2} + 5 \times 2 \frac{m}{2} \frac{k}{2} + 5 \times 2 \frac{k}{2} \frac{n}{2} + 12 \times 2 \frac{m}{2} \frac{n}{2} \right) \cdot \tau_a$$

- Two-level Strassen (ABC, AB, Naïve)

- 49 submatrix multiplications
- 95 extra additions of submatrices of A and B
- 154 extra additions of submatrices of C

$$T_a = \left(49 \times 2 \frac{m}{4} \frac{n}{4} \frac{k}{4} + 95 \times 2 \frac{m}{4} \frac{k}{4} + 95 \times 2 \frac{k}{4} \frac{n}{4} + 154 \times 2 \frac{m}{4} \frac{n}{4} \right) \cdot \tau_a$$

$$\begin{aligned} M_0 &:= \alpha(A_{00} + A_{11})(B_{00} + B_{11}); & C_{00} &+= M_0; & C_{11} &+= M_0; \\ M_1 &:= \alpha(A_{10} + A_{11})B_{00}; & C_{10} &+= M_1; & C_{11} &-= M_1; \\ M_2 &:= \alpha A_{00}(B_{01} - B_{11}); & C_{01} &+= M_2; & C_{11} &+= M_2; \\ M_3 &:= \alpha A_{11}(B_{10} - B_{00}); & C_{00} &+= M_3; & C_{10} &+= M_3; \\ M_4 &:= \alpha(A_{00} + A_{01})B_{11}; & C_{01} &+= M_4; & C_{00} &-= M_4; \\ M_5 &:= \alpha(A_{10} - A_{00})(B_{00} + B_{01}); & C_{11} &+= M_5; \\ M_6 &:= \alpha(A_{01} - A_{11})(B_{10} + B_{11}); & C_{00} &+= M_6; \end{aligned}$$

$$\begin{aligned} M_0 &= (A_{0,0} + A_{2,2} + A_{1,1} + A_{3,3})(B_{0,0} + B_{2,2} + B_{1,1} + B_{3,3}); & C_{0,0} &+= M_0; & C_{1,1} &+= M_0; & C_{2,2} &+= M_0; & C_{3,3} &+= M_0; \\ M_1 &= (A_{1,0} + A_{3,2} + A_{1,1} + A_{3,3})(B_{0,0} + B_{2,2}); & C_{1,0} &+= M_1; & C_{1,1} &-= M_1; & C_{2,1} &+= M_1; & C_{3,3} &-= M_1; \\ M_2 &= (A_{0,0} + A_{2,2})(B_{0,1} + B_{2,3} + B_{1,1} + B_{3,3}); & C_{0,1} &+= M_2; & C_{1,1} &+= M_2; & C_{2,3} &+= M_2; & C_{3,3} &+= M_2; \\ M_3 &= (A_{1,1} + A_{3,3})(B_{1,0} + B_{3,2} + B_{0,0} + B_{2,2}); & C_{0,0} &+= M_3; & C_{1,0} &+= M_3; & C_{2,2} &+= M_3; & C_{3,2} &+= M_3; \\ M_4 &= (A_{0,0} + A_{2,2} + A_{0,1} + A_{2,3})(B_{1,1} + B_{3,3}); & C_{0,0} &-= M_4; & C_{0,1} &+= M_4; & C_{2,2} &-= M_4; & C_{2,3} &+= M_4; \\ M_5 &= (A_{1,0} + A_{3,2} + A_{0,1} + A_{2,2})(B_{0,0} + B_{2,2} + B_{0,1} + B_{2,3}); & C_{1,1} &+= M_5; & C_{3,3} &+= M_5; \\ M_6 &= (A_{0,1} + A_{2,3} + A_{0,0} + A_{2,2})(B_{1,0} + B_{3,2} + B_{1,1} + B_{3,3}); & C_{0,0} &+= M_6; & C_{2,2} &+= M_6; & C_{3,2} &+= M_6; \\ M_7 &= (A_{2,0} + A_{2,2} + A_{3,1} + A_{3,3})(B_{0,0} + B_{1,1}); & C_{0,1} &+= M_7; & C_{3,1} &+= M_7; & C_{2,2} &-= M_7; & C_{3,3} &-= M_7; \\ M_8 &= (A_{3,0} + A_{3,2} + A_{3,1} + A_{3,3})(B_{0,0}); & C_{3,0} &+= M_8; & C_{3,1} &-= M_8; & C_{3,2} &-= M_8; & C_{3,3} &+= M_8; \\ M_9 &= (A_{2,0} + A_{2,2})(B_{0,1} + B_{1,1}); & C_{2,1} &+= M_9; & C_{3,1} &+= M_9; & C_{3,2} &-= M_9; & C_{3,3} &-= M_9; \\ M_{10} &= (A_{3,1} + A_{3,3})(B_{1,0} + B_{0,0}); & C_{2,0} &+= M_{10}; & C_{3,0} &+= M_{10}; & C_{2,2} &-= M_{10}; & C_{3,2} &-= M_{10}; \end{aligned}$$

$$\begin{aligned} M_{11} &= (A_{3,0} + A_{1,0} + A_{2,0} + A_{0,0})(B_{0,0} + B_{0,2} + B_{0,1} + B_{0,3}); & C_{3,3} &+= M_{10}; \\ M_{40} &= (A_{2,1} + A_{0,1} + A_{3,1} + A_{1,1})(B_{1,0} + B_{1,2} + B_{1,1} + B_{1,3}); & C_{2,2} &+= M_{41}; \\ M_{41} &= (A_{0,2} + A_{2,2} + A_{1,3} + A_{3,3})(B_{2,0} + B_{2,2} + B_{3,1} + B_{3,3}); & C_{0,0} &+= M_{42}; \\ M_{42} &= (A_{1,2} + A_{3,2} + A_{1,3} + A_{3,3})(B_{2,0} + B_{2,2}); & C_{1,0} &+= M_{43}; & C_{1,1} &+= M_{42}; \\ M_{43} &= (A_{0,2} + A_{2,2} + A_{3,1} + A_{3,3})(B_{2,0} + B_{2,2}); & C_{0,1} &+= M_{44}; & C_{1,1} &-= M_{43}; \\ M_{44} &= (A_{0,2} + A_{2,2})(B_{2,1} + B_{2,3} + B_{3,1} + B_{3,3}); & C_{0,0} &+= M_{44}; & C_{1,1} &+= M_{44}; \\ M_{45} &= (A_{1,3} + A_{3,3})(B_{3,0} + B_{2,2} + B_{2,0} + B_{2,2}); & C_{0,0} &+= M_{45}; & C_{1,0} &+= M_{45}; \\ M_{46} &= (A_{0,2} + A_{2,2} + A_{0,3} + A_{2,3})(B_{3,0}, B_{3,3}); & C_{0,0} &-= M_{46}; & C_{1,0} &+= M_{46}; \\ M_{47} &= (A_{1,2} + A_{3,2} + A_{0,2} + A_{2,3})(B_{2,0} + B_{2,2} + B_{2,1} + B_{2,3}); & C_{1,1} &+= M_{47}; \\ M_{48} &= (A_{0,3} + A_{2,3} + A_{1,3} + A_{3,3})(B_{3,0}, B_{3,2} + B_{3,1} + B_{3,3}); & C_{0,0} &+= M_{48}; \end{aligned}$$

Memory Operations

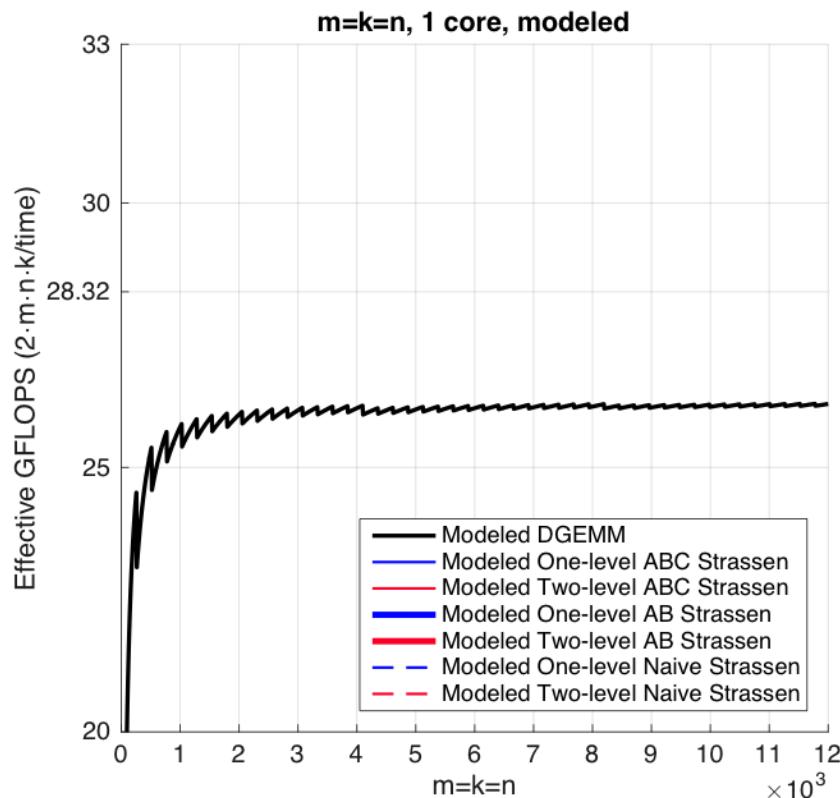
$$T_m = N_m^{A \times} \cdot T_m^{A \times} + N_m^{B \times} \cdot T_m^{B \times} + N_m^{C \times} \cdot T_m^{C \times} + N_m^{A+} \cdot T_m^{A+} + N_m^{B+} \cdot T_m^{B+} + N_m^{C+} \cdot T_m^{C+}$$

- DGEMM $T_m = (1 \cdot mk \lceil \frac{n}{n_c} \rceil + 1 \cdot nk + 1 \cdot 2\lambda mn \lceil \frac{k}{k_c} \rceil) \cdot \tau_b$
- One-level
 - ABC Strassen $T_m = (12 \cdot \frac{m}{2} \frac{k}{2} \lceil \frac{n/2}{n_c} \rceil + 12 \cdot \frac{n}{2} \frac{k}{2} + 12 \cdot 2\lambda \frac{m}{2} \frac{n}{2} \lceil \frac{k/2}{k_c} \rceil) \cdot \tau_b$
 - AB Strassen $T_m = (12 \cdot \frac{m}{2} \frac{k}{2} \lceil \frac{n/2}{n_c} \rceil + 12 \cdot \frac{n}{2} \frac{k}{2} + 7 \cdot 2\lambda \frac{m}{2} \frac{n}{2} \lceil \frac{k/2}{k_c} \rceil + 36 \cdot \frac{m}{2} \frac{n}{2}) \cdot \tau_b$
 - Naïve Strassen $T_m = (7 \cdot \frac{m}{2} \frac{k}{2} \lceil \frac{n/2}{n_c} \rceil + 7 \cdot \frac{n}{2} \frac{k}{2} + 7 \cdot 2\lambda \frac{m}{2} \frac{n}{2} \lceil \frac{k/2}{k_c} \rceil + 19 \cdot \frac{m}{2} \frac{k}{2} + 19 \cdot \frac{n}{2} \frac{k}{2} + 36 \cdot \frac{m}{2} \frac{n}{2}) \cdot \tau_b$
- Two-level
 - ABC Strassen $T_m = (194 \cdot \frac{m}{4} \frac{k}{4} \lceil \frac{n/4}{n_c} \rceil + 194 \cdot \frac{n}{4} \frac{k}{4} + 154 \cdot 2\lambda \frac{m}{4} \frac{n}{4} \lceil \frac{k/4}{k_c} \rceil) \cdot \tau_b$
 - AB Strassen $T_m = (194 \cdot \frac{m}{4} \frac{k}{4} \lceil \frac{n/4}{n_c} \rceil + 194 \cdot \frac{n}{4} \frac{k}{4} + 49 \cdot 2\lambda \frac{m}{4} \frac{n}{4} \lceil \frac{k/4}{k_c} \rceil + 462 \cdot \frac{m}{4} \frac{n}{4}) \cdot \tau_b$
 - Naïve Strassen $T_m = (49 \cdot \frac{m}{4} \frac{k}{4} \lceil \frac{n/4}{n_c} \rceil + 49 \cdot \frac{n}{4} \frac{k}{4} + 49 \cdot 2\lambda \frac{m}{4} \frac{n}{4} \lceil \frac{k/4}{k_c} \rceil + 293 \cdot \frac{m}{4} \frac{k}{4} + 293 \cdot \frac{n}{4} \frac{k}{4} + 462 \cdot \frac{m}{4} \frac{n}{4}) \cdot \tau_b$

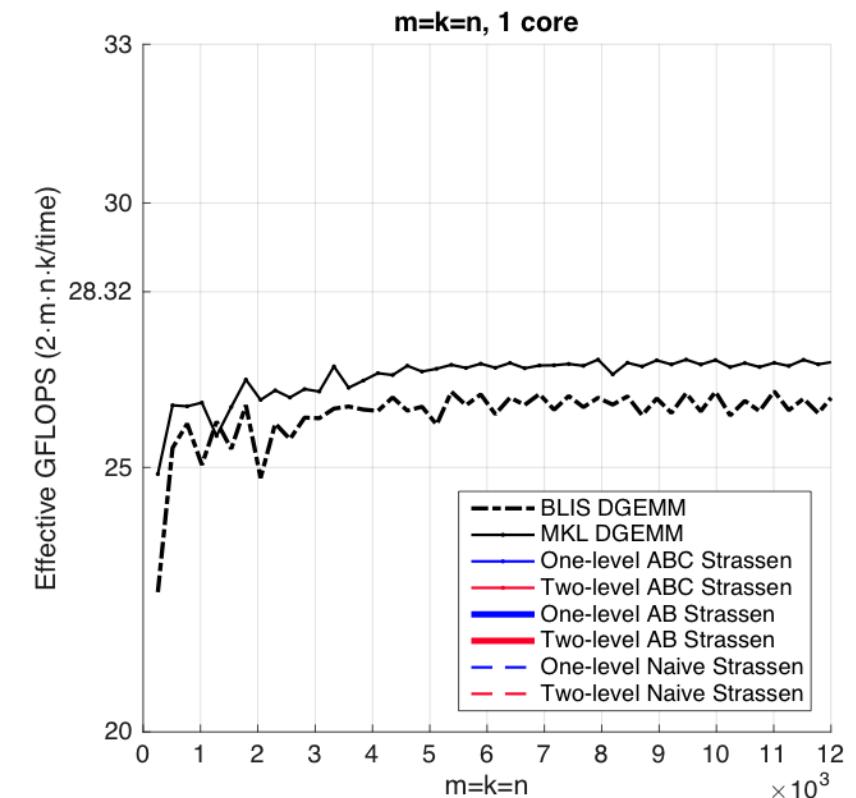
Modeled and Actual Performance on Single Core

Observation (Square Matrices)

Modeled Performance

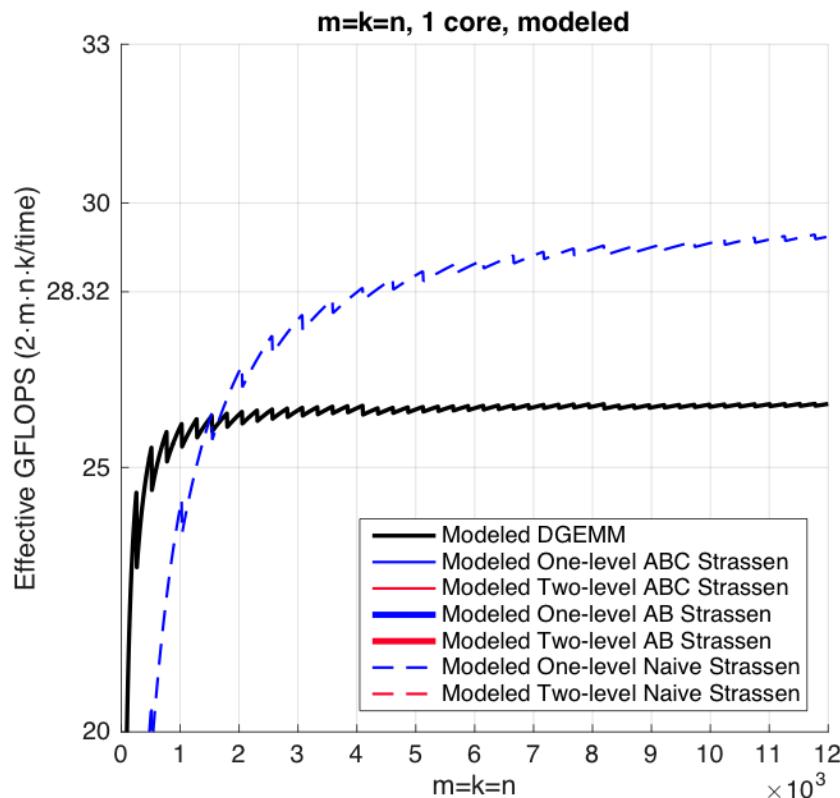


Actual Performance

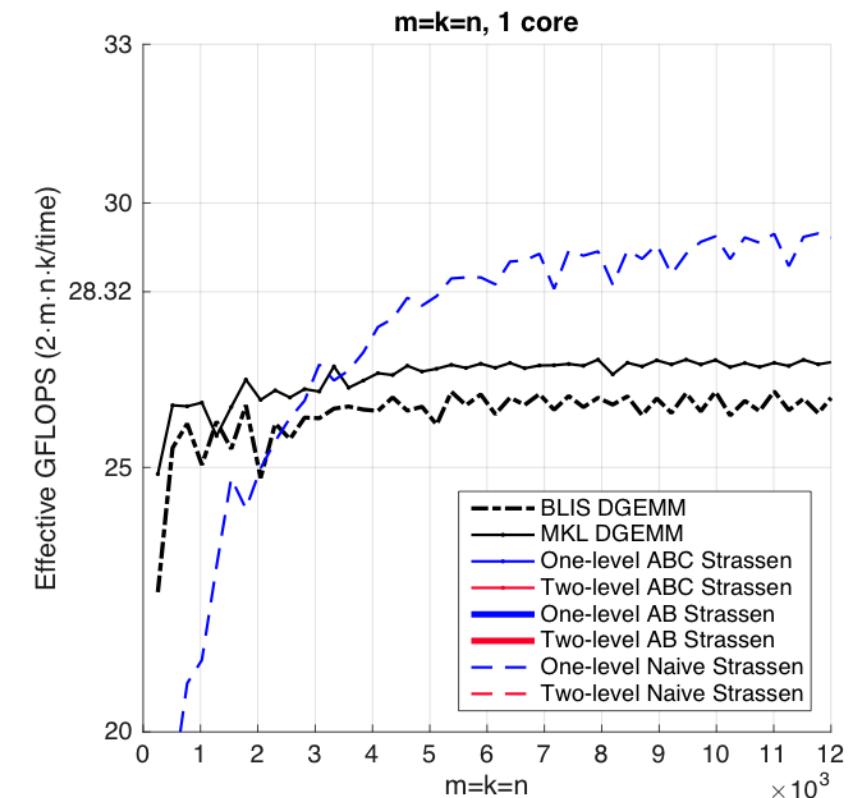


Observation (Square Matrices)

Modeled Performance

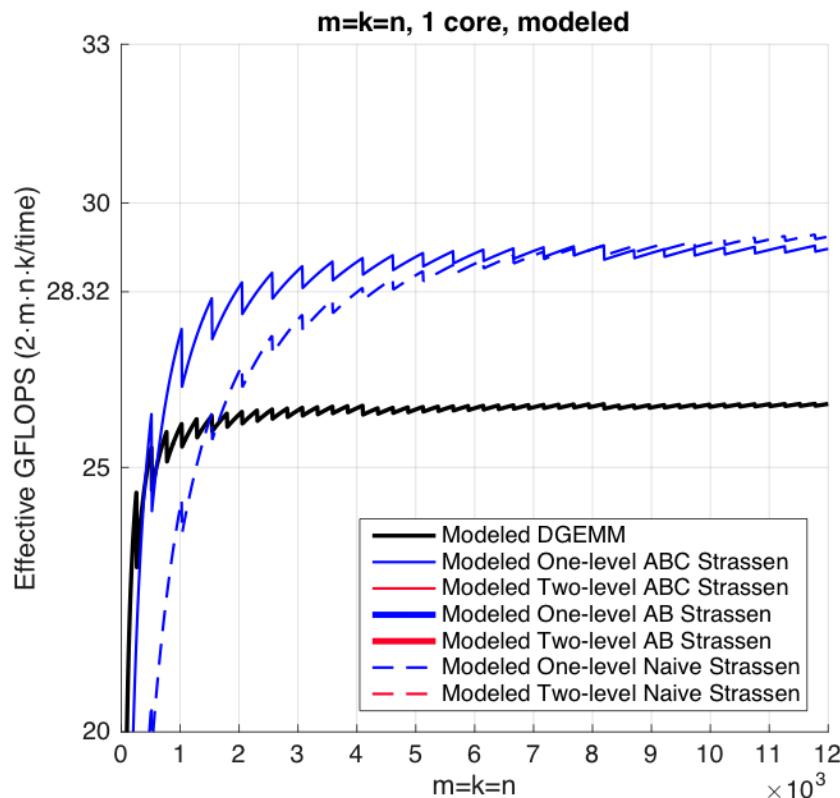


Actual Performance

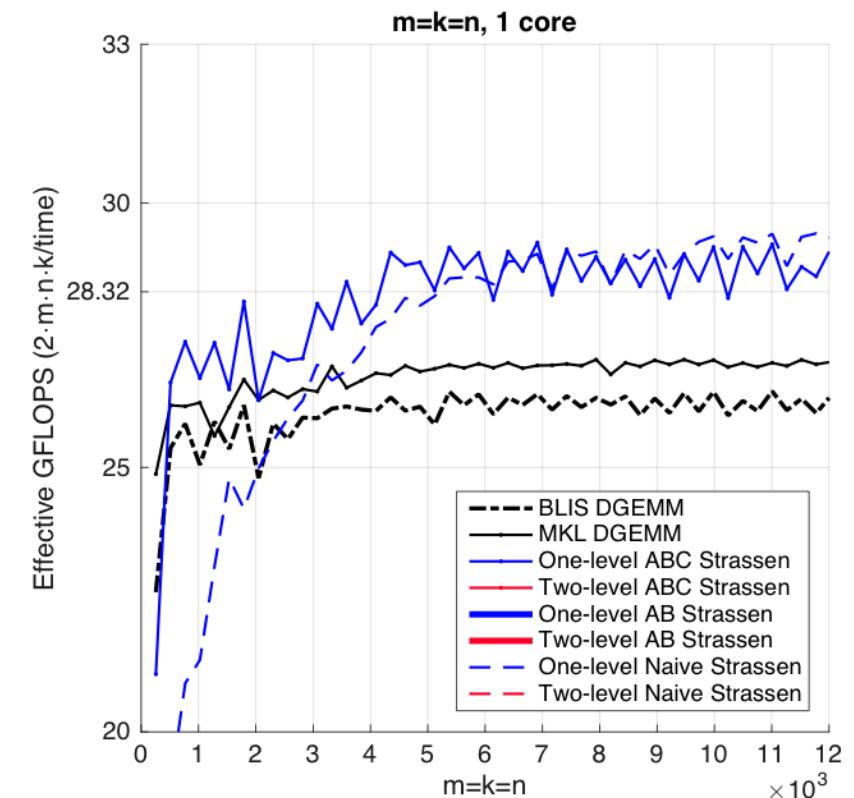


Observation (Square Matrices)

Modeled Performance

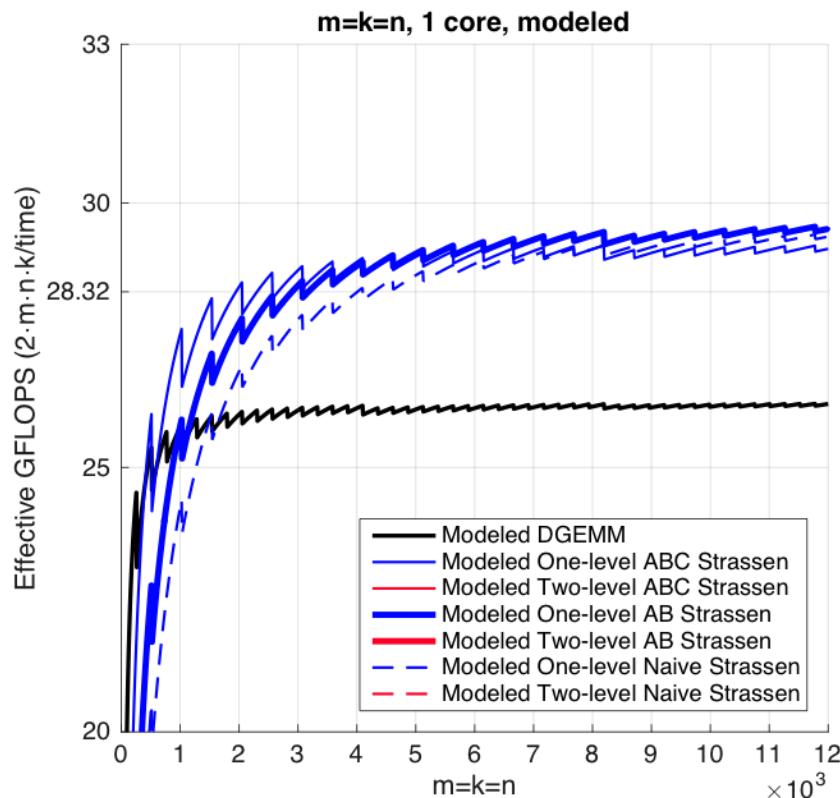


Actual Performance

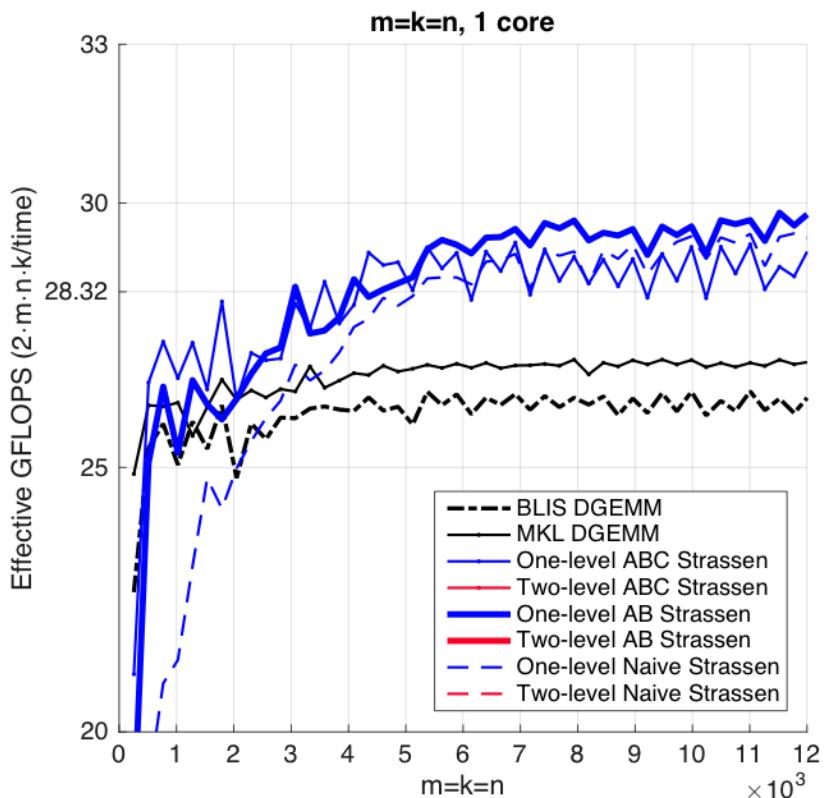


Observation (Square Matrices)

Modeled Performance

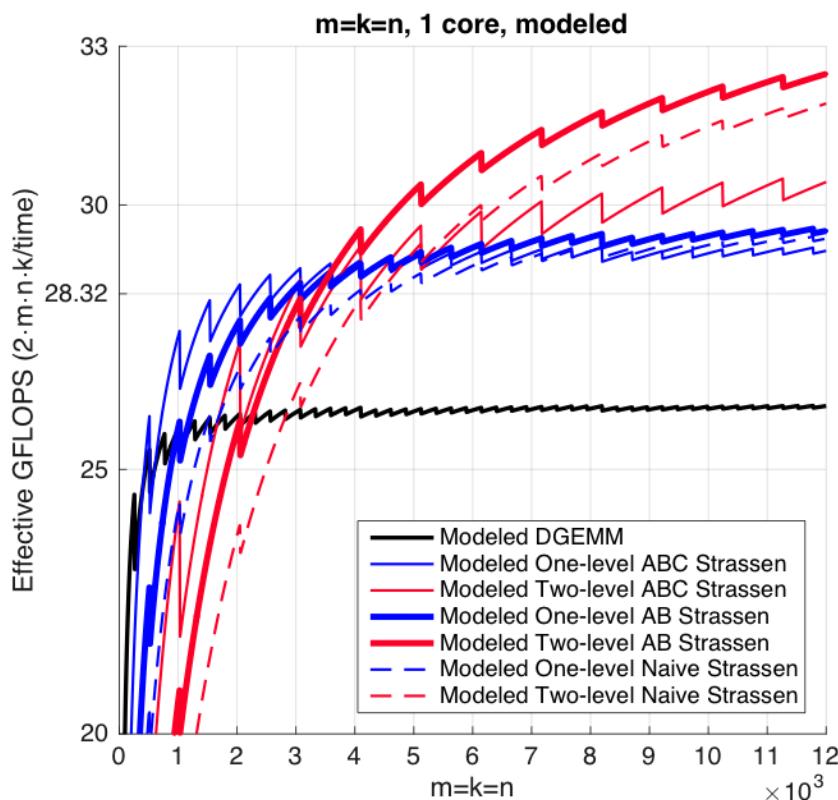


Actual Performance

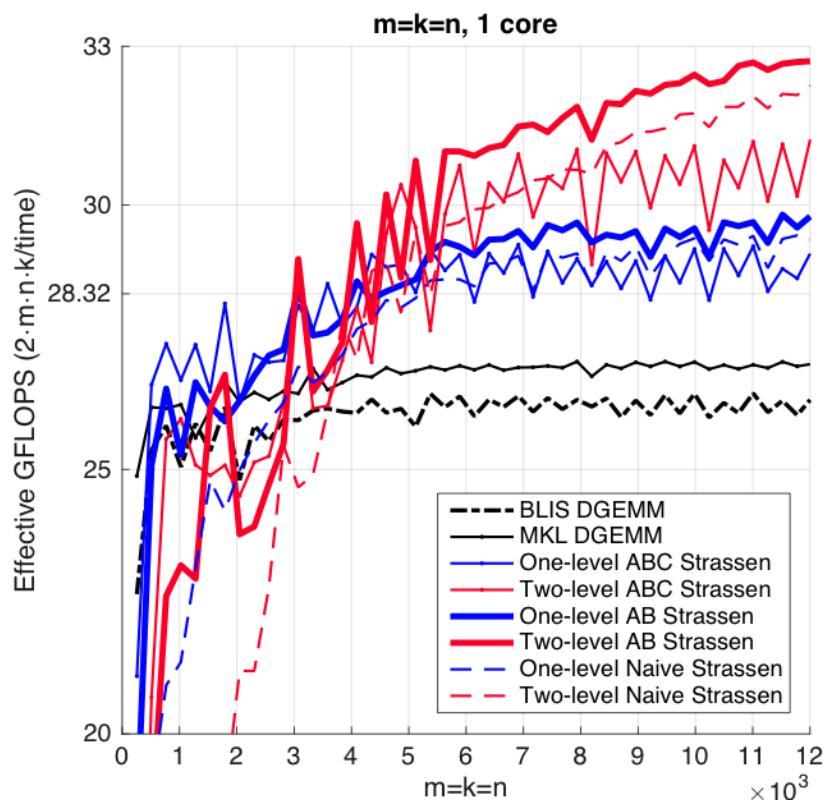


Observation (Square Matrices)

Modeled Performance

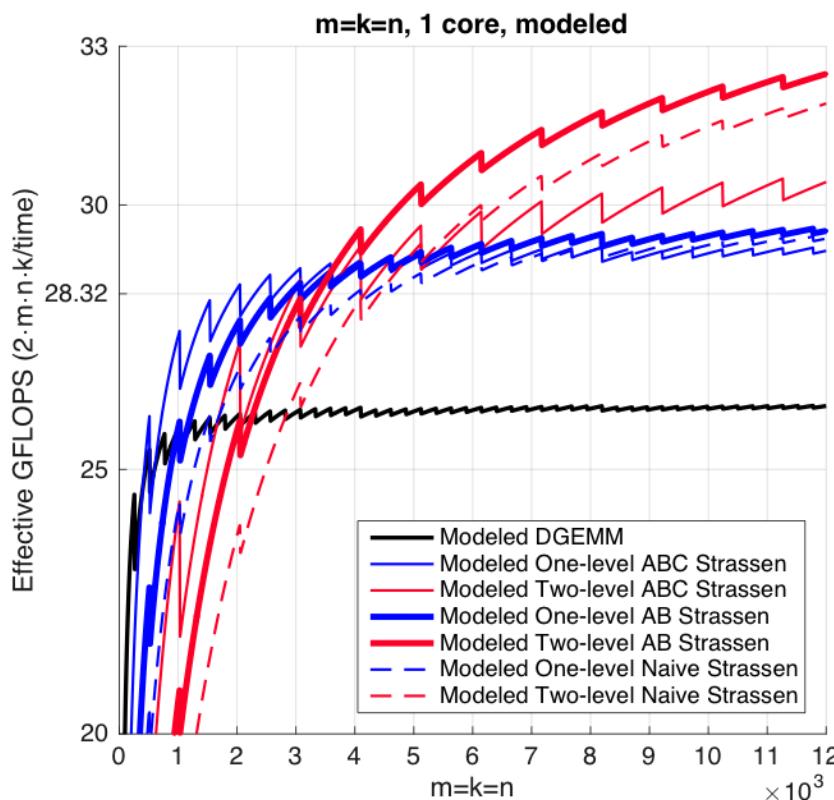


Actual Performance

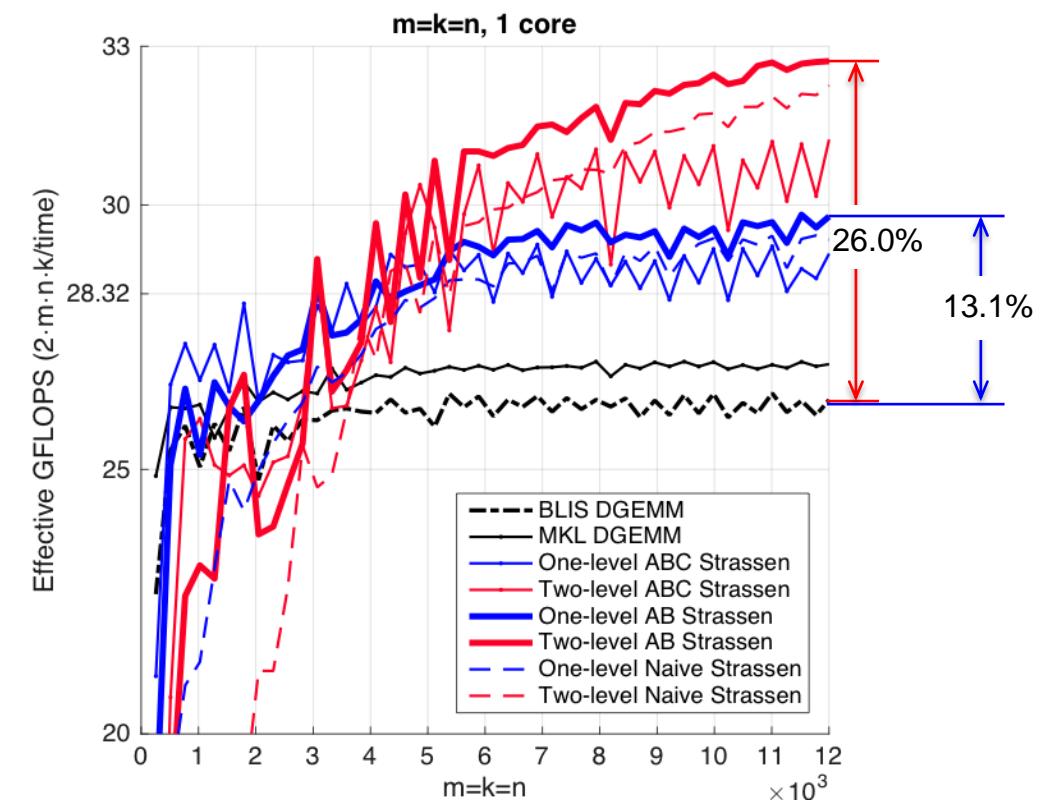


Observation (Square Matrices)

Modeled Performance



Actual Performance

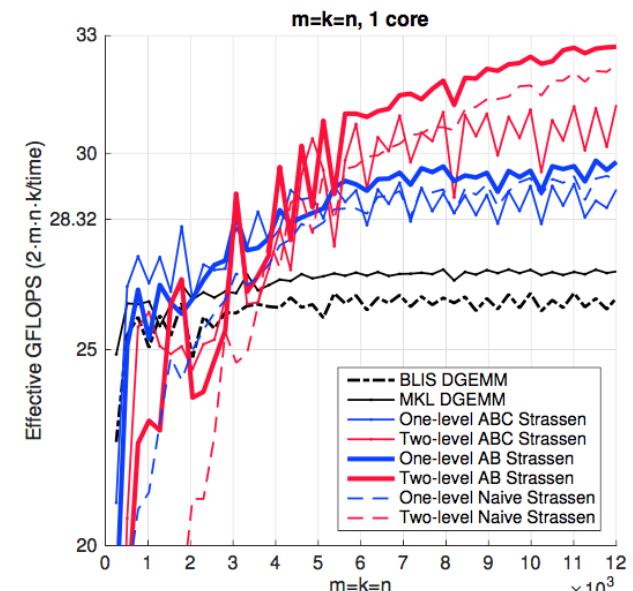
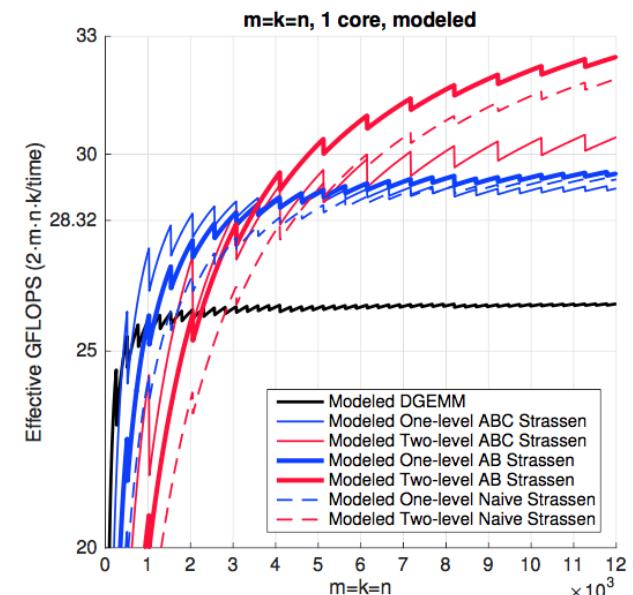


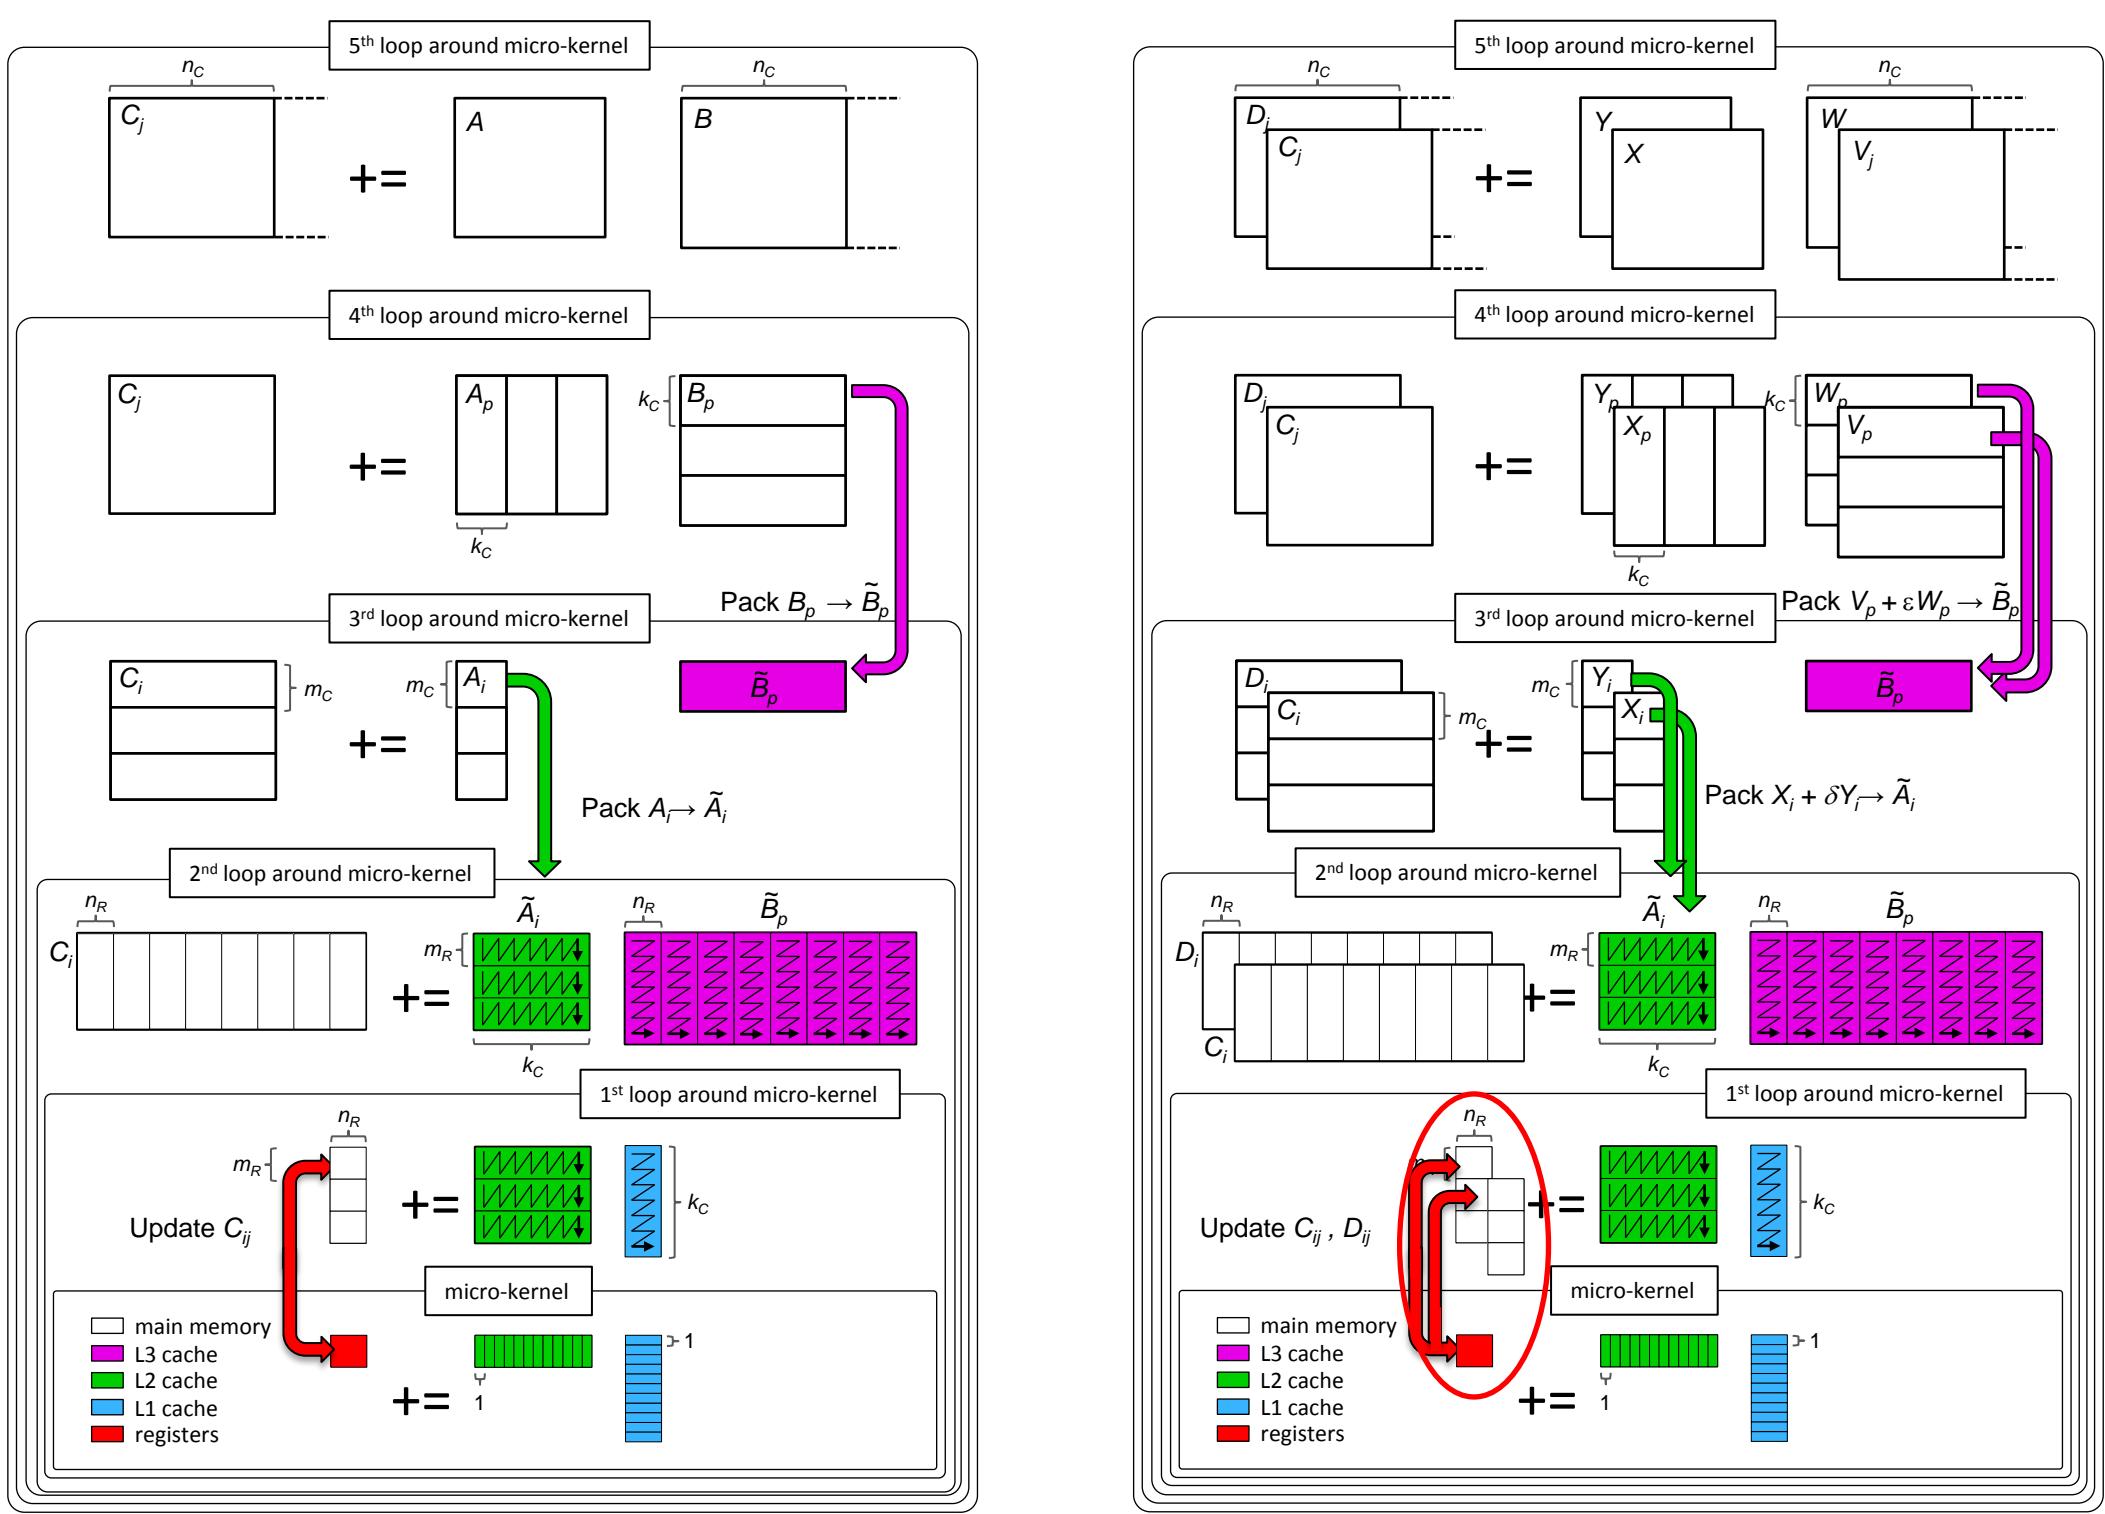
Theoretical Speedup over DGEMM

- One-level Strassen (1+14.3% speedup)
 - 8 multiplications → 7 multiplications;
- Two-level Strassen (1+30.6% speedup)
 - 64 multiplications → 49 multiplications;

Observation (Square Matrices)

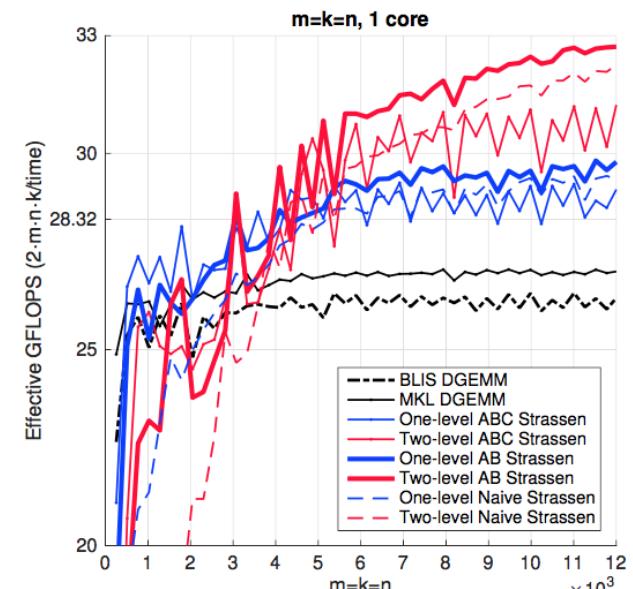
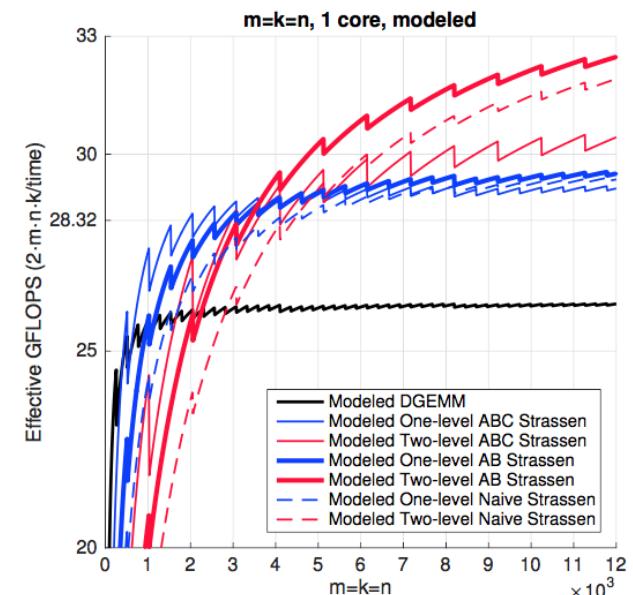
- Both one-level and two-level
 - For small square matrices, **ABC Strassen** outperforms **AB Strassen**
 - For larger square matrices, this trend reverses
- Reason
 - **ABC Strassen** avoids storing M (M resides in the register) 😊
 - **ABC Strassen** increases the number of times for updating submatrices of C 😞





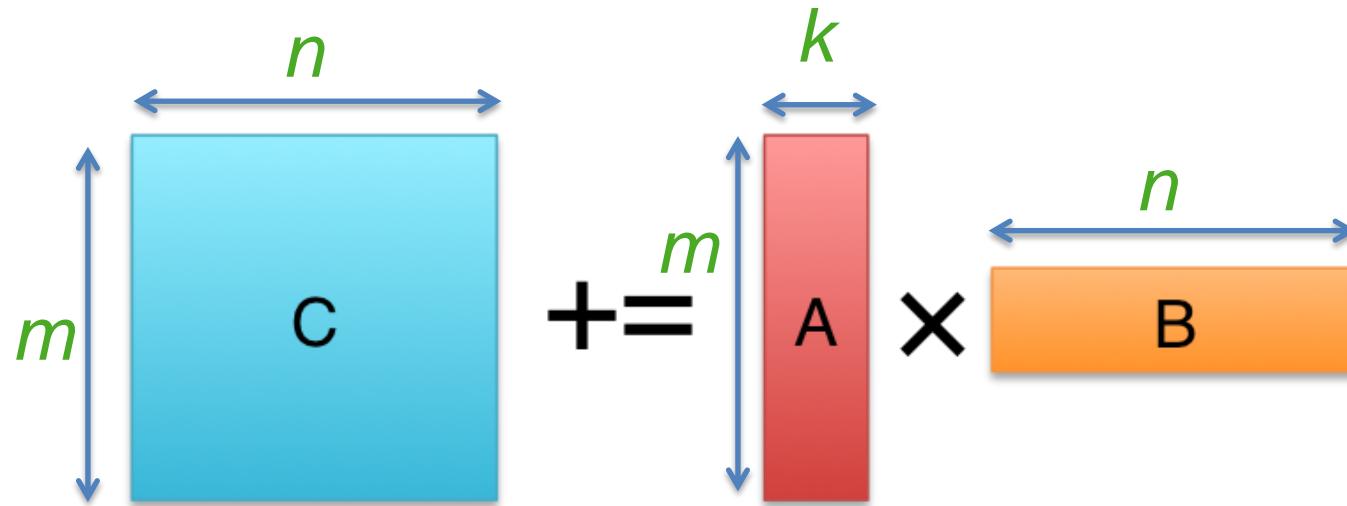
Observation (Square Matrices)

- Both one-level and two-level
 - For small square matrices, **ABC Strassen** outperforms **AB Strassen**
 - For larger square matrices, this trend reverses
- Reason
 - **ABC Strassen** avoids storing M (M resides in the register) 😊
 - **ABC Strassen** increases the number of times for updating submatrices of C 😞



Observation (Rank-k Update)

- What is Rank-k update?



Observation (Rank-k Update)

- Importance of Rank-k update

Numer. Math. 13, 354–356 (1969)

Gaussian Elimination is not Optimal

VOLKER STRASSEN*

Received December 12, 1968

1. Below we will give an algorithm which computes the coefficients of the product of two square matrices A and B of order n from the coefficients of A and B with less than $4.7 \cdot n^{\log 7}$ arithmetical operations (all logarithms in this paper are for base 2, thus $\log 7 \approx 2.8$; the usual method requires approximately $2n^3$ arithmetical operations). The algorithm induces algorithms for inverting a matrix of order n , solving a system of n linear equations in n unknowns, computing a determinant of order n etc. all requiring less than $\text{const } n^{\log 7}$ arithmetical operations.

This fact should be compared with the result of KLYUYEV and KOKOVKIN-SCHERBAK [1] that Gaussian elimination for solving a system of linear equations is optimal if one restricts oneself to operations upon rows and columns as a whole. We also note that WINOGRAD [2] modifies the usual algorithms for matrix multiplication and inversion and for solving systems of linear equations, trading roughly half of the multiplications for additions and subtractions.

It is a pleasure to thank D. BRILLINGER for inspiring discussions about the present subject and St. COOK and B. PARLETT for encouraging me to write this paper.

2. We define algorithms $\alpha_{m,k}$ which multiply matrices of order $m2^k$, by induction on k : $\alpha_{m,0}$ is the usual algorithm for matrix multiplication (requiring m^3 multiplications and $m^2(m-1)$ additions). $\alpha_{m,k}$ already being known, define $\alpha_{m,k+1}$ as follows:

If A, B are matrices of order $m2^{k+1}$ to be multiplied, write

$$A = \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix}, \quad B = \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix}, \quad AB = \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix},$$

where the A_{ik}, B_{ik}, C_{ik} are matrices of order $m2^k$. Then compute

$$\begin{aligned} I &= (A_{11} + A_{22})(B_{11} + B_{22}), \\ II &= (A_{21} + A_{22})B_{11}, \\ III &= A_{11}(B_{12} - B_{22}), \\ IV &= A_{22}(-B_{11} + B_{21}), \\ V &= (A_{11} + A_{12})B_{22}, \\ VI &= (-A_{11} + A_{21})(B_{11} + B_{12}), \\ VII &= (A_{12} - A_{22})(B_{21} + B_{22}), \end{aligned}$$

Blocked LU with partial pivoting (**getrf**)

Algorithm: $[A, p] := \text{LUPIV_BLK}(A)$

Partition

$$A \rightarrow \left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right), \quad p \rightarrow \left(\begin{array}{c} p_T \\ p_B \end{array} \right)$$

where A_{TL} is 0×0 , p_T has 0 elements

while $n(A_{TL}) < n(A)$ **do**

Determine block size b

Repartition

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \rightarrow \left(\begin{array}{c|c|c} A_{00} & A_{01} & A_{02} \\ \hline A_{10} & A_{11} & A_{12} \\ \hline A_{20} & A_{21} & A_{22} \end{array} \right),$$

$$\left(\begin{array}{c} p_T \\ p_B \end{array} \right) \rightarrow \left(\begin{array}{c} p_0 \\ p_1 \\ p_2 \end{array} \right)$$

where A_{11} is $b \times b$, p_1 is $b \times 1$

$$\left[\left(\begin{array}{c} A_{11} \\ \hline A_{21} \end{array} \right), p_1 \right] := \text{LUPIV_UNB} \left(\begin{array}{c} A_{11} \\ \hline A_{21} \end{array} \right)$$

$$\left(\begin{array}{c|c} A_{10} & A_{12} \\ \hline A_{20} & A_{22} \end{array} \right) := \text{PIV} \left(p_1, \left(\begin{array}{c|c} A_{10} & A_{12} \\ \hline A_{20} & A_{22} \end{array} \right) \right)$$

$$A_{12} := L_{11}^{-1} A_{12}$$

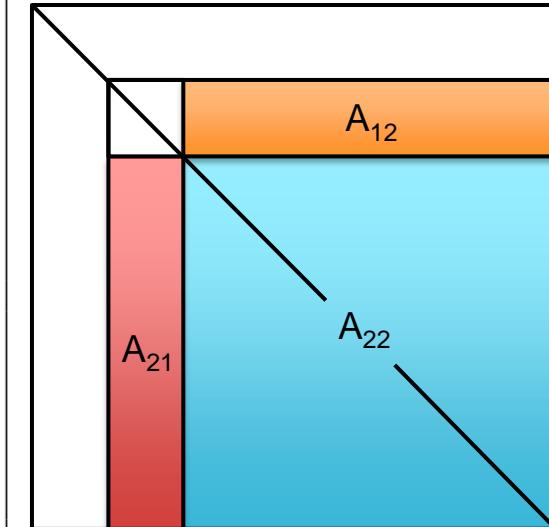
$$A_{22} := A_{22} - A_{21} A_{12}$$

Continue with

$$\left(\begin{array}{c|c} A_{TL} & A_{TR} \\ \hline A_{BL} & A_{BR} \end{array} \right) \leftarrow \left(\begin{array}{c|c|c} A_{00} & A_{01} & A_{02} \\ \hline A_{10} & A_{11} & A_{12} \\ \hline A_{20} & A_{21} & A_{22} \end{array} \right),$$

$$\left(\begin{array}{c} p_T \\ p_B \end{array} \right) \leftarrow \left(\begin{array}{c} p_0 \\ p_1 \\ p_2 \end{array} \right)$$

endwhile

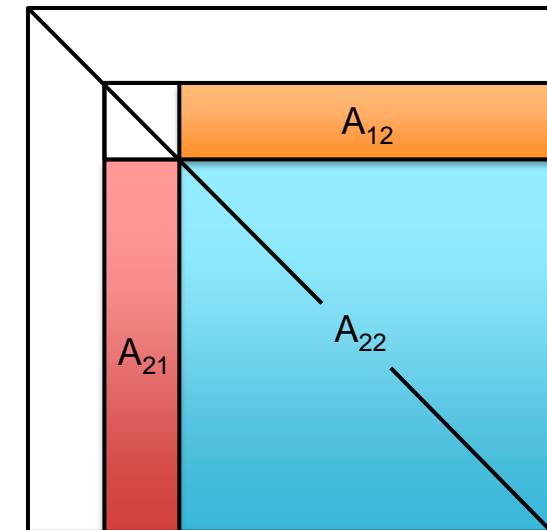


* The results have been found while the author was at the Department of Statistics of the University of California, Berkeley. The author wishes to thank the National Science Foundation for their support (NSF GP-7454).

Observation (Rank-k Update)

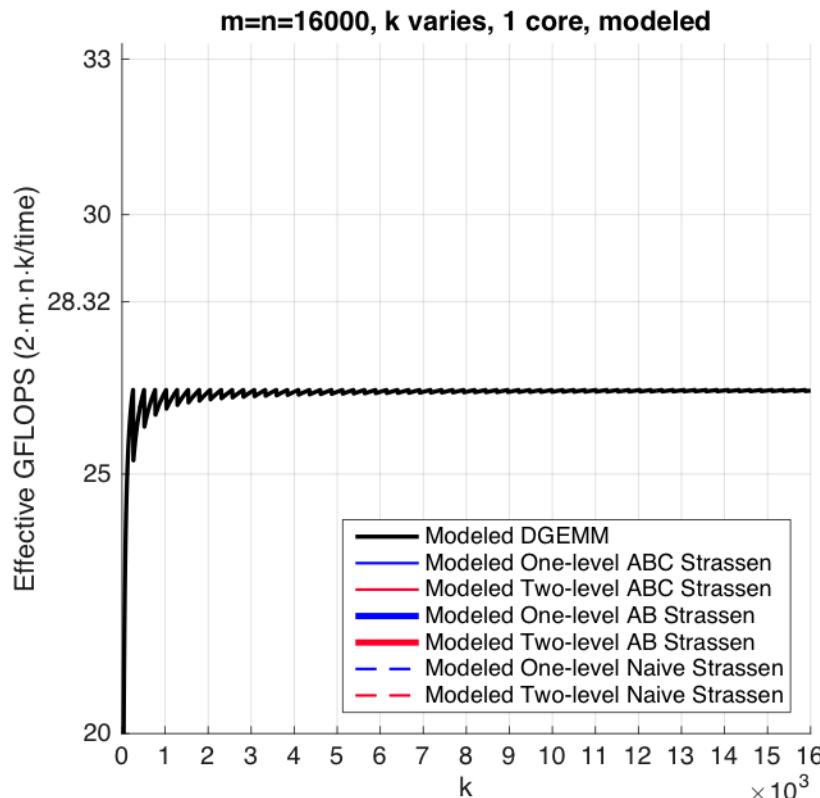
- Importance of Rank-k update

$+=$ \times

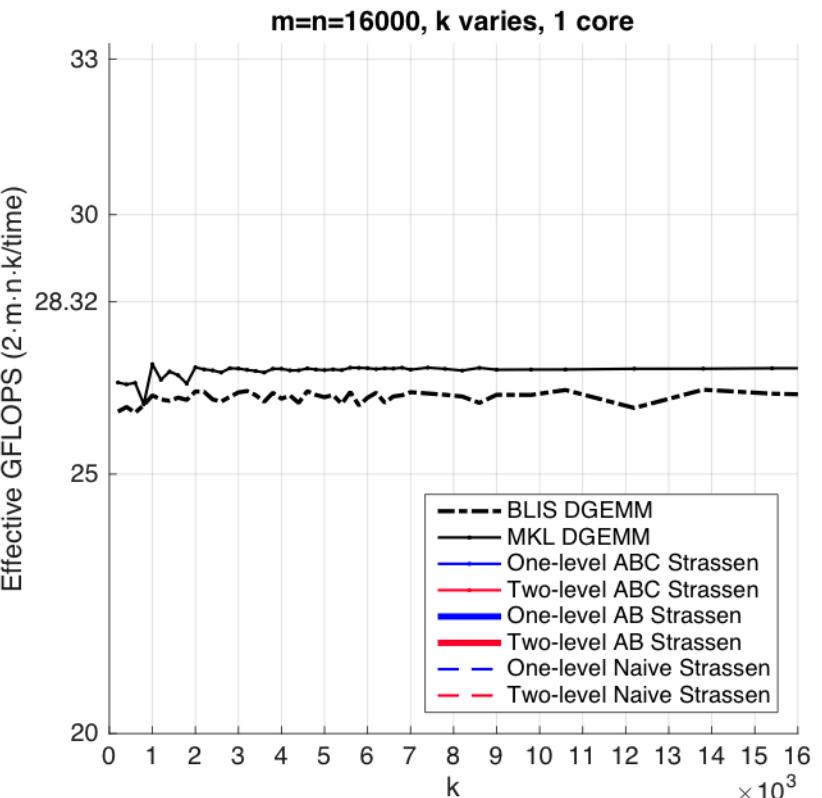


Observation (Rank-k Update)

Modeled Performance

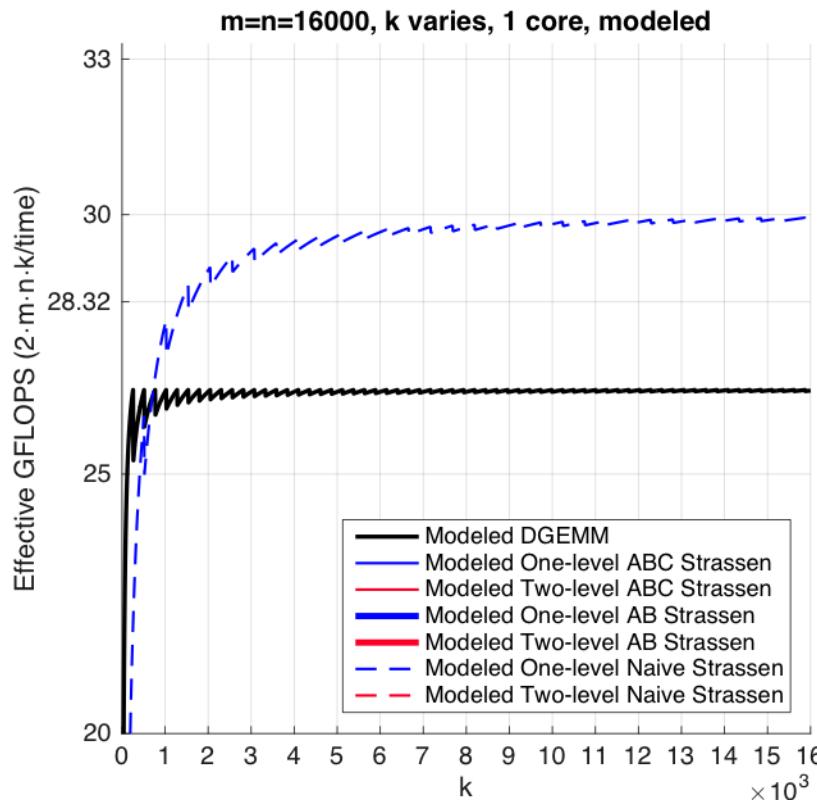


Actual Performance

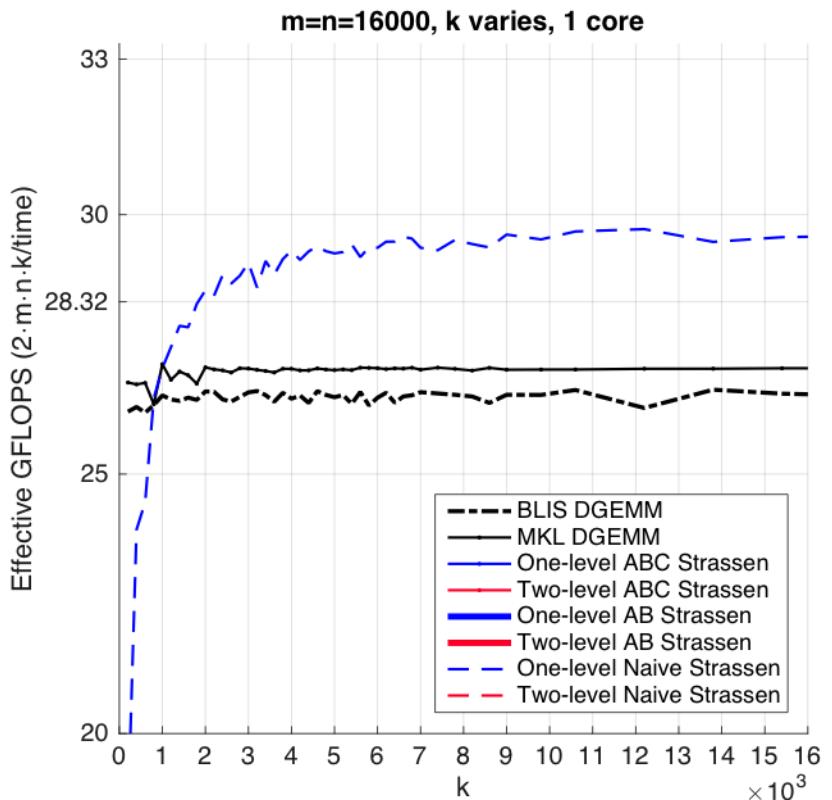


Observation (Rank-k Update)

Modeled Performance

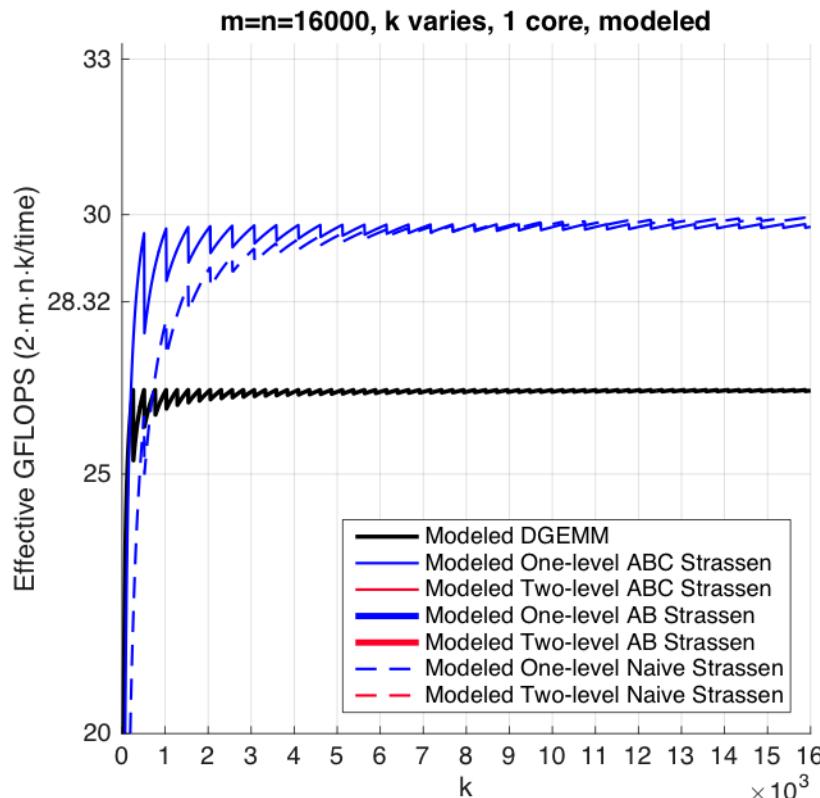


Actual Performance

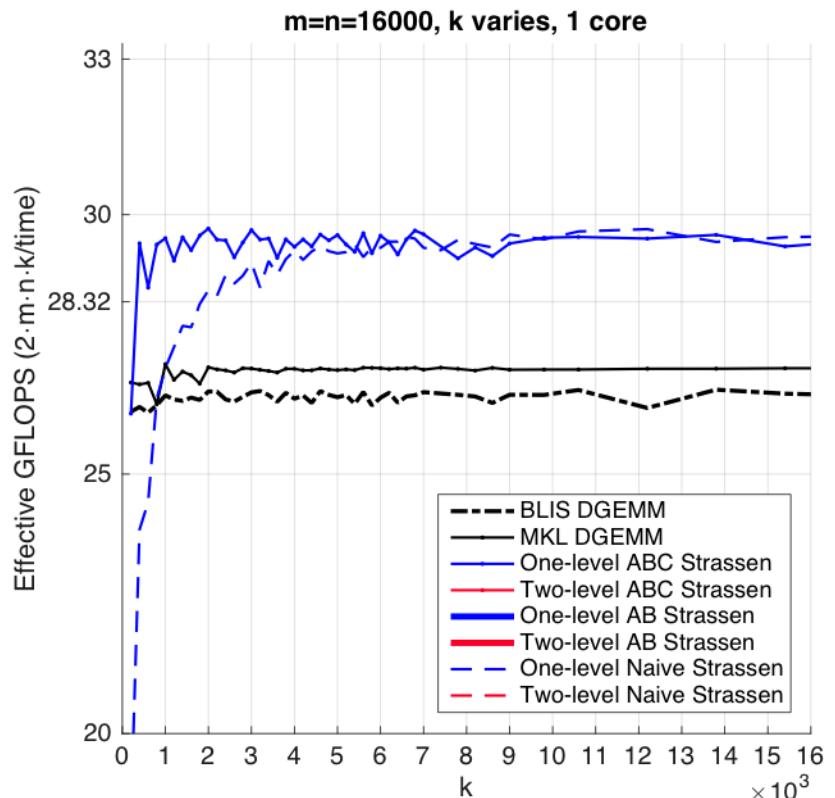


Observation (Rank-k Update)

Modeled Performance

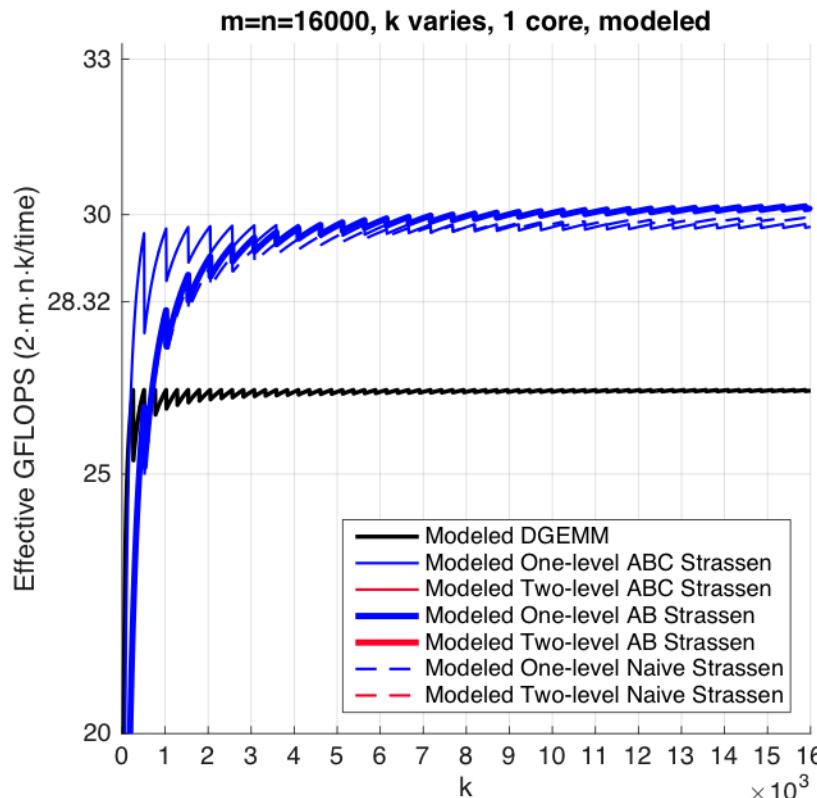


Actual Performance

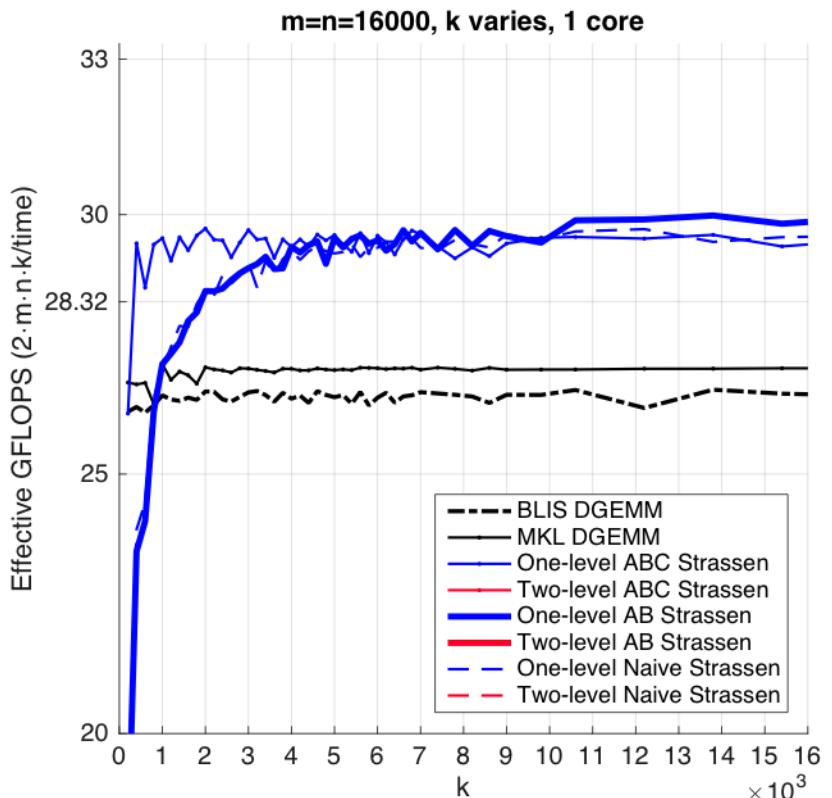


Observation (Rank-k Update)

Modeled Performance

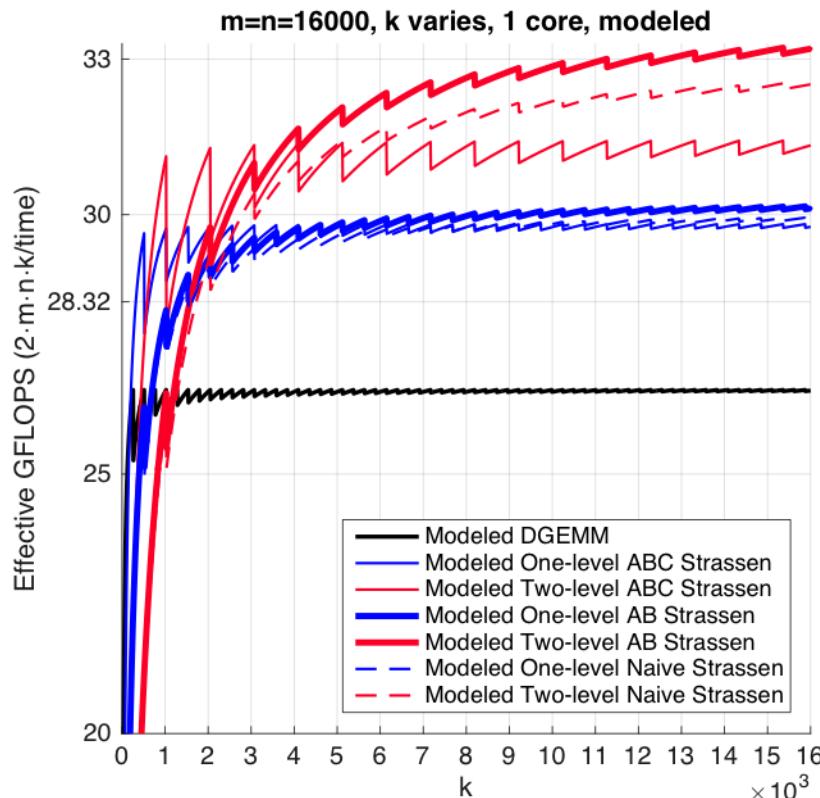


Actual Performance

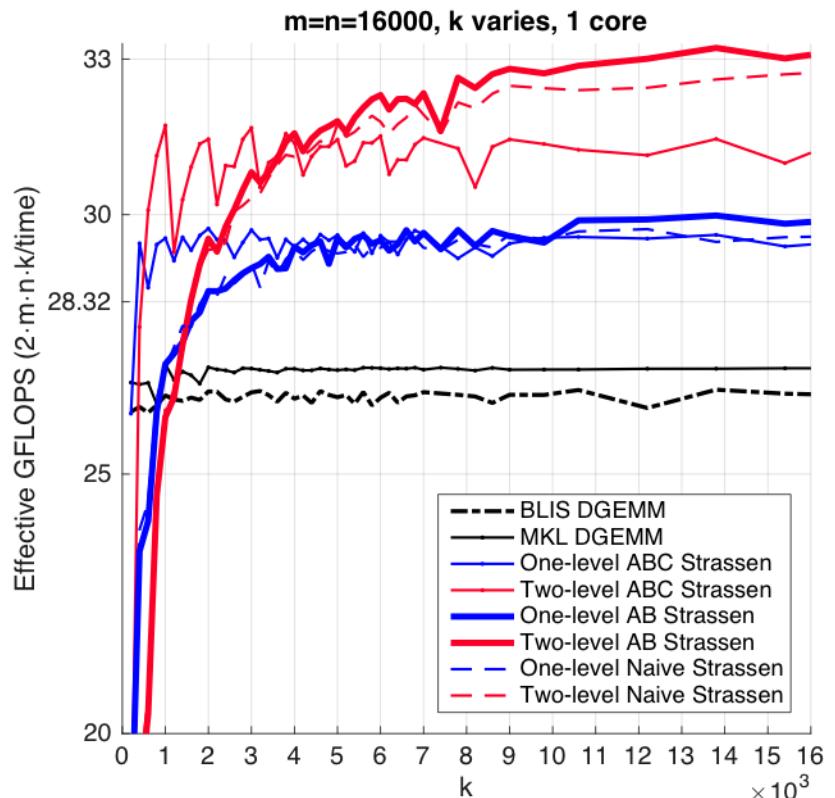


Observation (Rank-k Update)

Modeled Performance

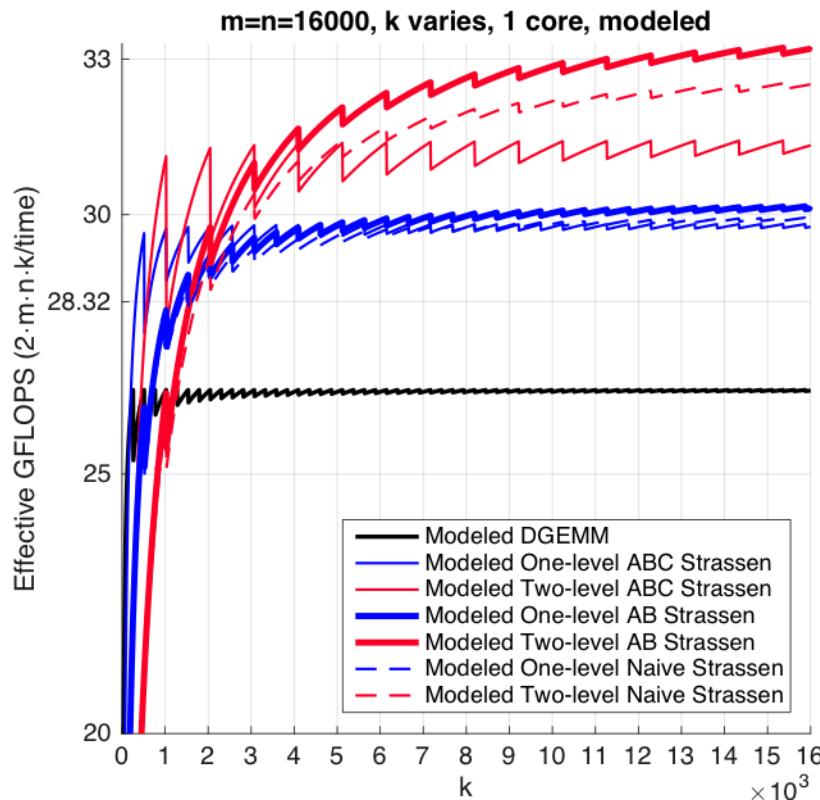


Actual Performance

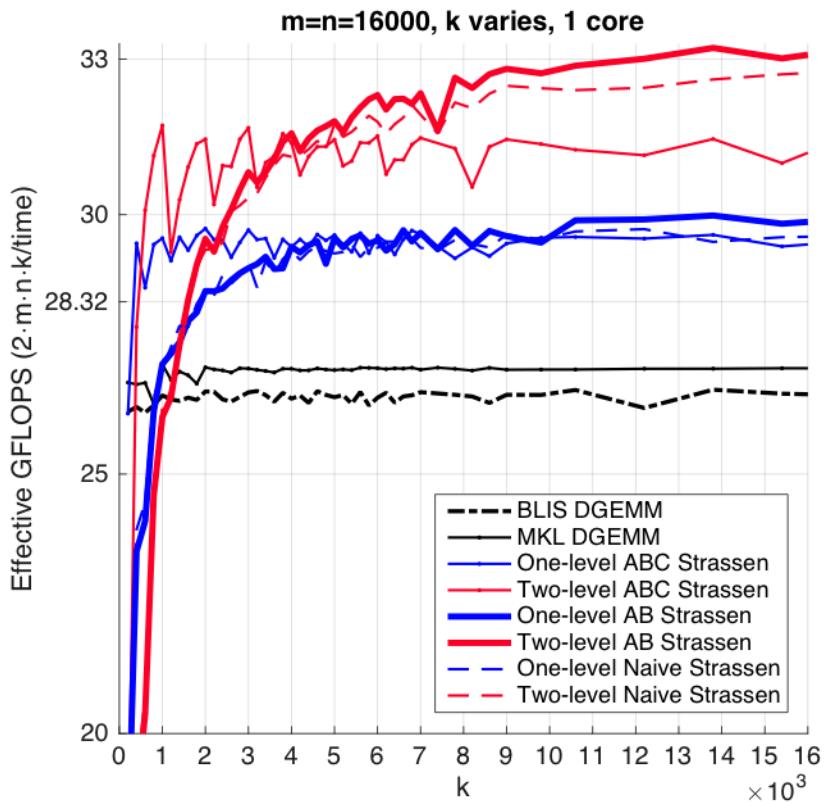


Observation (Rank-k Update)

Modeled Performance



Actual Performance

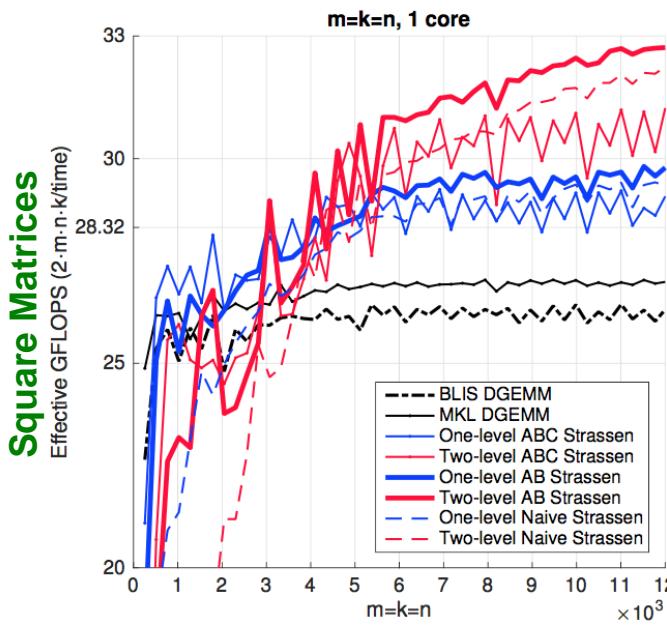


- Reason:

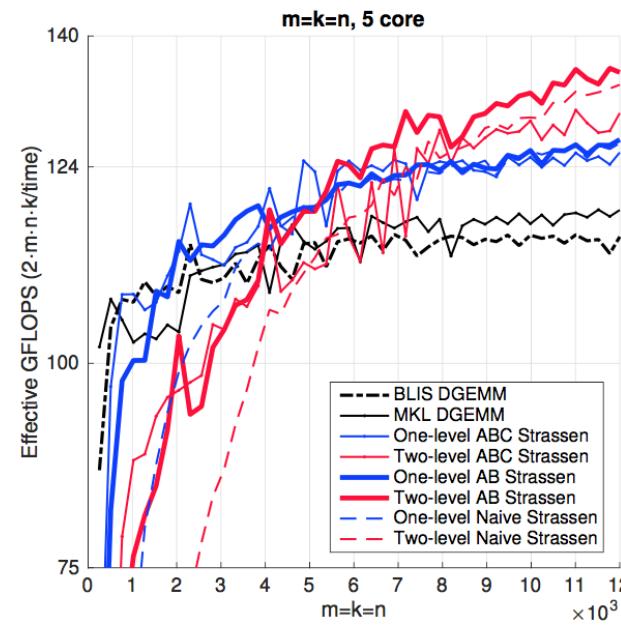
ABC Strassen avoids forming the temporary matrix M explicitly in the memory (M resides in register), especially important when $m, n \gg k$.

Single Node Experiment

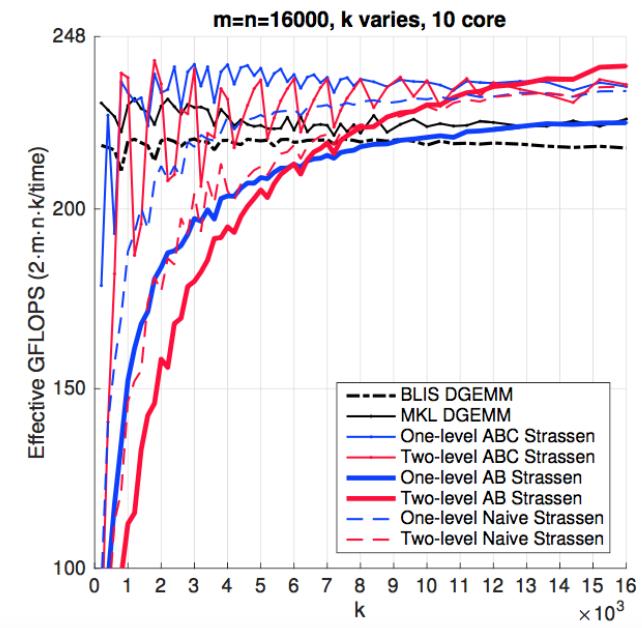
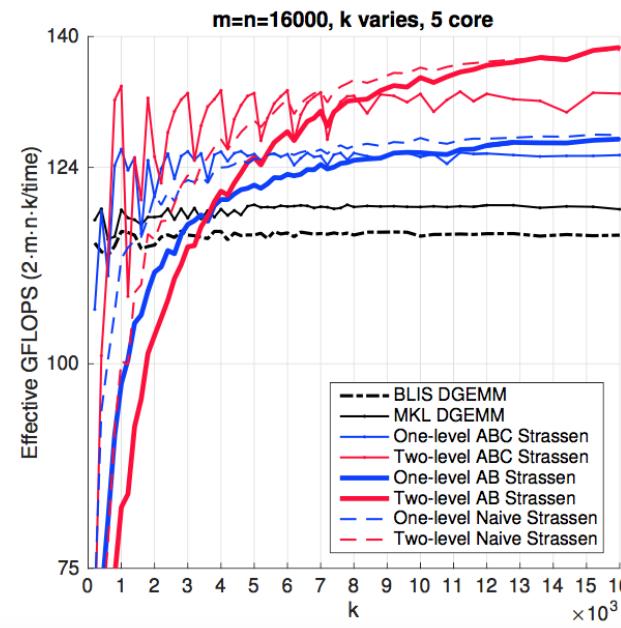
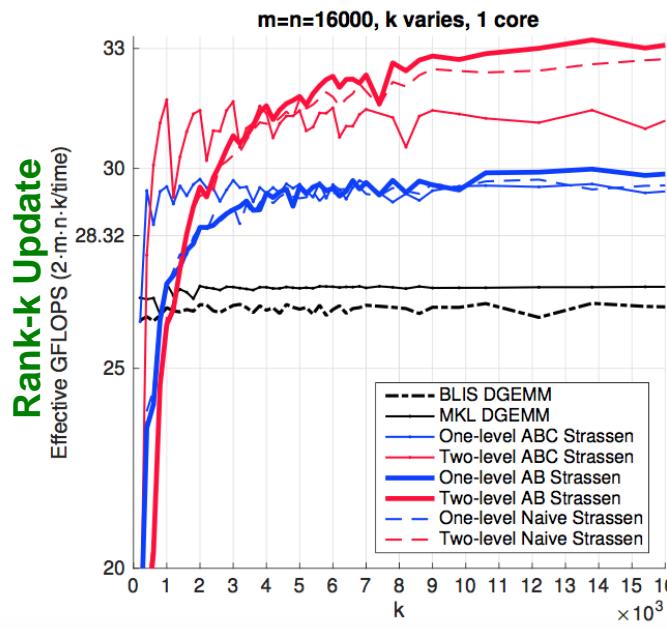
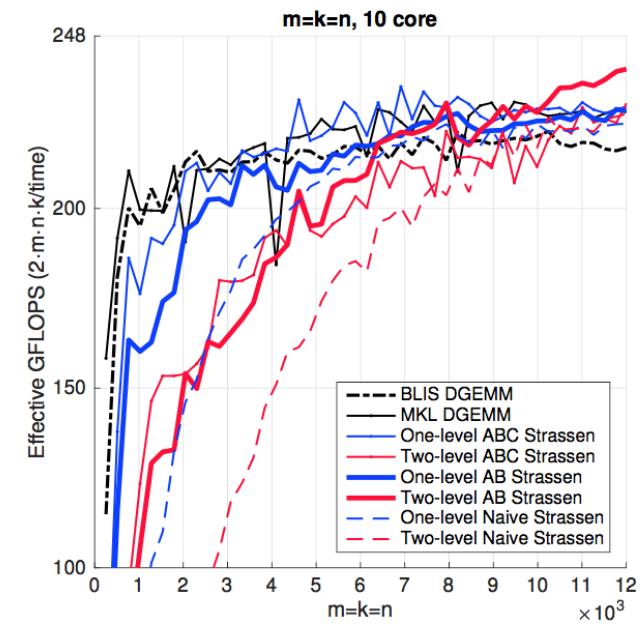
1 core



5 core

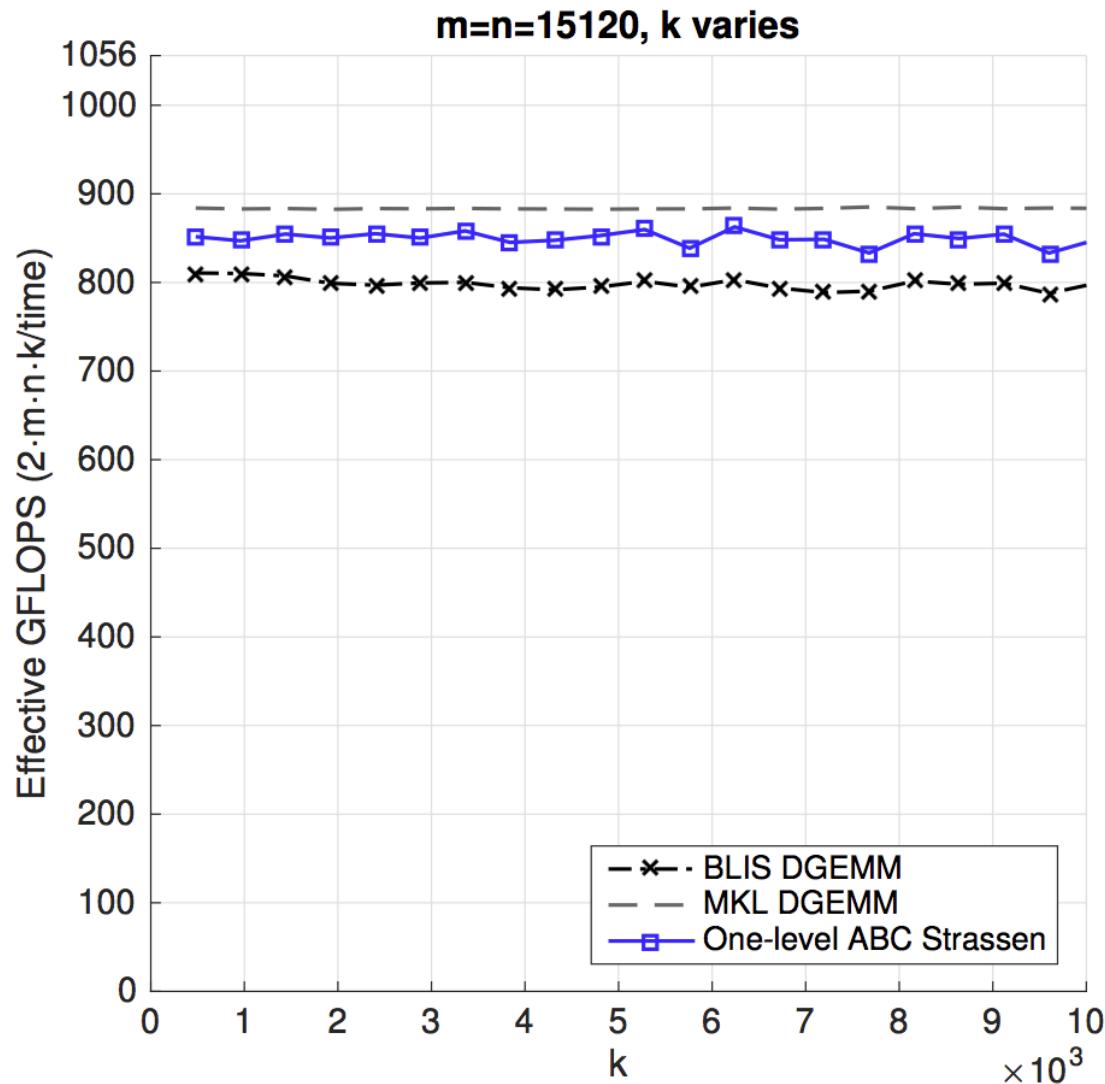


10 core

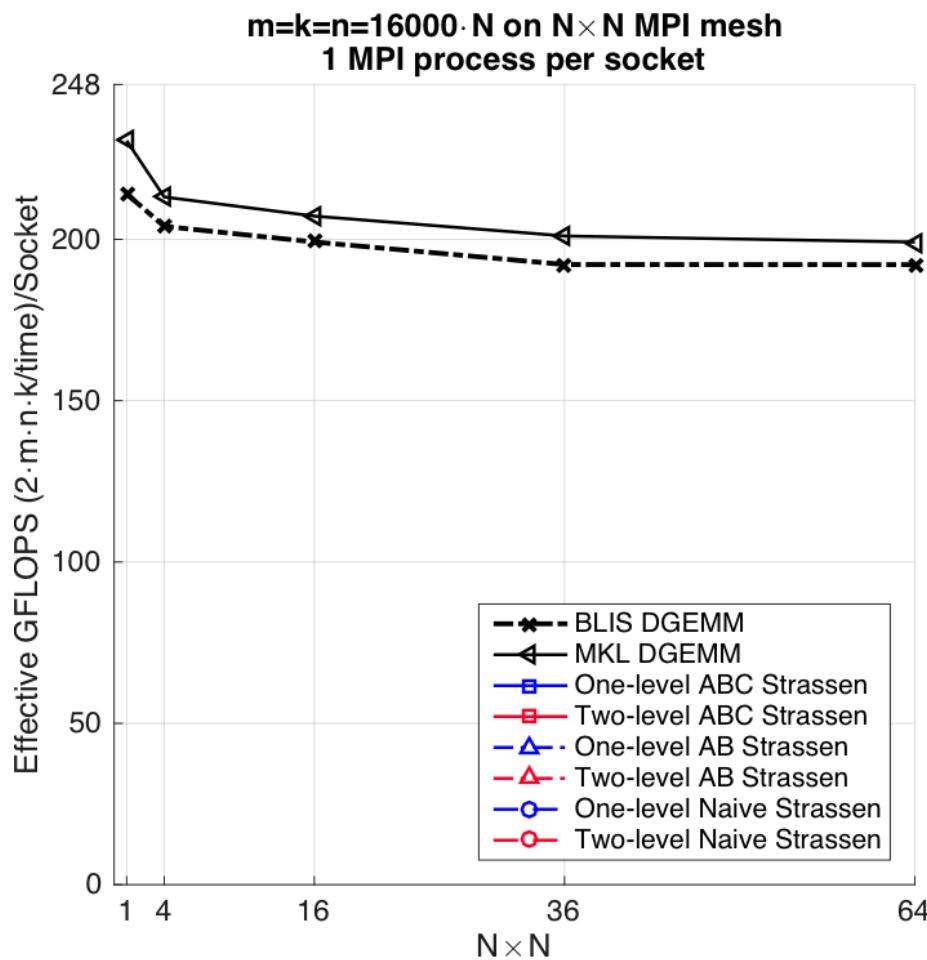
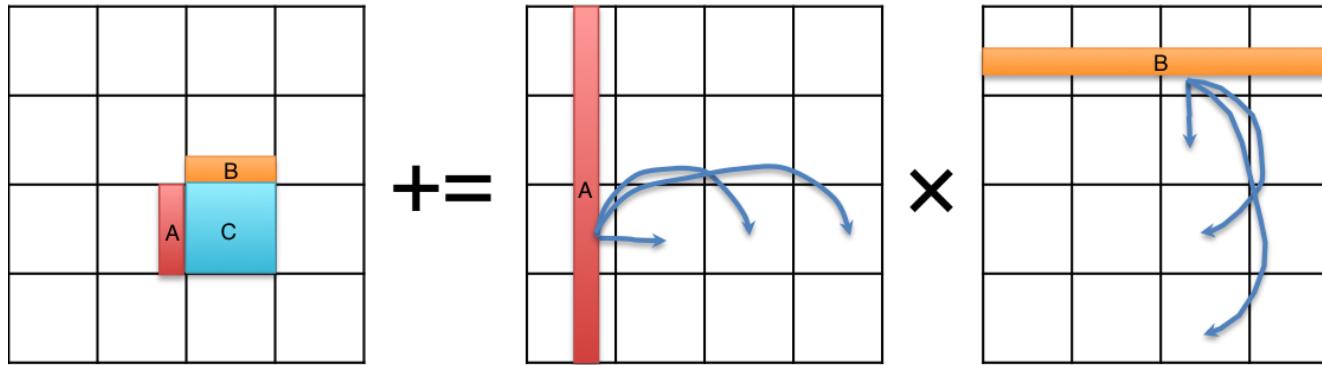


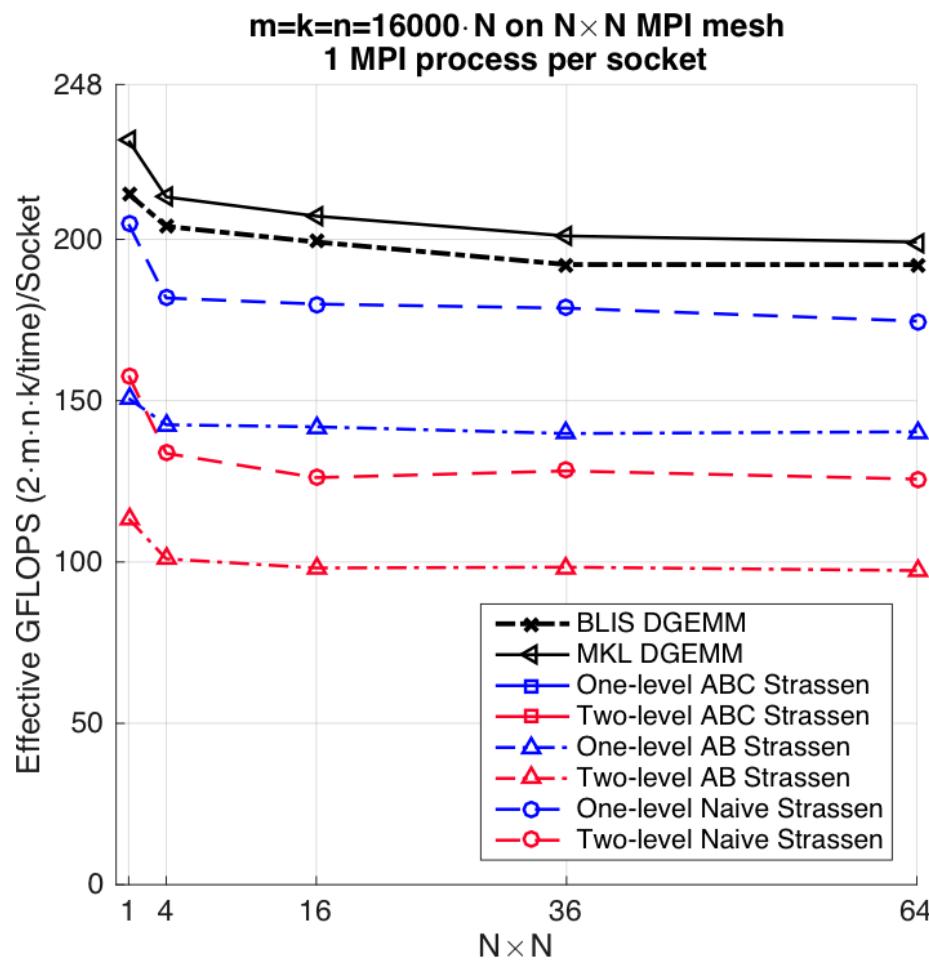
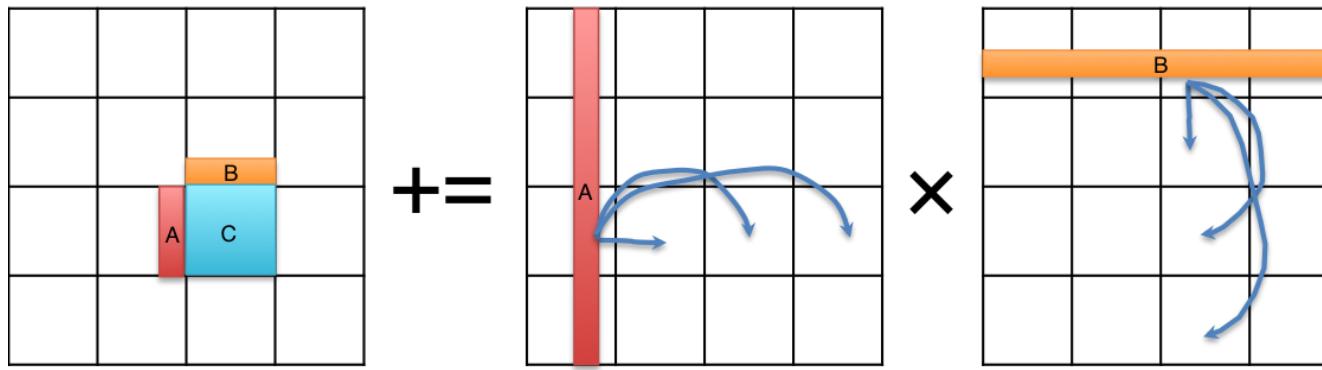
Many-core Experiment

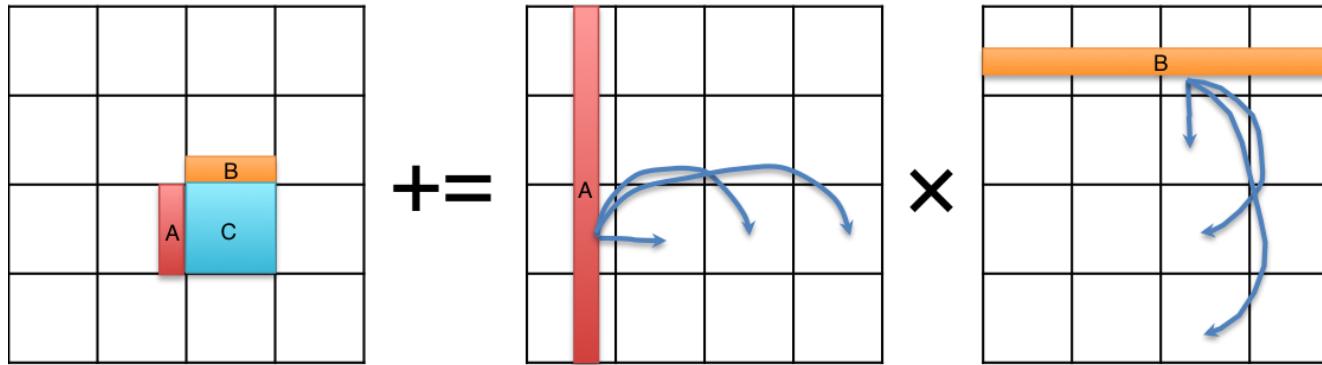
Intel® Xeon Phi™ coprocessor (KNC)



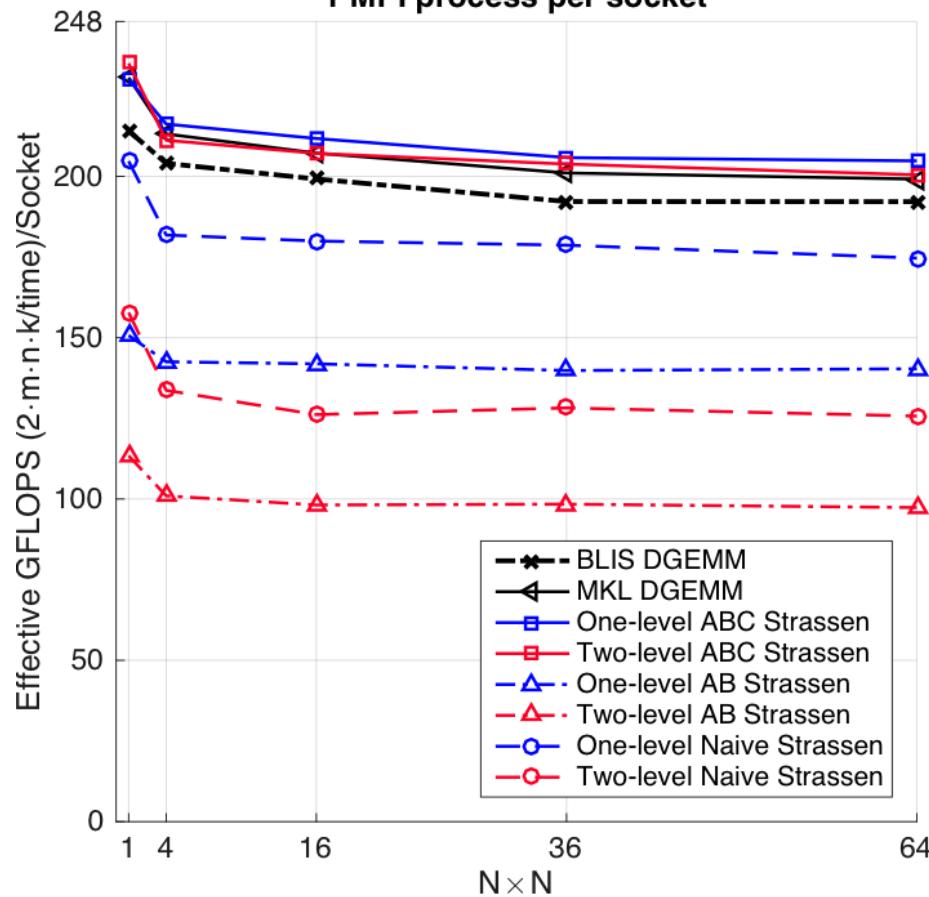
Distributed Memory Experiment







$m=k=n=16000 \cdot N$ on $N \times N$ MPI mesh
1 MPI process per socket



Outline

- Review of State-of-the-art GEMM in BLIS
- Strassen's Algorithm Reloaded
- Theoretical Model and Practical Performance
- Extension to Other BLAS-3 Operation
- Extension to Other Fast Matrix Multiplication
- Conclusion

Level-3 BLAS Symmetric Matrix-Matrix Multiplication (**SYMM**)

- Symmetric matrix-matrix multiplication (SYMM) is supported in the level-3 BLAS* interface as

```
dsymm( side, uplo, m, n,  
       alpha, A, lda, B, ldb,  
       beta, C, ldc )
```

- SYMM computes

$$C := \alpha AB + \beta C;$$

Level-3 BLAS Symmetric Matrix-Matrix Multiplication (**SYMM**)

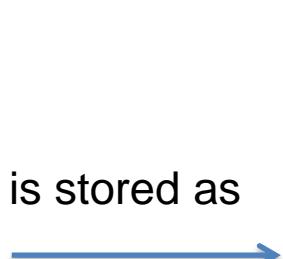
$$C := \alpha AB + \beta C;$$

- A is a symmetric matrix (square and equal to its transpose);
- Assume only the lower triangular half is stored;
- Follow the same algorithm we used for GEMM by modifying the packing routine for matrix A to account for the symmetric nature of A .

Example Matrix A

1	2	3	4	5	6	7	8
2	3	4	5	6	7	8	9
3	4	5	6	7	8	9	10
4	5	6	7	8	9	10	11
5	6	7	8	9	10	11	12
6	7	8	9	10	11	12	13
7	8	9	10	11	12	13	14
8	9	10	11	12	13	14	15

is stored as



1	0	0	0	0	0	0	0
2	3	0	0	0	0	0	0
3	4	5	0	0	0	0	0
4	5	6	7	0	0	0	0
5	6	7	8	9	0	0	0
6	7	8	9	10	11	0	0
7	8	9	10	11	12	13	0
8	9	10	11	12	13	14	15

SYMM with One-Level Strassen

- When partitioning A for one-level Strassen operations, we integrate the symmetric nature of A .

$$\begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 2 & 3 & 0 & 0 & 0 & 0 & 0 & 0 \\ 3 & 4 & 5 & 0 & 0 & 0 & 0 & 0 \\ 4 & 5 & 6 & 7 & 0 & 0 & 0 & 0 \\ \hline 5 & 6 & 7 & 8 & 9 & 0 & 0 & 0 \\ 6 & 7 & 8 & 9 & 10 & 11 & 0 & 0 \\ 7 & 8 & 9 & 10 & 11 & 12 & 13 & 0 \\ 8 & 9 & 10 & 11 & 12 & 13 & 14 & 15 \end{bmatrix} \xrightarrow{\hspace{10em}} \begin{array}{l} A_{00} \\ A_{01} \\ A_{10} \\ A_{11} \end{array}$$
$$A_{00} = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 2 & 3 & 0 & 0 \\ 3 & 4 & 5 & 0 \\ 4 & 5 & 6 & 7 \end{bmatrix} \quad A_{01} = \begin{bmatrix} 5 & 6 & 7 & 8 \\ 6 & 7 & 8 & 9 \\ 7 & 8 & 9 & 10 \\ 8 & 9 & 10 & 11 \end{bmatrix}$$
$$A_{10} = \begin{bmatrix} 5 & 6 & 7 & 8 \\ 6 & 7 & 8 & 9 \\ 7 & 8 & 9 & 10 \\ 8 & 9 & 10 & 11 \end{bmatrix} \quad A_{11} = \begin{bmatrix} 9 & 0 & 0 & 0 \\ 10 & 11 & 0 & 0 \\ 11 & 12 & 13 & 0 \\ 12 & 13 & 14 & 15 \end{bmatrix}$$

SYMM Implementation Platform: BLISlab*

What is BLISlab?

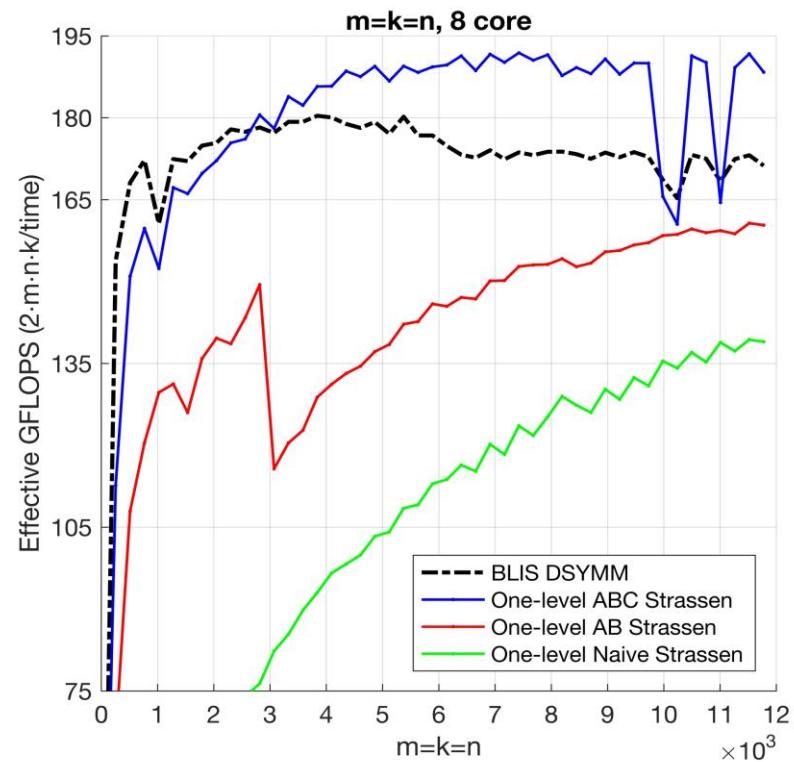
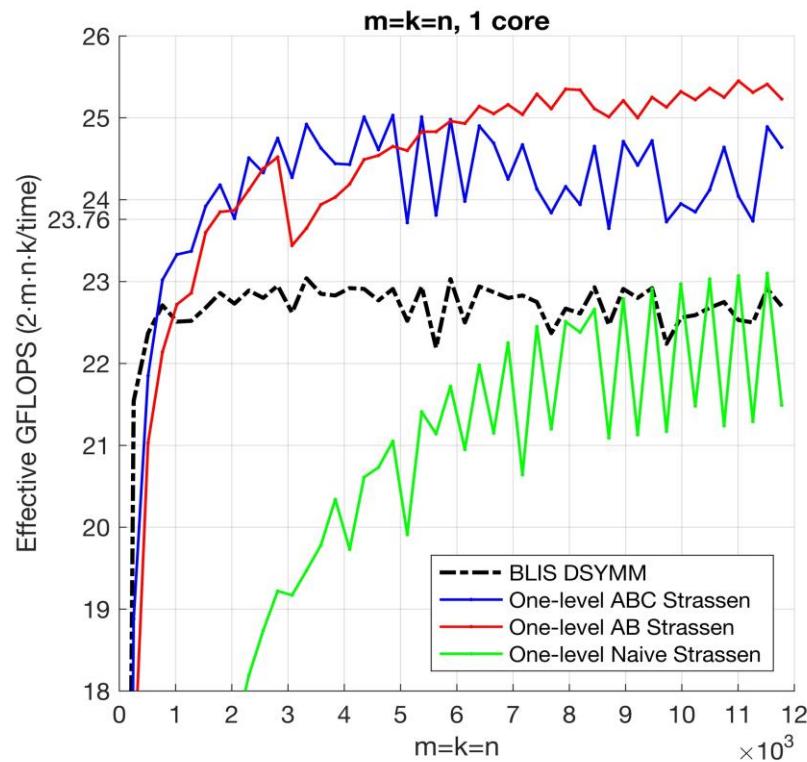
- A Sandbox for Optimizing GEMM that mimics implementation in BLIS
- A set of exercises that use GEMM to show how high performance can be attained on modern CPUs with hierarchical memories
- Used in Prof. Robert van de Geijn's CS 383C Numerical Linear Algebra

Contributions

- Added DSYMM
- Applied Strassen's algorithm to DGEMM and DSYMM
- 3 versions: Naive, AB, ABC
- Strassen improves performance in both DGEMM and DSYMM

* <https://github.com/flame/blislab>

DSYMM Results in BLISlab

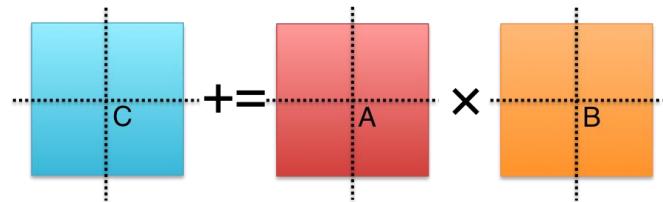


Outline

- Review of State-of-the-art GEMM in BLIS
- Strassen's Algorithm Reloaded
- Theoretical Model and Practical Performance
- Extension to Other BLAS-3 Operation
- **Extension to Other Fast Matrix Multiplication**
- Conclusion

(2,2,2) Strassen Algorithm

$M_0 := (A_{00} + A_{11})(B_{00} + B_{11}); \quad C_{00} += M_0; \quad C_{11} += M_0;$
 $M_1 := (A_{10} + A_{11})B_{00}; \quad C_{10} += M_1; \quad C_{11} -= M_1;$
 $M_2 := A_{00}(B_{01} - B_{11}); \quad C_{01} += M_2; \quad C_{11} += M_2;$
 $M_3 := A_{11}(B_{10} - B_{00}); \quad C_{00} += M_3; \quad C_{10} += M_3;$
 $M_4 := (A_{00} + A_{01})B_{11}; \quad C_{01} += M_4; \quad C_{00} -= M_4;$
 $M_5 := (A_{10} - A_{00})(B_{00} + B_{01}); \quad C_{11} += M_5;$
 $M_6 := (A_{01} - A_{11})(B_{10} + B_{11}); \quad C_{00} += M_6;$



$$\mathbf{U} = \begin{bmatrix} 1 & 0 & 1 & 0 & 1 & -1 & 0 \\ 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 & 0 & 1 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & -1 \end{bmatrix}$$

$$\mathbf{V} = \begin{bmatrix} 1 & 1 & 0 & -1 & 0 & 1 & 0 \\ 0 & 0 & 1 & 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 & 0 & 0 & 1 \\ 1 & 0 & -1 & 0 & 1 & 0 & 1 \end{bmatrix}$$

$$\mathbf{W} = \begin{bmatrix} 1 & 0 & 0 & 1 & -1 & 0 & 1 \\ 0 & 0 & 1 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 \\ 1 & -1 & 1 & 0 & 0 & 1 & 0 \end{bmatrix}$$

(3,2,3) Fast Matrix Multiplication

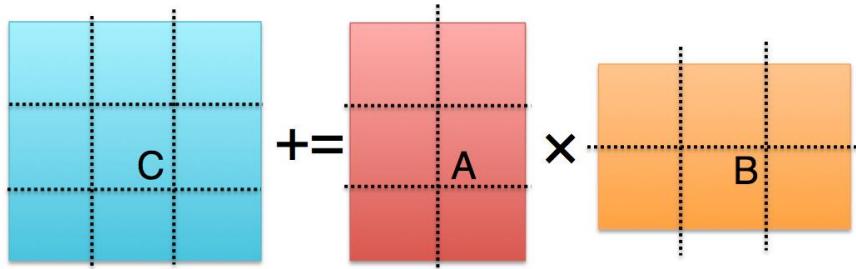
$M_0 := (-A_{01} + A_{11} + A_{21})(B_{11} + B_{12});$
 $M_1 := (-A_{00} + A_{10} + A_{21})(B_{00} - B_{12});$
 $M_2 := (-A_{00} + A_{10})(-B_{01} - B_{11});$
 $M_3 := (A_{10} - A_{20} + A_{21})(B_{02} - B_{10} + B_{12});$
 $M_4 := (-A_{00} + A_{01} + A_{10})(B_{00} + B_{01} + B_{11});$
 $M_5 := (A_{10})(B_{00});$
 $M_6 := (-A_{10} + A_{11} + A_{20} - A_{21})(B_{10} - B_{12});$
 $M_7 := (A_{11})(B_{10});$
 $M_8 := (-A_{00} + A_{10} + A_{20})(B_{01} + B_{02});$
 $M_9 := (-A_{00} + A_{01})(-B_{00} - B_{01});$
 $M_{10} := (A_{21})(B_{02} + B_{12});$
 $M_{11} := (A_{20} - A_{21})(B_{02});$
 $M_{12} := (A_{01})(B_{00} + B_{01} + B_{10} + B_{11});$
 $M_{13} := (A_{10} - A_{20})(B_{00} - B_{02} + B_{10} - B_{12});$
 $M_{14} := (-A_{00} + A_{01} + A_{10} - A_{11})(B_{11});$

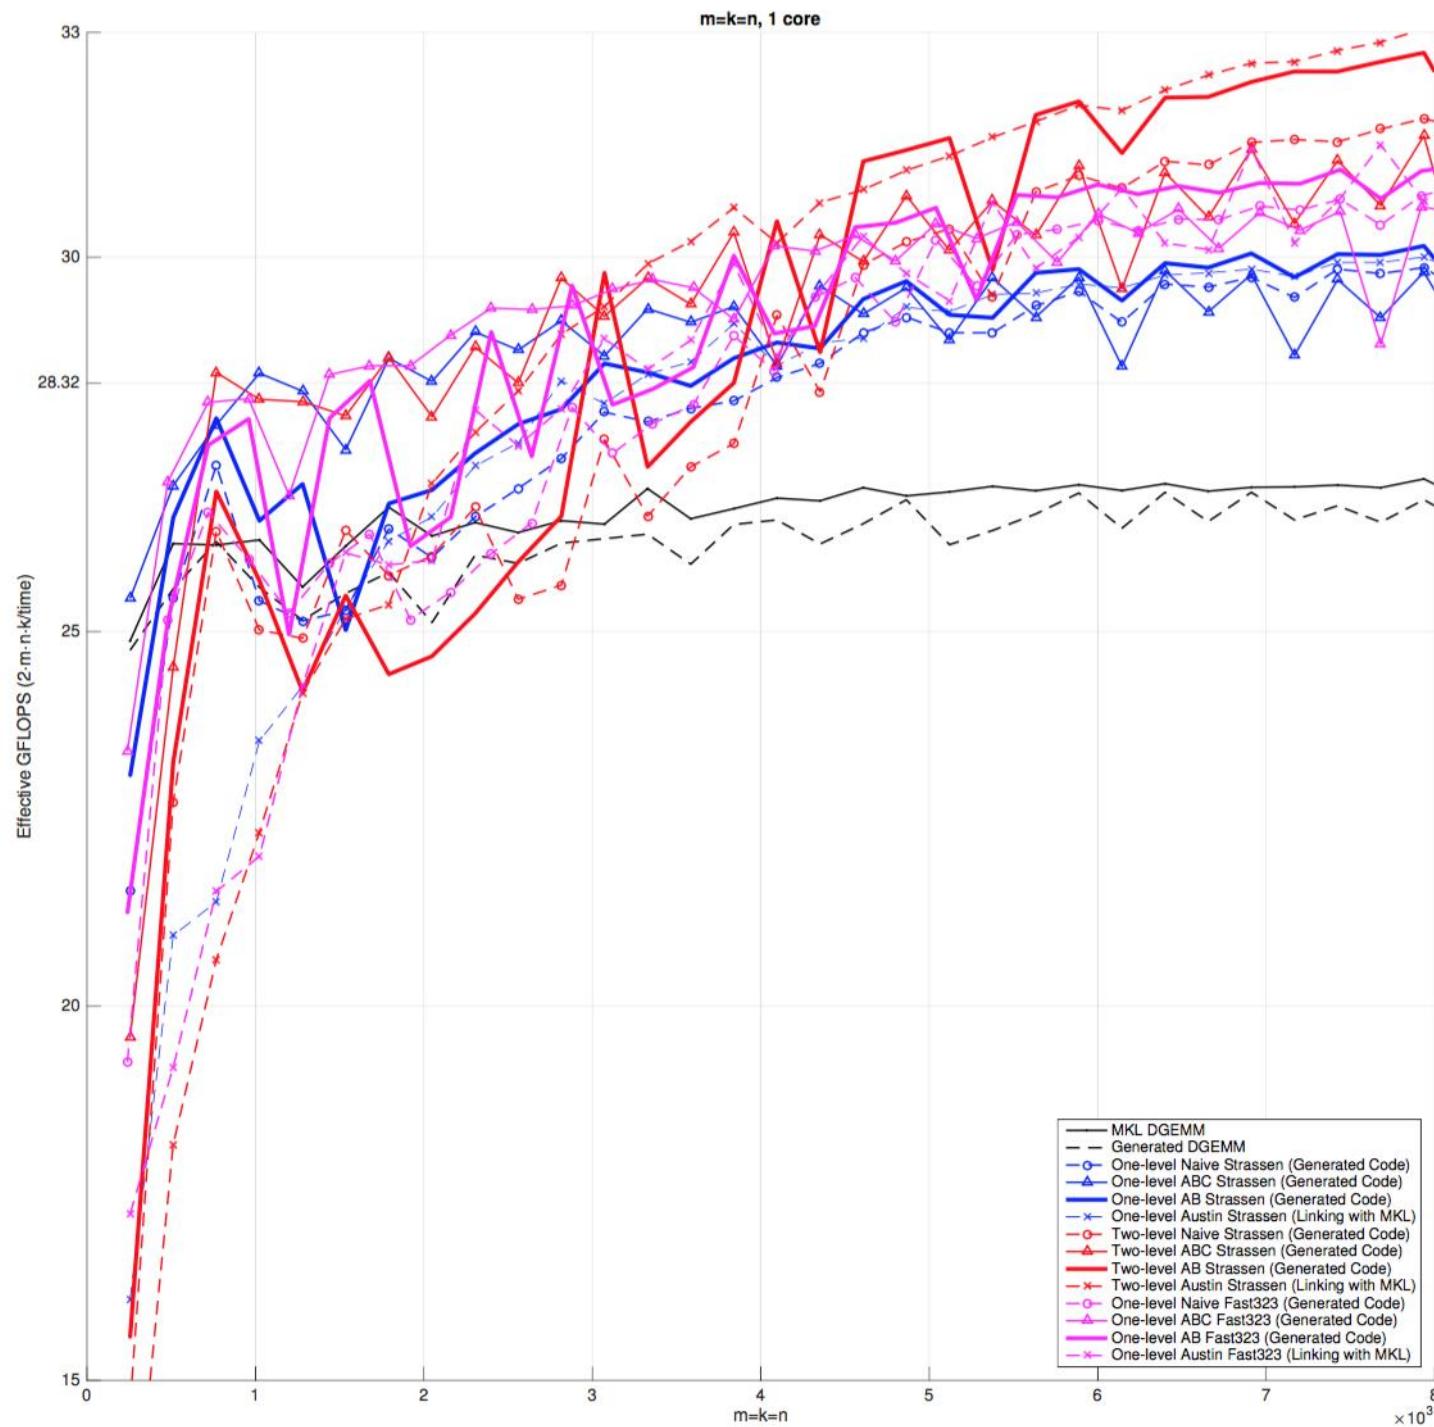
$C_{02} := M_0;$
 $C_{02} := M_1; C_{21} := M_1;$
 $C_{00} := M_2; C_{01} := M_2; C_{21} := M_2;$
 $C_{02} := M_3; C_{12} := M_3; C_{20} := M_3;$
 $C_{00} := M_4; C_{01} := M_4; C_{11} := M_4;$
 $C_{00} := M_5; C_{01} := M_5; C_{10} := M_5; C_{11} := M_5; C_{20} := M_5;$
 $C_{02} := M_6; C_{12} := M_6;$
 $C_{02} := M_7; C_{10} := M_7; C_{12} := M_7;$
 $C_{21} := M_8;$
 $C_{01} := M_9; C_{11} := M_9;$
 $C_{02} := M_{10}; C_{20} := M_{10}; C_{22} := M_{10};$
 $C_{02} := M_{11}; C_{12} := M_{11}; C_{21} := M_{11}; C_{22} := M_{11};$
 $C_{00} := M_{12};$
 $C_{20} := M_{13};$
 $C_{02} := M_{14}; C_{11} := M_{14};$

$$\mathbf{U} = \begin{bmatrix} 0 & -1 & -1 & 0 & -1 & 0 & 0 & 0 & -1 & -1 & 0 & 0 & 0 & 0 & 0 & -1 \\ -1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 1 & 1 \\ 0 & 1 & 1 & 1 & 1 & 1 & -1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 1 \\ 1 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & -1 \\ 0 & 0 & 0 & -1 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 1 & 0 & 0 \\ 1 & 1 & 0 & 1 & 0 & 0 & -1 & 0 & 0 & 0 & 0 & 0 & 1 & -1 & 0 & 0 \end{bmatrix}$$

$$\mathbf{V} = \begin{bmatrix} 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & -1 & 0 & 0 & 1 & 1 & 0 & 0 \\ 0 & 0 & -1 & 0 & 1 & 0 & 0 & 0 & 0 & 1 & -1 & 0 & 0 & 1 & 0 & 0 \\ 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 1 & 0 & -1 & 0 \\ 0 & 0 & 0 & -1 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 \\ 1 & -1 & -1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 \\ 1 & 0 & 0 & 1 & 0 & 0 & -1 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & -1 \end{bmatrix}$$

$$\mathbf{W} = \begin{bmatrix} 0 & 0 & -1 & 0 & -1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 1 & 0 & 1 & -1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 \\ -1 & -1 & 0 & 1 & 0 & 0 & -1 & 1 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & -1 \\ 0 & 0 & 0 & 0 & 0 & 1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 1 & -1 & 0 & 0 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & -1 \\ 0 & 0 & 0 & 1 & 0 & 0 & -1 & 1 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & 0 \\ 0 & 0 & 0 & -1 & 0 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & -1 \\ 0 & -1 & 1 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 0 & 0 & -1 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 & 1 & 1 & 0 & 0 & 0 & 0 \end{bmatrix}$$

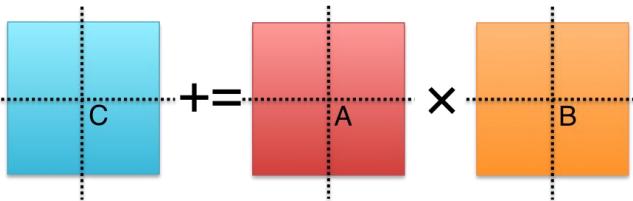




Outline

- Review of State-of-the-art GEMM in BLIS
- Strassen's Algorithm Reloaded
- Theoretical Model and Practical Performance
- Extension to Other BLAS-3 Operation
- Extension to Other Fast Matrix Multiplication
- Conclusion

To achieve practical high performance of Strassen's algorithm.....



Conventional
Implementations

Our
Implementations

Matrix Size

Must be large



$$c \quad += \quad A \times B$$



Matrix Shape

Must be square



$$C \quad += \quad A \times B$$



No Additional
Workspace



Parallelism

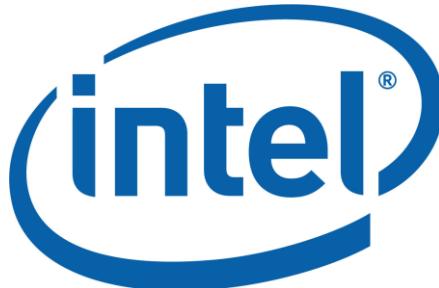
Usually task parallelism



Can be data parallelism



Acknowledgement



- NSF grants ACI-1148125/1340293, CCF-1218483.
- Intel Corporation through an Intel Parallel Computing Center (IPCC).
- Access to the Maverick and Stampede supercomputers administered by TACC.

We thank Field Van Zee, Chenhan Yu, Devin Matthews, and the rest of the SHPC team (<http://shpc.ices.utexas.edu>) for their supports.

Any opinions, findings, and conclusions or recommendations expressed in this material are those of the author(s) and do not necessarily reflect the views of the National Science Foundation.

Thank you!